

## CHAPTER

# 17

## Switch Expressions, Records, and Other Recently Added Features

A key attribute of Java has been its ability to adapt to the increasing demands of the modern computing environment. Over the years, Java has incorporated many new features, each responding to changes in the computing environment or to innovations in computer language design. This ongoing process has enabled Java to remain one of the world's most important and popular computer languages. As explained earlier, this book has been updated for JDK 17, which is a long-term support (LTS) version of Java. JDK 17 incorporates a number of new language features that have been added to Java since the previous LTS version, which was JDK 11. A few of the smaller additions have been described in the preceding chapters. In this chapter, several major additions are examined. They are

- Enhancements to **switch**
- Text blocks
- Records
- Patterns in **instanceof**
- Sealed classes and interfaces

Here is a brief description of each. The **switch** has been enhanced in a number of ways, the most impacting of which is the *switch expression*. A **switch** expression enables a **switch** to produce a value. *Text blocks* allow a string literal to occupy more than a single line. Supported by the new keyword **record**, records enable you to create a class that is specifically designed to hold a group of values. A second form of **instanceof** has been added that uses a type pattern. With this form, you can specify a variable that receives an instance of the type being tested if **instanceof** succeeds. It is now possible to specify a *sealed* class or interface. A sealed class can be inherited by only explicitly specified subclasses. A sealed interface can be implemented by only explicitly specified classes or extended by only explicitly specified interfaces. Thus, sealing a class or interface gives you fine-grained control over its inheritance and implementation.

## Enhancements to switch

The **switch** statement has been part of Java since the start. It is a crucial element of Java's program control statements and provides for a multiway branch. Moreover, **switch** is so fundamental to programming that it is found in one form or another in other popular programming languages. The traditional form of **switch** was described in Chapter 5. This is the form of **switch** that has always been part of Java. Beginning with JDK 14, **switch** has been substantially enhanced with the addition of four new features, shown here:

- The **switch** expression
- The **yield** statement
- Support for a list of **case** constants
- The **case** with an arrow

Each new feature is examined in detail in the discussions that follow, but here is a brief description: The **switch** expression is, essentially, a **switch** that produces a value. Thus, a **switch** expression can be used on the right side of an assignment, for example. The **yield** statement specifies a value that is produced by a **switch** expression. It is now possible to have more than one **case** constant in a **case** statement through the use of a list of constants. A second form of **case** has been added that uses an arrow (->) instead of a colon. The arrow gives **case** new capabilities.

Collectively, the enhancements to **switch** represent a fairly significant change to the Java language. Not only do they provide new capabilities, but in some situations, they also offer superior alternatives to traditional approaches. Because of this, a solid understanding of the "how" and "why" behind the new **switch** features is important.

One of the best ways to understand the **switch** enhancements is to start with an example that uses a traditional **switch** and then gradually incorporate each new feature. This way, the use and benefit of the enhancements will be clearly apparent. To begin, imagine some device that produces integer codes that indicate various events and you want to associate a priority level with each event code. Most events will have a normal priority, but a few will have a higher priority. Here is a program that uses a traditional **switch** statement to supply a priority level given an event code:

```
// Use a traditional switch to set a priority level based on which
// event code is matched.
class TraditionalSwitch {

    public static void main(String[] args) {
        int priorityLevel;

        int eventCode = 6010;

        // A traditional switch that supplies a value associated
        // with a case.
        switch(eventCode) {
            case 1000: // In this traditional switch, case stacking is used.
```

```
case 1205:  
case 8900:  
    priorityLevel = 1;  
    break;  
case 2000:  
case 6010:  
case 9128:  
    priorityLevel = 2;  
    break;  
case 1002:  
case 7023:  
case 9300:  
    priorityLevel = 3;  
    break;  
default: // normal priority  
    priorityLevel = 0;  
}  
  
System.out.println("Priority level for event code " + eventCode +  
                    " is " + priorityLevel);  
}  
}
```

The output is shown here:

```
Priority level for event code 6010 is 2
```

There is certainly nothing wrong with using a traditional **switch** as shown in the program, and this is the way Java code has been written for more than two decades. However, as the following sections will show, in many cases, the traditional **switch** can be improved by use of the enhanced **switch** features.

## Use a List of case Constants

We begin with one of the easiest ways to modernize a traditional **switch**: by use of a list of **case** constants. In the past, when two constants were both handled by the same code sequence, **case stacking** was employed, and this is the approach used by the preceding program. For example, here are how the **cases** for 1000, 1205, and 8900 are handled:

```
case 1000: // In this traditional switch, case stacking is used.  
case 1205:  
case 8900:  
    priorityLevel = 1;  
    break;
```

The stacking of **case** statements enable all three **case** statements to use the same code sequence to set **priorityLevel** to 1. As explained in Chapter 5, in a traditional-style **switch**, the stacking of **cases** is made possible because execution falls through each **case** until a **break** is encountered. Although this approach works, a more elegant solution can be achieved by use of a **case constant list**.

Beginning with JDK 14, you can specify more than one **case** constant in a single **case**. To do so, simply separate each constant with a comma. For example, here is a more compact way to code the **case** for 1000, 1205, and 8900:

```
case 1000, 1205, 8900: // use a case list
    priorityLevel = 1;
    break;
```

Here is the entire **switch**, rewritten to use lists of **case** constants:

```
// This switch uses a list of constants with each case.
switch(eventCode) {
    case 1000, 1205, 8900:
        priorityLevel = 1;
        break;
    case 2000, 6010, 9128:
        priorityLevel = 2;
        break;
    case 1002, 7023, 9300:
        priorityLevel = 3;
        break;
    default: // normal priority
        priorityLevel = 0;
}
```

As you can see, the number of **case** statements has been reduced by six, making the **switch** easier to read and a bit more manageable. Although support for a **case** constant list does not by itself add any fundamentally new functionality to the **switch**, it does help streamline your code. In many situations, it also offers an easy way to improve existing code—especially when extensive **case**-stacking was previously employed. Thus, it is a feature that you can put to work immediately, with minimal code rewriting.

## Introducing the **switch** Expression and the **yield** Statement

Of the enhancements to **switch**, the one that will have the most profound impact is the **switch expression**. A **switch** expression is, essentially, a **switch** that returns a value. Thus, it has all of the capabilities of a traditional **switch** statement, plus the ability to produce a result. This added capability makes the **switch** expression one of the more important additions to Java in recent years.

One way to supply the value of a **switch** expression is with the new **yield** statement. It has this general form:

```
yield value;
```

Here, *value* is the value produced by the **switch**, and it can be any expression compatible with the type of value required. A key point to understand about **yield** is that it immediately terminates the **switch**. Thus, it works somewhat like **break**, with the added capability of supplying a value. It is important to point out that **yield** is a context-sensitive keyword. This means that outside its use in a **switch** expression, **yield** is simply an identifier with no special meaning. However, if you use a method called **yield()**, it must be qualified. For example, if **yield()** is a non-**static** method within its class, you must use **this.yield()**.

It is very easy to specify a **switch** expression. Simply use the **switch** in a context in which a value is required, such as on the right side of an assignment statement, an argument to a method, or a return value. For example, this line indicates that a **switch** expression is being employed:

```
int x = switch(y) { // ...
```

Here, the **switch** result is being assigned to the **x** variable. A key point about using a **switch** expression is that each **case** (plus **default**) must produce a value (unless it throws an exception). In other words, each path through a **switch** expression must produce a result.

The addition of the **switch** expression simplifies the coding of situations in which each **case** sets the value of some variable. Such situations can occur in a number of different ways. For example, each **case** might set a **boolean** variable that indicates the success or failure of some action taken by the **switch**. Often, however, setting a variable is the *primary purpose* of the **switch**, as is the case with the **switch** used by the preceding program. Its job is to produce the priority level associated with an event code. With a traditional **switch** statement, each **case** statement must individually assign a value to the variable, and this variable becomes the de facto result of the **switch**. This is the approach used by the preceding programs, in which the value of the variable **priorityLevel** is set by each **case**. Although this approach has been used in Java programs for decades, the **switch** expression offers a better solution because the desired value is produced by the **switch** itself.

The following version of the program puts the preceding discussion into action by changing the **switch** statement into a **switch** expression:

```
// Convert a switch statement into a switch expression.
class SwitchExpr {

    public static void main(String[] args) {
        int eventCode = 6010;

        // This is a switch expression. Notice how its value is assigned
        // to the priorityLevel variable. Also notice how the value of the
        // switch is supplied by the yield statement.
        int priorityLevel = switch(eventCode) {
            case 1000, 1205, 8900:
                yield 1;
            case 2000, 6010, 9128:
                yield 2;
            case 1002, 7023, 9300:
                yield 3;
            default: // normal priority
                yield 0;
        };

        System.out.println("Priority level for event code " + eventCode +
                           " is " + priorityLevel);
    }
}
```

Look closely at the **switch** in the program. Notice that it differs in important ways from the one used in the previous examples. Instead of each **case** assigning a value to **priorityLevel** individually, this version assigns the outcome of the **switch** itself to the **priorityLevel** variable.

Thus, only one assignment to **priorityLevel** is required, and the length of the **switch** is reduced. Using a **switch** expression also ensures that each **case** yields a value, thus avoiding the possibility of forgetting to give **priorityLevel** a value in one of the **cases**. Notice that the value of the **switch** is produced by the **yield** statement inside each **case**. As explained, **yield** causes immediate termination of the **switch**, so no fall-through from **case** to **case** will occur. Thus, no **break** statement is required, or allowed. One other thing to notice is the semicolon after the closing brace of the **switch**. Because this **switch** is used in an assignment, it must be terminated by a semicolon.

There is an important restriction that applies to a **switch** expression: the **case** statements must handle all of the values that might occur. Thus, a **switch** expression must be *exhaustive*. For example, if its controlling expression is of type **int**, then all **int** values must be handled by the **switch**. This would, of course, constitute a very large number of **case** statements! For this reason, most **switch** expressions will have a **default** statement. The exception to this rule is when an enumeration is used, and each value of the enumeration is matched by a **case**.

## Introducing the Arrow in a **case** Statement

Although the use of **yield** in the preceding program is a perfectly valid way to specify a value for a **switch** expression, it is not the only way to do so. In many situations, an easier way to supply a value is through the use of a new form of the **case** that substitutes **->** for the colon in a **case**. For example, this line:

```
case 'X': // ...
```

can be rewritten using the arrow like this:

```
case 'X' -> // ...
```

To avoid confusion, in this discussion we will refer to a **case** with an arrow as an *arrow case* and the traditional, colon-based form as a *colon case*. Although both forms will match the character X, the precise action of each style of **case** statement differs in three very important ways.

First, one arrow **case** *does not* fall through to the next **case**. Thus, there is no need to use **break**. Execution simply terminates at the end of an arrow **case**. Although the fall-through nature of a traditional colon-based **case** has always been part of Java, fall-through has been criticized because it can be a source for bugs, such as when the programmer forgets to add a **break** statement to prevent fall-through when fall-through is not desired. The arrow **case** avoids this situation. Second, the arrow **case** provides a “shorthand” way to supply a value when used in a **switch** expression. For this reason, the arrow **case** is often used in **switch** expressions. Third, the target of an arrow **case** must be either an expression, a block, or throw an exception. It cannot be a statement sequence, as is allowed with a traditional **case**. Thus, the arrow **case** will have one of these general forms:

```
case constant -> expression;
case constant -> { block-of-statements }
case constant -> throw ...
```

Of course, the first two forms represent the primary uses.

Arguably, the most common use of an arrow **case** is in a **switch** expression, and the most common target of the arrow **case** is an expression. Thus, it is here that we will begin. When the target of an arrow **case** is an expression, the value of that expression becomes the value of the **switch** expression when that **case** is matched. Thus, it provides a very efficient alternative to the **yield** statement in many situations. For example, here is the first **case** in the event code example rewritten to use an arrow **case**:

```
case 1000, 1205, 8900 -> 1;
```

Here, the value of the expression (which is 1) automatically becomes the value produced by the **switch** when this **case** is matched. In other words, the expression becomes the value yielded by the **switch**. Notice that this statement is quite compact, yet clearly expresses the intent to supply a value.

In the following program, the entire **switch** expression has been completely rewritten to use the arrow **case**:

```
// Use the arrow "shorthand" to supply the priority level.
class SwitchExpr2 {

    public static void main(String[] args) {
        int eventCode = 6010;

        // In this switch expression, notice how the value is supplied
        // by use of an arrow case. Notice that no break statement is
        // required (or allowed) to prevent fall-through.
        int priorityLevel = switch(eventCode) {
            case 1000, 1205, 8900 -> 1;
            case 2000, 6010, 9128 -> 2;
            case 1002, 7023, 9300 -> 3;
            default -> 0; // normal priority
        };

        System.out.println("Priority level for event code " + eventCode +
                           " is " + priorityLevel);
    }
}
```

This produces the same output as before. Looking at the code, it is easy to see why this form of the arrow **case** is so appropriate for many types of **switch** expressions. It is compact and eliminates the need for a separate **yield** statement. Because the arrow **case** does not fall through, there is no need for a **break** statement. Each **case** terminates by yielding the value of its expression. Furthermore, if you compare this final version of the **switch** to the original, traditional **switch** shown at the start of this discussion, it is readily apparent how streamlined and expressive this version is. In combination, the **switch** enhancements offer a truly impressive way to improve the clarity and resiliency of your code.

## A Closer Look at the Arrow case

The arrow **case** provides considerable flexibility. First, when using its expression form, the expression can be of any type. For example, the following is a valid **case** statement:

```
case -1 -> getErrorCode();
```

Here, the result of the call to **getErrorCode()** becomes the value of the enclosing **switch** expression. Here is another example:

```
case 0 -> normalCompletion = true;
```

In this case, the result of the assignment, which is **true**, becomes the value yielded. The key point is that any valid expression can be used as the target of the arrow **case** as long as it is compatible with the type required by the **switch**.

As mentioned, the target of the `->` can also be a block of code. You will need to use a block as the target of an arrow **case** whenever you need more than a single expression. For example, each **case** in this version of the event code program sets the value of a **boolean** variable called **stopNow** to indicate if immediate termination is required and then yields the priority level.

```
// Use blocks with an arrow.
class BlockArrowCase {

    public static void main(String[] args) {
        boolean stopNow;

        int eventCode = 9300;

        // Use code blocks with an arrow. Again, notice
        // that no break statement is required (or allowed)
        // to prevent fall through. Because the target of an
        // arrow is a block, yield must be used to supply a value.
        int priorityLevel = switch(eventCode) {
            case 1000, 1205, 8900 -> { // use a block with an arrow
                stopNow = false;
                System.out.println("Alert");
                yield 1;
            }
            case 2000, 6010, 9128 -> {
                stopNow = false;
                System.out.println("Warning");
                yield 2;
            }
            case 1002, 7023, 9300 -> {
                stopNow = true;
                System.out.println("Danger");
                yield 3;
            }
        }
    }
}
```

```
    default -> {
        stopNow = false;
        System.out.println("Normal.");
        yield 0;
    }
};

System.out.println("Priority level for event code " + eventCode +
    " is " + priorityLevel);
if(stopNow) System.out.println("Stop required.");
}
```

Here is the output:

```
Danger
Priority level for event code 9300 is 3
Stop required.
```

As this example shows, when using a block, you must use **yield** to supply a value to a **switch** expression. Furthermore, even though block targets are used, each path through the **switch** expression must still provide a value.

Although the preceding program provides a simple illustration of a block target of an arrow **case**, it also raises an interesting question. Notice that each **case** in the **switch** sets the value of two variables. The first is **priorityLevel**, which is the value yielded. The second is **stopNow**. Is there a way for a **switch** expression to yield more than one value? In a direct sense, the answer is “no” because only one value can be produced by the **switch**. However, it is possible to encapsulate two or more values within a class and yield an object of that class. Beginning with JDK 16, Java provides an especially streamlined and efficient way to accomplish this: the **record**. Described later in this chapter, a **record** aggregates two or more values into a single logical unit. As it relates to this example, a **record** could hold both the **priorityLevel** and the **stopNow** values, and this **record** could be yielded by the **switch** as a unit. Thus, a **record** offers a convenient way for a **switch** to yield more than a single value.

Although the arrow **case** is very helpful in a **switch** expression, it is important to emphasize that it is not limited to that use. The arrow **case** can also be used in a **switch** statement, which enables you to write **switches** in which no **case** fall-through can occur. In this situation, no **yield** statement is required (or allowed), and no value is produced by the **switch**. In essence, it works much like a traditional **switch** but without the fall-through. Here is an example:

```
// Use case arrows with a switch statement
class StatementSwitchWithArrows {

    public static void main(String[] args) {
        int up = 0;
        int down = 0;
        int left = 0;
        int right = 0;

        char direction = 'R';
```

```

// Use arrows with a switch statement. Notice that
// no value is produced.
switch(direction) {
    case 'L' -> {
        System.out.println("Turning Left");
        left++;
    }
    case 'R' -> {
        System.out.println("Turning Right");
        right++;
    }
    case 'U' -> {
        System.out.println("Moving Up");
        up++;
    }
    case 'D' -> {
        System.out.println("Moving Down");
        down++;
    }
}
System.out.println(right);
}
}

```

In this program, the **switch** is a statement, not an expression. This is because of two reasons. First, no value is produced. Second, it is not exhaustive because no **default** statement is included. (Recall that **switch** expressions must be exhaustive, but not **switch** statements.) Notice, however, that because no fall-through occurs with an arrow **case**, no **break** statement is needed. As a point of interest, because each **case** increases the value of a different variable, it would not be possible to transform this **switch** into an expression. What value would it produce? All four **cases** increment a different variable.

One last point: you cannot mix arrow **cases** with traditional, colon **cases** in the same **switch**. You must choose one or the other. For example, this sequence is invalid:

```

// This won't work! You cannot mix a colon case with an arrow case.
switch(direction) {
    case 'L' -> {
        System.out.println("Turning Left");
        left++;
    }
    case 'R' : // Wrong! Can't mix case styles.
        System.out.println("Turning Right");
        right++;
        break;
    case 'U' -> {
        System.out.println("Moving Up");
        up++;
    }
    case 'D' -> {
        System.out.println("Moving Down");
        down++;
    }
}

```

## Another switch Expression Example

To conclude this overview of the **switch** enhancements, another example is presented. It uses a **switch** expression to determine whether a letter is an English-language vowel. It makes use of all of the new **switch** features. Pay special attention to the way Y is handled. In English, Y can be a vowel or a consonant. The program lets you specify which way you want the Y interpreted by the way the **yIsVowel** variable is set. To handle this special case, a block is used as the target of the **->**.

```
// Use a switch expression to determine if a character is an English vowel.  
// Notice the use of a block as the target of an arrow case for Y.  
  
class Vowels {  
  
    public static void main(String[] args) {  
  
        // If Y is to be counted as a vowel, set this  
        // variable to true.  
        boolean yIsVowel = true;  
  
        char ch = 'Y';  
  
        boolean isVowel = switch(ch) {  
            case 'a', 'e', 'i', 'o', 'u', 'A', 'E', 'I', 'O', 'U' -> true;  
            case 'y', 'Y' -> { if(yIsVowel) yield true; else yield false; }  
            default -> false;  
        };  
  
        if(isVowel) System.out.println(ch + " is a vowel.");  
    }  
}
```

As an experiment, try rewriting this program using a traditional **switch**. As you will find, doing so results in a much longer, less manageable version. The new **switch** enhancements often provide a superior approach.

## Text Blocks

Beginning with JDK 15, Java provides support for *text blocks*. A text block is a new kind of literal that is comprised of a sequence of characters that can occupy more than one line. A text block reduces the tedium programmers often face when creating long string literals because newline characters can be used in a text block without the need for the **\n** escape sequence. Furthermore, tab and double quote characters can also be entered directly, without using an escape sequence, and the indentation of a multiline string can be preserved. Although text blocks may, at first, seem to be a relatively small addition to Java, they may well become one of the most popular features.

## Text Block Fundamentals

A text block is supported by a new delimiter, which is three double-quote characters: **"""**. A text block is created by enclosing a string within a set of these delimiters. Specifically, a text block

begins immediately following the newline after the opening """". Thus, the line containing the opening delimiter must end with a newline. The text block begins on the next line. A text block ends at the first character of the closing """". Here is a simple example:

```
"""
Text blocks make
multiple lines easy.
"""


```

This example creates a string in which the line "Text blocks make" is separated from "multiple lines easy." by a newline. It is not necessary to use the \n escape sequence to obtain the newline. Thus, the text block automatically preserves the newlines in the text. Again, the text block begins after the newline *following* the opening delimiter and ends at the start of the closing delimiter. Therefore, the newline after the second line is also preserved.

It is important to emphasize that even though a text block uses the """" delimiter, it is of type **String**. Thus, the preceding text block could be assigned to a **String** variable, as shown here:

```
String str = """
Text blocks make
multiple lines easy.
""";
```

When **str** is output using this statement:

```
System.out.println(str);
```

the following is displayed:

```
Text blocks make
multiple lines easy.
```

Notice something else about this example. Because the last line ends with a newline, that newline will also be in the resulting string. If you don't want a trailing newline, then put the closing delimiter at the end of the last line, like this:

```
String str = """
Text blocks make
multiple lines easy."""; // now, no newline at the end
```

## Understanding Leading Whitespace

In the preceding example, the text in the block was placed flush left. However, this is not required. You can have leading whitespace in a text block. There are two primary reasons that you might want leading whitespace. First, it will enable the text to be better aligned with the indentation level of the code around it. Second, it supports one or more levels of indentation within the text block itself.

In general, leading whitespace in a text block is automatically removed. However, the number of leading whitespaces to remove from each line is determined by the number of leading whitespaces in the line with the least indentation. For example, if all lines are flush

left, then no whitespace is removed. If all lines are indented two spaces, then two spaces are removed from each line. However, if one line is indented two spaces, the next four spaces, and the third six spaces, then only two spaces are removed from the start of each line. This removes unwanted leading space while preserving the indentation of text within the block. This mechanism is illustrated by the following program:

```
// Demonstrate indentation in a text block.

class TextBlockDemo {

    public static void main(String[] args) {
        String str = """
            Text blocks support strings that
            span two or more lines and preserve
                indentation. They reduce the
                    tedium associated with the
                        entry of long or complicated
                            strings into a program.
            """;

        System.out.println(str);
    }
}
```

This program produces the following output:

```
Text blocks support strings that
span two or more lines and preserve
    indentation. They reduce the
        tedium associated with the
            entry of long or complicated
                strings into a program.
```

As you can see, leading whitespace has been removed up to, but not beyond, the level of the leftmost lines. Thus, the text block can be indented in the program to better fit the indentation level of the code, with the leading whitespace removed when the string is created. However, any whitespace after the indentation level of the block is preserved.

One other key point: The closing """" participates in determining the amount of whitespace to remove because it, too, can set the indentation level. Thus, if the closing delimiter is flush left, no whitespace is removed. Otherwise, whitespace is removed up to the first text character or when the closing delimiter is encountered. For example, consider this sequence:

```
String str = """
    A
        B
    C
"""; // this will determine the start of indent
```

```

String str2 = """
    A
        B
    C
"""; // this has no effect

String str3 = """
    A
        B
    C
"""; // this removes whitespace up to the """

System.out.print(str);
System.out.print(str2);
System.out.print(str3);

```

This sequence displays the following:

```

A
    B
C
A
    B
C
    A
        B
    C

```

Pay special attention to the placement of the closing delimiter for **str2**. Because the number of preceding spaces for the lines containing A and C are fewer than that preceding the "'''", the closing delimiter has no effect on the number of spaces removed.

## Use Double Quotes in a Text Block

Another major advantage to text blocks is the ability to use double quotes without the need for the \" escape sequence. In a text block, double quotes are treated like any other character. For example, consider the following program:

```

// Use double quotes in a text block.

class TextBlockDemo2 {

    public static void main(String[] args) {

        String str = """
                    A text block can use double quotes without
                    the need for escape sequences. For example:

                    He said, "The cat is on the roof."
                    She asked, "How did it get up there?"
                    """;

        System.out.println(str);
    }
}

```

The output is shown here:

```
A text block can use double quotes without  
the need for escape sequences. For example:
```

```
He said, "The cat is on the roof."  
She asked, "How did it get up there?"
```

As you can see, there was no need to use the \" escape sequence. Furthermore, because double quotes are treated as “normal” characters, there is also no need for them to be balanced within a text block. For example

```
"""  
"xyz"  
"""
```

is perfectly acceptable. Just remember that three double quotes as a unit defines the text block delimiter.

## Escape Sequences in Text Blocks

The escape sequences, such as \n or \t, can be used in a text block. However, because double quotes, newlines, and tabs can be entered directly, they will not often be needed. That said, with the addition of text blocks, two new escape sequences were added to Java. The first is \s, which specifies a space. Thus, it can be used to indicate trailing spaces. The second is \endofline. Because the \ must be followed by a line terminator, it must be used only at the end of the line. In a text block, the \ prevents a newline character from being included at the end of its line. Thus, the \ is essentially a line continuation indicator. For example,

```
String str = """  
    one \  
    two  
    three \  
    four  
    """;  
System.out.println(str);
```

Because the newline is suppressed after one and three, the output will be as shown here:

```
one two  
three four
```

It is important to point out that \endofline can only be used in a text block. It cannot be used to continue a traditional string literal, for example.

One final point before leaving the topic of text blocks: Because text blocks provide a better, easier way to enter many types of strings into your program, it is anticipated that they will be widely used by the Java programmer community. It is likely that you will begin to encounter them in code that you work on, or use them in code that you create. Just remember, your JDK release must be at least 15 or later.

## Records

Beginning with JDK 16, Java supports a special-purpose class called a *record*. A record is designed to provide an efficient, easy-to-use way to hold a group of values. For example, you might use a record to hold a set of coordinates; bank account numbers and balances; the length, width, and height of a shipping container; and so on. Because it holds a group of values, a record is commonly referred to as an *aggregate* type. However, the record is more than simply a means of grouping data, because records also have some of the capabilities of a class. In addition, a record has unique features that simplify its declaration and streamline access to its values. As a result, records make it much easier to work with groups of data. Records are supported by the new context-sensitive keyword **record**.

One of the central motivations for records is the reduction of the effort required to create a class whose primary purpose is to organize two or more values into a single unit. Although it has always been possible to use **class** for this purpose, doing so can entail writing a number of lines of code for constructors, getter methods, and possibly (depending on use) overriding one or more of the methods inherited from **Object**. As you will see, by creating a data aggregate by using **record**, these elements are handled automatically for you, greatly simplifying your code. Another reason for the addition of records is to enable a program to clearly indicate that the intended purpose of a class is to hold a grouping of data, rather than act as a full-featured class. Because of these advantages, records are a much welcomed addition to Java.

## Record Basics

As stated, a record is a narrowly focused, specialized class. It is declared by use of the **record** context-sensitive keyword. As such, **record** is a keyword only in the context of a **record** declaration. Otherwise, it is treated as a user-defined identifier with no special meaning. Thus, the addition of **record** does not impact or break existing code.

The general form of a basic **record** declaration is shown here:

```
record recordName(component-list) {
    // optional body statements
}
```

As the general form shows, a **record** declaration has significant differences from a **class** declaration. First, notice that the **record name** is immediately followed by a comma-separated list of parameter declarations called a *component list*. This list defines the data that the record will hold. Second, notice that the **body** is optional. This is made possible because the compiler will automatically provide the elements necessary to store the data; construct a record; create **getter methods** to access the data; and override **toString()**, **equals()**, and **hashCode()** inherited from **Object**. As a result, for many uses of a record, no body is required because the **record** declaration itself fully defines the record.

Here is an example of a simple **record** declaration:

```
record Employee(String name, int idNum) {}
```

Here, the record name is **Employee** and it has two components: the string **name** and the integer **idNum**. It specifies no statements in its body, so its body is empty. As the names imply, such a record could be used to store the name and ID number of an employee.

Given the **Employee** declaration just shown, a number of elements are automatically created. First, private final fields for **name** and **idNum** are declared as type **String** and **int**, respectively. Second, public read-only accessor methods (getter methods) that have the same names and types as the record components **name** and **idNum** are provided. Therefore, these getter methods are called **name()** and **idNum()**. In general, each record component will have a corresponding private final field and a read-only public getter method automatically created by the compiler.

Another element created automatically by the compiler will be the record's *canonical constructor*. This constructor has a parameter list that contains the same elements, in the same order, as the component list in the record declaration. The values passed to the constructor are automatically assigned to the corresponding fields in the record. In a **record**, the canonical constructor takes the place of the default constructor used by a **class**.

A **record** is instantiated by use of **new**, just the way you create an instance of a **class**. For example, this creates a new **Employee** object, with the name "Doe, John" and ID number 1047:

```
Employee emp = new Employee("Doe, John", 1047);
```

After this declaration executes, the private fields **name** and **idNum** for **emp** will contain the values "Doe, John" and 1047, respectively. Therefore, you can use the following statement to display the information associated with **emp**:

```
System.out.println("The employee ID for " + emp.name() + " is " +
    emp.idNum());
```

The resulting output is shown here:

```
The employee ID for Doe, John is 1047
```

A key point about a record is that its data is held in private final fields and only getter methods are provided. Thus, the data that a record holds is immutable. In other words, once you construct a record, its contents cannot be changed. However, if a record holds a reference to some object, you can make a change to that object, but you cannot change to what object the reference in the record refers. Thus, in Java terms, records are said to be *shallowly immutable*.

The following program puts the preceding discussion into action. Records are often used as elements in a list. For example, a business might maintain a list of **Employee** records to store an employee's name along with his or her corresponding ID number. The following program shows a simple example of such usage. It creates a small array of **Employee** records. It then cycles through the array, displaying the contents of each record.

```
// A simple Record example.

// Declare an employee record. This automatically creates a
// record class with private, final fields called name and idNum,
// and with read-only accessors called name() and idNum().
record Employee(String name, int idNum) {}

class RecordDemo {
    public static void main(String[] args) {
        // Create an array of Employee records.
        Employee[] empList = new Employee[4];
```

```
// Create a list of employees that uses the Employee record.  
// Notice how each record is constructed. The arguments  
// are automatically assigned to the name and idNum fields in  
// record that is being created.  
empList[0] = new Employee("Doe, John", 1047);  
empList[1] = new Employee("Jones, Robert", 1048);  
empList[2] = new Employee("Smith, Rachel", 1049);  
empList[3] = new Employee("Martin, Dave", 1050);  
  
// Use the record accessors to display names and IDs.  
for(Employee e: empList)  
    System.out.println("The employee ID for " + e.name() + " is " +  
                      e.idNum());  
}  
}
```

The output is shown here:

```
The employee ID for Doe, John is 1047  
The employee ID for Jones, Robert is 1048  
The employee ID for Smith, Rachel is 1049  
The employee ID for Martin, Dave is 1050
```

Before continuing, it is important to mention some key points related to records. First, a **record** cannot inherit another class. However, all records implicitly inherit **java.lang.Record**, which specifies abstract overrides of the **equals()**, **hashCode()**, and **toString()** methods declared by **Object**. Implicit implementations of these methods are automatically created, based on the record declaration. A **record** type cannot be extended. Thus, all **record** declarations are considered final. Although a **record** cannot extend another class, it can implement one or more interfaces. With the exception of **equals**, you cannot use the names of methods defined by **Object** as names for a **record**'s components. Aside from the fields associated with a **record**'s components, any other fields must be **static**. Finally, a **record** can be generic.

## Create Record Constructors

Although you will often find that the **automatically supplied canonical constructor** is precisely what you want, you can also declare one or more of your own constructors. You can also define your own implementation of the canonical constructor. You might want to declare a **record** constructor for a number of reasons. For example, the constructor could check that a value is **within a required range**, verify that a value is in the **proper format**, ensure that an object implements optional functionality, or confirm that an argument is not **null**. For a **record**, there are two general types of constructors that you can explicitly create: canonical and non-canonical, and there are some differences between the two. The creation of each type is examined here, beginning with defining your own implementation of the canonical constructor.

## Declare a Canonical Constructor

Although the canonical constructor has a specific, predefined form, there are two ways that you can code your own implementation. First, you can explicitly declare the full form of the canonical constructor. Second, you can use what is called a *compact record constructor*. Each approach is examined here, beginning with the full form.

To define your own implementation of a canonical constructor, simply do so as you would with any other constructor, specifying the record's name and its parameter list. It is important to emphasize that for the canonical constructor, the types and parameter names must be the same as those specified by the **record** declaration. This is because the parameter names are linked to the automatically created fields and accessor methods defined by the **record** declaration. Thus, they must agree in both type and name. Furthermore, each component must be fully initialized upon completion of the constructor. The following restrictions also apply: the constructor must be at least as accessible as its **record** declaration. Thus, if the access modifier for the **record** is **public**, the constructor must also be specified **public**. A constructor cannot be generic, and it cannot include a **throws** clause. It also cannot invoke another constructor defined for the record.

Here is an example of the **Employee** record that explicitly defines the canonical constructor. It uses this constructor to remove any leading or trailing whitespace from a name. This ensures that names are stored in a consistent manner.

```
record Employee(String name, int idNum) {  
  
    // Use a canonical constructor to remove any leading and trailing spaces  
    // that might be in the name string. This ensures that names are stored  
    // in a consistent manner.  
    public Employee(String name, int idNum) {  
        // Remove any leading and trailing spaces.  
        this.name = name.trim();  
        this.idNum = idNum;  
    }  
}
```

In the constructor, leading and/or trailing whitespace is removed by a call to **trim()**. Defined by the **String** class, **trim()** deletes all leading and trailing whitespace from a string and returns the result. (If there are no leading or trailing spaces, the original string is returned unaltered.) The resulting string is assigned to the field **this.name**. Thus, no **Employee** record **name** will contain leading or trailing spaces. Next, the value of **idNum** is assigned to **this.idNum**. Because the identifiers **name** and **idNum** are the same for both fields corresponding to the **Employee** components and for the names used by the canonical constructor's parameters, the field names must be qualified by **this**.

Although there is certainly nothing wrong with creating a canonical constructor as just shown, there is often an easier way: through the use of a *compact constructor*. A compact record constructor is declared by specifying the name of the record, but without parameters. The compact constructor implicitly has parameters that are the same as the record's components, and its components are automatically assigned the values of the arguments passed to the constructor. Within the compact constructor you can, however, alter one or more of the arguments prior to their value being assigned to the record.

The following example converts the previous canonical constructor into its compact form:

```
// Use a compact canonical constructor to remove any leading and
// trailing spaces from the name string.
public Employee {
    // Remove any leading and trailing spaces.
    name = name.trim();
}
```

Here, the result of `trim()` is called on the `name` parameter (which is implicitly declared by the compact constructor) and the result is assigned back to the `name` parameter. At the end of the compact constructor, the value of `name` is automatically assigned to its corresponding field. The value of the implicit `idNum` parameter is also assigned to its corresponding field at the end of the constructor. Because the parameters are implicitly assigned to their corresponding fields when the constructor ends, there is no need to initialize the fields explicitly. Moreover, it would not be legal to do so.

Here is a reworked version of the previous program that demonstrates the compact canonical constructor:

```
// Use a compact record constructor.

// Declare an employee record.
record Employee(String name, int idNum) {

    // Use a compact canonical constructor to remove any leading and
    // trailing spaces from the name string.
    public Employee {
        // Remove any leading and trailing spaces.
        name = name.trim();
    }
}

class RecordDemo2 {
    public static void main(String[] args) {
        Employee[] empList = new Employee[4];

        // Here, the name has no leading or trailing spaces.
        empList[0] = new Employee("Doe, John", 1047);

        // The next three names have leading and/or trailing spaces.
        empList[1] = new Employee(" Jones, Robert", 1048);
        empList[2] = new Employee("Smith, Rachel ", 1049);
        empList[3] = new Employee(" Martin, Dave ", 1050);

        // Use the record accessors to display names and IDs.
        // Notice that all leading and/or trailing spaces have been
        // removed from the name component by the constructor.
        for(Employee e: empList)
            System.out.println("The employee ID for " + e.name() + " is " +
                               e.idNum());
    }
}
```

The output is shown here:

```
The employee ID for Doe, John is 1047
The employee ID for Jones, Robert is 1048
The employee ID for Smith, Rachel is 1049
The employee ID for Martin, Dave is 1050
```

As you can see, the names have been standardized with leading and trailing spaces removed. To prove to yourself that the call to `trim()` is necessary to achieve this result, simply remove the compact constructor, recompile, and run the program. The leading and trailing spaces will still be in the names.

## Declare a Non-canonical Constructor

Although the canonical constructor will often be sufficient, you can declare other constructors. The key requirement is that any non-canonical constructor must first call another constructor in the record via `this`. The constructor invoked will often be the canonical constructor. Doing this ultimately ensures that all fields are assigned. Declaring a non-canonical constructor enables you to create special-case records. For example, you might use such a constructor to create a record in which one or more of the components is given a default placeholder value.

The following program declares a non-canonical constructor for `Employee` that initializes the name to a known value, but gives the `idNum` field the special value `pendingID` (which is `-1`) to indicate an ID value is not available when the record is created:

```
// Use a non-canonical constructor in a record.

// Declare an employee record that explicitly declares both
// a canonical and non-canonical constructor.
record Employee(String name, int idNum) {

    // Use a static field in a record.
    static int pendingID = -1;

    // Use a compact canonical constructor to remove any leading and
    // trailing spaces from the name string.
    public Employee {
        // Remove any leading and trailing spaces.
        name = name.trim();
    }

    // This is a non-canonical constructor. Notice that it is
    // not passed an ID number. Instead, it passes pendingID to the
    // canonical constructor to create the record.
    public Employee(String name) {
        this(name, pendingID);
    }
}

class RecordDemo3 {
    public static void main(String[] args) {
        Employee[] empList = new Employee[4];
```

```
// Create a list of employees that uses the Employee record.  
empList[0] = new Employee("Doe, John", 1047);  
empList[1] = new Employee("Jones, Robert", 1048);  
empList[2] = new Employee("Smith, Rachel", 1049);  
  
// Here, the ID number is pending.  
empList[3] = new Employee("Martin, Dave");  
  
// Display names and IDs.  
for(Employee e: empList) {  
    System.out.print("The employee ID for " + e.name() + " is ");  
    if(e.idNum() == Employee.pendingID) System.out.println("Pending");  
    else System.out.println(e.idNum());  
}  
}
```

Pay special attention to the way that the record for Martin, Dave is created by use of the non-canonical constructor. That constructor passes the **name** argument to the canonical constructor, but specifies the value **pendingID** as the **idNum** value. This enables a placeholder record to be created without having to specify an ID number. One other point: Notice that the value **pendingID** is declared as a **static** field in **Employee**. As explained earlier, instance fields are not allowed in a **record** declaration, but a **static** field is legal.

Notice that this version of **Employee** declares both a canonical constructor and a non-canonical constructor. This is perfectly valid. A record can define as many different constructors as its needs, as long as all adhere to the rules defined for **record**.

It is important to emphasize that records are immutable. As it relates to this example, it means that when an ID value for Martin, Dave is obtained, the old record must be replaced by a new record that contains the ID number. It is not possible to alter the record to update the ID. The immutability of records is a primary attribute.

## Another Record Constructor Example

Before leaving the topic of record constructors, we will look at one more example. Because a **record** is used to aggregate data, a common use of a **record** constructor is to verify the validity or applicability of an argument. For example, before constructing the record, the constructor may need to determine if a value is out of range, in an improper format, or otherwise unsuitable for use. If an invalid condition is found, the constructor could create a default or error instance. However, often it would be better for the constructor to throw an exception. This way, the user of the **record** would immediately be aware of the error and could take steps to correct it.

In the preceding **Employee** record examples, names have been specified using the common convention of *lastname,firstname*, such as Doe, John. However, there was no mechanism to verify or enforce that this format was being used. The following version of the compact canonical constructor provides a limited check that the name has the format *lastname,firstname*. It does so by confirming that there is one and only one comma in the name and that there is at least one character (other than space) before and after the comma.

Although a far more thorough, careful verification would be needed by a real-world program, this minimal check is sufficient to serve as an example of the validation role a record constructor might play.

Here is a version of the **Employee** record in which the compact canonical constructor throws an exception if the **name** component does not meet the minimal criteria required for the *lastname,firstname* format:

```
// Use a compact canonical constructor to remove any leading
// and trailing spaces in the name component. Also perform
// a basic check that the required format of lastname, firstname
// is passed to the name parameter.
public Employee {
    // Remove any leading and trailing spaces.
    name = name.trim();

    // Perform a minimalist check that name follows the
    // lastname, firstname format.
    //
    // First, confirm that name contains only one comma.
    int i = name.indexOf(',') ; // look for comma separating names.
    int j = name.lastIndexOf(',') ;
    if(i != j) throw
        new IllegalArgumentException("Multiple commas found.");

    // Next, confirm that a comma is present after
    // at least one leading character, and that at least one
    // character follows the comma.
    if(i < 1 | name.length() == i+1) throw
        new IllegalArgumentException("Required format: last, first");
}
```

When using this constructor, the following statement is still correct:

```
empList[0] = new Employee("Doe, John", 1047);
```

However, the following three are invalid and will result in an exception:

```
// No comma between last and first name.
empList[1] = new Employee("Jones Robert", 1048);

// Extra commas.
empList[1] = new Employee("Jones, ,Robert", 1048);

// Missing last name.
empList[1] = new Employee(", Robert", 1048);
```

As an aside, you might find it interesting to think of ways that you can improve the ability of the constructor to verify that the name uses the proper format. For example, you might want to explore an approach that uses a regular expression. (See Chapter 31.)

## Create Record Getter Methods

Although it is seldom necessary, it is possible to create your own implementation of a getter method. When you declare the getter, the implicit version is no longer supplied. One possible reason you might want to declare your own getter is to throw an exception if some condition is not met. For example, if a record holds a filename and a URL, the getter for the filename might throw a **FileNotFoundException** if the file is not present at the URL. There is a very important requirement, however, that applies to creating your getters: they must adhere to the principle that a record is immutable. Thus, a getter that returns an altered value is semantically questionable (and should be avoided) even though such code would be syntactically correct.

If you do declare a getter implementation, there are a number of rules that apply. A getter must have the same return type and name as the component that it obtains. It must also be explicitly declared public. (Thus, default accessibility is not sufficient for a getter declaration in a **record**.) No **throws** clause is allowed in a getter declaration. Finally, a getter must be non-generic and non-static.

A better alternative to overriding a getter in cases in which you want to obtain a modified value of a component is to declare a separate method with its own name. For example, assuming the **Employee** record, you might want to obtain only the last name from the **name** component. Using the standard getter to do this would entail modifying the value obtained from the component. Doing this is a bad idea because it would violate the immutability aspect of the record. However, you could declare another method, called **lastName()**, that returns only the last name. The following program demonstrates this approach. It also uses the format-checking constructor from the previous section to ensure that names are stored as *lastname,firstname*.

```
// Use an instance method in a record.

// This version of Employee provides a method called lastName()
// that returns only the last name of the name component.
// It also includes the version of the compact constructor that
// checks for the conventional lastname, firstname format.
record Employee(String name, int idNum) {

    // Use a compact canonical constructor to remove any leading
    // and trailing spaces in the name component. Also perform
    // a basic check that the required format of lastname, firstname.
    // is passed to the name parameter.
    public Employee {
        // Remove any leading and trailing spaces.
        name = name.trim();

        // Perform a minimalist check that name follows the
        // lastname, firstname format.
        //
        // First, confirm that name contains only one comma.
        int i = name.indexOf(',') ; // look for comma separating names.
        int j = name.lastIndexOf(',');
        if(i != j) throw
            new IllegalArgumentException("Multiple commas found.");
    }
}
```

```
// Next, confirm that a comma is present after
// at least one leading character, and that at least one
// character follows the comma.
if(i < 1 | name.length() == i+1) throw
    new IllegalArgumentException("Required format: last, first");
}

// An instance method that returns only the last name
// without the first name.
String lastName() {
    return name.substring(0, name.trim().indexOf(',') );
}
}

class RecordDemo4 {
    public static void main(String[] args) {
        Employee emp = new Employee("Jones, Robert", 1048);

        // Display the name component unmodified.
        System.out.println("Employee full name is " + emp.name());

        // Display only last name.
        System.out.println("Employee last name is " + emp.lastName());
    }
}
```

The output is shown here:

```
Employee full name is Jones, Robert
Employee last name is Jones
```

As the output shows, the implicit getter for the **name** component returns the name unaltered. The instance method **lastName()** obtains only the last name. With this approach, the immutable attribute of the **Employee** record is preserved, while still providing a convenient means of obtaining the last name.

## Pattern Matching with instanceof

The traditional form of the **instanceof** operator was introduced in Chapter 13. As explained there, **instanceof** evaluates to **true** if an object is of a specified type, or can be cast to that type. Beginning with JDK 16, a second form of the **instanceof** operator has been added to Java that supports the **new pattern matching** feature. In general terms, *pattern matching* defines a mechanism that determines if a **value** fits a **general form**. As it relates to **instanceof**, pattern matching is used to test the type of a value (which must be a reference type) against a specified type. This kind of pattern is called a **type pattern**. If the pattern matches, a **pattern variable** will receive a reference to the object matched by the pattern.

The pattern matching form of **instanceof** is shown here:

*objref* instanceof *type pattern-var*

If **instanceof** succeeds, *pattern-var* will be created and contain a reference to the object that matches the pattern. If it fails, *pattern-var* is never created. This form of **instanceof** succeeds

if the object referred to by *objref* can be cast to *type* and the static type of *objref* is not a subtype of *type*.

For example, the following fragment creates a **Number** reference called **myOb** that refers to an **Integer** object. (Recall that **Number** is a superclass of all numeric primitive-type wrappers.) It then uses the **instanceof** operator to confirm that the object referred to by **myOb** is an **Integer**, which it will be in this example. This results in an object called **iObj** of type **Integer** being instantiated that contains the matched value.

```
Number myOb = Integer.valueOf(9);

// Use the pattern matching version of instanceof.
if(myOb instanceof Integer iObj) {
    // iObj is known and in scope here.
    System.out.println("iObj refers to an integer: " + iObj);
}
// iObj does not exist here
```

As the comments indicate, **iObj** is known only within the scope of the **if** clause. It is not known outside of the **if**. It also would not be known within an **else** clause, should one have been included. It is crucial to understand that the pattern variable **iObj** is created only if the pattern matching succeeds.

The primary advantage of the pattern matching form of **instanceof** is that it reduces the amount of code that was typically needed by the traditional form of **instanceof**. For example, consider this functionally equivalent version of the preceding example that uses the traditional approach:

```
// Use a traditional instanceof.
if(myOb instanceof Integer) {
    // Use an explicit cast to obtain iObj.
    Integer iObj = (Integer) myOb;
    System.out.println("iObj refers to an integer: " + iObj);
}
```

With the traditional form, a separate declaration statement and explicit cast are required to create the **iObj** variable. The pattern matching form of **instanceof** streamlines the process.

## Pattern Variables in a Logical AND Expression

An **instanceof** can be used in a logical AND expression. However, you need to remember that the pattern variable is only in scope after it has been created. For example, the following **if** succeeds only when **myOb** refers to an **Integer** and its value is nonnegative. Pay special attention to the expression in the **if**:

```
if((myOb instanceof Integer iObj) && (iObj >= 0)) { // is OK
    // myOb is both an Integer and nonnegative.
    // ...
}
```

The **iObj** pattern variable is created only if the left side of the **&&** (the part that contains the **instanceof** operator) is true. Notice that **iObj** is also used by the right side. This is possible

because the short-circuit form of the AND logical operator is used, and the right side is evaluated *only if* the left succeeds. Thus, if the right-side operand is evaluated, **iObj** will be in scope. However, if you tried to write the preceding statement using the **&** operator like this:

```
if((myOb instanceof Integer iObj) & (iObj >= 0)) { // Error!
    // myOb is both an Integer and nonnegative.
    // ...
}
```

a compilation error would occur because **iObj** will not be in scope if the left side fails. Recall that the **&** operator causes both sides of the expression to be evaluated, but **iObj** is only in scope if the left side is true. This error is caught by the compiler. A related situation occurs with this fragment:

```
int count = 10;
if((count < 100) && myOb instanceof Integer iObj) { // is OK
    // myOb is both an Integer and nonnegative, and count is less than 100.

    iObj = count;
    // ...
}
```

This fragment compiles because the **if** block will execute only when both sides of the **&&** are true. Thus, the use of **iObj** in the **if** block is valid. However, a compilation error will result if you tried to use the **&** rather than the **&&**, as shown here:

```
if((count < 100) & myOb instanceof Integer iObj) { // Error!
```

In this case, the compiler cannot know whether or not **iObj** will be in scope in the **if** block because the right side of the **&** will not necessarily be evaluated.

One other point: A logical expression cannot introduce the same pattern variable more than once. For example, in a logical AND, it is an error if both operands create the same pattern variable.

## Pattern Matching in Other Statements

Although a frequent use of the pattern matching form of **instanceof** is in an **if** statement, it is by no means limited to that use. It can also be employed in the conditional portion of the loop statements. As an example, imagine that you are processing a collection of objects, perhaps contained in an array. Furthermore, at the start of the array are several strings, and you want to process those strings, but not any of the remaining objects in the list. The following sequence accomplishes this task with a **for** loop in which the condition uses **instanceof** to confirm that an object in the array is a **String** and to obtain that string for processing within the loop. Thus, pattern matching is used to control the execution of a **for** loop and to obtain the next value for processing.

```
Object[] someObjs = {
    new String("Alpha"),
    new String("Beta"),
    new String("Omega"),
```

```

        Integer.valueOf(10)
    };

int i;

// This loop iterates until an element is not a String, or the end
// of the array is reached.
for(i = 0; (someObjs[i] instanceof String str) && (i < someObjs.length); i++) {
    System.out.println("Processing " + str);
    // ...
}

System.out.println("The first " + i + " entries in the list are strings.");

```

The output from this fragment is shown here:

```

Processing Alpha
Processing Beta
Processing Omega
The first 3 entries in the list are strings.

```

The pattern matching form of **instanceof** can also be useful in a **while** loop. For example, here is the preceding **for** loop, recoded as a **while**:

```

i = 0;
while((someObjs[i] instanceof String str) && (i < someObjs.length)) {
    System.out.println("Processing " + str);
    i++;
}

```

Although it is technically possible to use the pattern matching **instanceof** in the conditional portion of a **do** loop, such use is severely limited because the pattern variable cannot be used in body of the loop because it will not be in scope until the **instanceof** operator is executed.

## Sealed Classes and Interfaces

Beginning with JDK 17, it is possible to declare a class that can be inherited by only specific subclasses. Such a class is called *sealed*. Prior to the advent of sealed classes, inheritance was an “all or nothing” situation. A class could either be extended by any subclass or marked as **final**, which prevented its inheritance entirely. Sealed classes fall between these two extremes because they enable you to specify precisely what subclasses a superclass will allow. In similar fashion, it is also possible to declare a sealed interface in which you specify only those classes that implement the interface and/or those interfaces that extend the sealed interface. Together, sealed classes and interfaces give you significantly greater control over inheritance, which can be especially important when designing class libraries.

## Sealed Classes

To declare a sealed class, precede the declaration with **sealed**. Then, after the class name, include a **permits** clause that specifies the allowed subclasses. Both **sealed** and **permits** are context-sensitive keywords that have special meaning only in a class or interface declaration. Outside of a class or interface declaration, **sealed** and **permits** are unrestricted and have no special meaning. Here is a simple example of a sealed class:

```
public sealed class MySealedClass permits Alpha, Beta {
    // ...
}
```

Here, the sealed class is called **MySealedClass**. It allows only two subclasses: **Alpha** and **Beta**. If any other class attempts to inherit **MySealedClass**, a compile-time error will occur.

Here are **Alpha** and **Beta**, the two subclasses of **MySealedClass**:

```
public final class Alpha extends MySealedClass {
    // ...
}

public final class Beta extends MySealedClass {
    // ...
}
```

Notice that each is specified as **final**. In general, a subclass of a sealed class must be declared as either **final**, **sealed**, or **non-sealed**. Let's look at each option. First, in this example, each subclass is declared **final**. This means that the only subclasses of **MySealedClass** are **Alpha** and **Beta**, and no subclasses of either of those can be created. Therefore, the inheritance chain ends with **Alpha** and **Beta**.

To indicate that a subclass is itself sealed, it must be declared **sealed** and its permitted subclasses must be specified. For example, this version of **Alpha** permits **Gamma**:

```
public sealed class Alpha extends MySealedClass permits Gamma {
    // ...
}
```

Of course, the class **Gamma** must then itself be declared either **sealed**, **final**, or **non-sealed**.

At first it might seem a bit surprising, but you can unseal a subclass of a sealed class by declaring it **non-sealed**. This context-sensitive keyword was added by JDK 17. It unlocks the subclass, enabling it to be inherited by any other class. For example, **Beta** could be coded like this:

```
public non-sealed class Beta extends MySealedClass {
    // ...
}
```

Now, any class may inherit **Beta**. However, the only direct subclasses of **MySealedClass** remain **Alpha** and **Beta**. A primary reason for **non-sealed** is to enable a superclass to specify a limited set of direct subclasses that provide a baseline of well-defined functionality but allow those subclasses to be freely extended.

If a class is specified in a **permits** clause for a sealed class, then that class *must* directly extend the sealed class. Otherwise, a compile-time error will result. Thus, a sealed class and its subclasses define a mutually dependent logical unit. Additionally, it is illegal to declare a class that does not extend a sealed class as **non-sealed**.

A key requirement of a sealed class is that every subclass that it permits must be accessible. Furthermore, if a sealed class is contained in a named module, then each subclass must also be in the same named module. In this case, a subclass can be in a different package from the sealed class. If the sealed class is in the unnamed module, then the sealed class and all permitted subclasses must be in the same package.

In the preceding discussion, the superclass **MySealedClass** and its subclasses **Alpha** and **Beta** would have been stored in separate files because they are all public classes. However, it is also possible for a sealed class and its subclasses to be stored in a single file (formally, a compilation unit) as long as the subclasses have default package access. In cases such as this, no **permits** clause is required for a sealed class. For example, here all three classes are in the same file:

```
// Because this is all in one file, MySealedClass does not require
// a permits clause.
public sealed class MySealedClass {
    // ...
}

final class Alpha extends MySealedClass {
    // ...
}

final class Beta extends MySealedClass {
    // ...
}
```

One last point: An abstract class can also be sealed. There is no restriction in this regard.

## Sealed Interfaces

A sealed interface is declared in the same way as a sealed class, by the use of **sealed**. A sealed interface uses its **permits** clause to specify the classes allowed to implement it and/or the interfaces allowed to extend it. Thus, a class that is not part of the **permits** clause cannot implement a sealed interface, and an interface not included in the **permits** clause cannot extend it.

Here is a simple example of a sealed interface that permits only the classes **Alpha** and **Beta** to implement it:

```
public sealed interface MySealedIF permits Alpha, Beta {
    void myMeth();
}
```

A class that implements a sealed interface must itself be specified as either **final**, **sealed**, or **non-sealed**. For example, here **Alpha** is marked **non-sealed** and **Beta** is specified as **final**:

```
public non-sealed class Alpha implements MySealedIF {
    public void myMeth() { System.out.println("In Alpha's myMeth()."); }
    // ...
}

public final class Beta implements MySealedIF {
    public void myMeth() { System.out.println("Inside Beta's myMeth()."); }
    // ...
}
```

Here is a key point: Any class specified in a sealed interface's **permits** clause *must* implement the interface. Thus, a sealed interface and its implementing classes form a logical unit.

A sealed interface can also specify which other interfaces can extend the sealed interface. For example, here, **MySealedIF** specifies that **MyIF** is permitted to extend it:

```
// Notice that MyIF is added to the permits clause.
public sealed interface MySealedIF permits Alpha, Beta, MyIF {
    void myMeth();
}
```

Because **MyIF** is part of the **MySealedIF permits** clause, it must be marked as either **non-sealed** or **sealed** and it must extend **MySealedIF**. For example,

```
public non-sealed interface MyIF extends MySealedIF {
    // ...
}
```

As you might expect, it is possible for a class to inherit a sealed class *and* implement a sealed interface. For example, here **Alpha** inherits **MySealedClass** and implements **MySealedIF**:

```
public non-sealed class Alpha extends MySealedClass implements MySealedIF {
    public void myMeth() { System.out.println("In Alpha's myMeth()."); }
    // ...
}
```

In the preceding examples, each class and interface are declared **public**. Thus, each is in its own file. However, as is the case with sealed classes, it is also possible for a sealed interface and its implementing classes (and extending interfaces) to be stored in a single file as long as the classes and interfaces have default package access. In cases such as this, no **permits** clause is required for a sealed interface. For example, here **MySealedIF** does not include a **permits** clause because **Alpha** and **Beta** are declared in the same file in the unnamed module:

```
public sealed interface MySealedIF {
    void myMeth();
}
```

```
non-sealed class Alpha extends MySealedClass implements MySealedIF {  
    public void myMeth() { System.out.println("In Alpha's myMeth() ."); }  
    // ...  
}  
  
final class Beta extends MySealedClass implements MySealedIF {  
    public void myMeth() { System.out.println("In Beta's myMeth() ."); }  
    // ...  
}
```

One final point: Sealed classes and interfaces are most applicable to developers of API libraries in which subclasses and subinterfaces must be strictly controlled.

## Future Directions

Beginning with JDK 12, Java releases may, and often do, include *preview features*. A preview feature is a new, fully developed enhancement to Java. However, a preview feature is *not yet* formally part of Java. Instead, a feature is previewed to allow programmers time to experiment with the feature and, if desired, communicate their thoughts and opinions prior to the feature being made permanent. This process enables a new feature to be improved or optimized based on actual developer use. As a result, a preview feature is *subject to change*. It can even be withdrawn. This means that a preview feature should not be used for code that you intend to publicly release. That said, it is expected that most preview features will ultimately become part of Java, possibly after a period of refinement. Preview features chart the course of Java's future direction.

JDK 17 includes one preview feature: Pattern Matching for **switch** (JEP 406). It adds pattern matching capabilities to **switch**. As described earlier in this chapter, pattern matching was first introduced by the enhancement of **instanceof** in JDK 16. Adding pattern matching to **switch** continues the process. Because this is a preview feature that is subject to change, it is not discussed further in this book.

Java releases may also include *incubator modules*, which preview a new API or tool that is undergoing development. Like a preview feature, an incubator feature is subject to change. Furthermore, an incubator feature can be removed in the future. Thus, there is no guarantee that an incubating module will formally become part of Java in the future. Incubator features give developers an opportunity to experiment with the API or tool, and possibly supply feedback. JDK 17 includes two incubator modules. The first is Foreign Function and Memory API (JEP 412). The second is Vector API (JEP 414).

It is important to emphasize that preview features and incubator modules can be introduced in any Java release. Therefore, you will want to watch for them in each new version of Java. They give you a chance to try a new enhancement before it potentially becomes a formal part of Java. Perhaps more importantly, preview features and incubator modules give you advance information on where Java's development is headed.