

CHAPTER 1



Introduction to Computing with Python

This book is about using Python for numerical computing. Python is a high-level, general-purpose interpreted programming language widely used in scientific computing and engineering. As a general-purpose language, Python was not specifically designed for numerical computing, but many of its characteristics make it well-suited for this task. First and foremost, Python is well known for its clean and easy-to-read code syntax. Good code readability improves maintainability, reduces bugs, and leads to better applications overall. It also enables rapid code development, since readability and expressiveness are essential in exploratory and interactive computing, where fast turnaround for testing various ideas and models is important.

In computational problem-solving, it is important to consider algorithms' performance and implementations. It is natural to strive for efficient high-performance code, and optimal performance is crucial for many computational problems. In such cases, it may be necessary to use a low-level program language, such as C or Fortran, to obtain the best performance out of the hardware that runs the code. However, it is not always the case that optimal runtime performance should be the highest priority. It is also important to consider the development time required to solve a problem in a programming language or an environment. While the best possible runtime performance can be achieved in a low-level programming language, working in a high-level language such as Python reduces the development time and often results in more flexible and extensible code.

These conflicting objectives present a trade-off between high performance but long development time and lower performance but shorter development time. Figure 1-1 shows a schematic visualization of this concept. When choosing a computational environment for solving a particular problem, it is important to consider this trade-off and to decide whether person-hours spent on the development or CPU-hours spent on running the computations are more valuable. It is worth noting that CPU-hours are cheap and getting even cheaper, but person-hours are expensive. Your own time is a very valuable resource. This makes a strong case for minimizing development time rather than the computation runtime by using a high-level programming language and environment such as Python and its scientific computing libraries.




Figure 1-1. The trade-off between low- and high-level programming languages. While a low-level language typically gives the best performance when a significant amount of development time is invested in the implementation of a solution to a problem, the development time required to obtain a first runnable code that solves the problem is typically shorter in a high-level language such as Python

A solution that partially avoids the trade-off between high- and low-level languages is to use a multilanguage model, where a high-level language is used to interface libraries and software packages written in low-level languages. In a high-level scientific computing environment, this type of interoperability with software packages written in low-level languages (e.g., Fortran, C, or C++) is an important requirement. Python excels at this type of integration, and as a result, Python has become a popular “glue language” used as an interface for setting up and controlling computations that use code written in low-level programming languages for time-consuming number crunching. This is an important reason why Python is a popular language for numerical computing. The multilanguage model enables rapid code development in a high-level language while retaining most of the performance of low-level languages.

Due to the multilanguage model, scientific and technical computing with Python involves much more than just the Python language itself. In fact, the Python language is only a piece of an entire ecosystem of software and solutions that provide a complete environment for scientific and technical computing. This ecosystem includes development tools and interactive programming environments, such as Spyder and IPython, which are designed particularly with scientific computing in mind. It also includes a vast collection of Python packages for scientific computing. This ecosystem of scientifically oriented libraries ranges from generic core libraries—such as NumPy, SciPy, and Matplotlib—to more specific libraries for problem domains. Another crucial layer in the scientific Python stack exists below the various Python modules. Many scientific Python libraries interface with low-level, high-performance scientific software packages, such as optimized LAPACK and BLAS libraries¹ for low-level vector, matrix, and linear algebra routines or other specialized libraries for specific computational tasks. These libraries are typically implemented in a compiled low-level language and can be highly optimized and efficient. Without the foundation that such libraries provide, scientific computing with Python would not be practical. Figure 1-2 is an overview of the various layers of the software stack for computing with Python.

¹For example, MKL, the Math Kernel Library from Intel at <https://software.intel.com/en-us/intel-mkl>; openBLAS at www.openblas.net; or ATLAS, the Automatically Tuned Linear Algebra Software at <http://math-atlas.sourceforge.net>

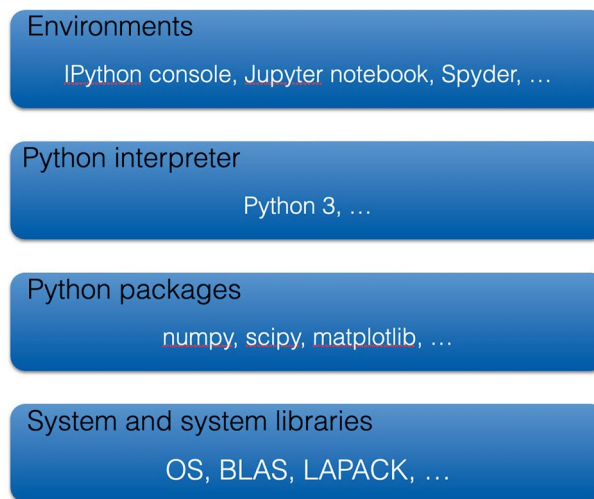


Figure 1-2. An overview of the components and layers in the scientific computing environment for Python, from a user's perspective from top to bottom. Users typically only interact with the top three layers, but the bottom layer constitutes a very important part of the software stack

■ **Tip** The SciPy organization (www.scipy.org) provides a centralized resource for information about the core packages in the scientific Python ecosystem, lists of additional specialized packages, and documentation and tutorials. It is a valuable resource when working with scientific and technical computing in Python. Another great resource is the Numeric and Scientific page on the official Python Wiki (<http://wiki.python.org/moin/NumericAndScientific>).

Besides the technical reasons why Python provides a good environment for computational work, it is also significant that it and its scientific computing libraries are **free and open source**. This eliminates economic constraints on when and how applications developed with the environment can be deployed and distributed by its users. Equally significant, it makes it possible for a dedicated user to obtain complete insight into how the language and the domain-specific packages are implemented and what methods are used. **For academic work where transparency and reproducibility are hallmarks, this is increasingly recognized as an important requirement of software used in research.** For commercial use, it provides freedom on how the environment is used and integrated into products and how such solutions are distributed to customers. All users benefit from the relief of not paying license fees, which may otherwise inhibit deployments on large computing environments, such as clusters and cloud computing platforms.

The social component of the scientific computing ecosystem for Python is another important aspect of its success. Vibrant user communities have emerged around the core packages and many domain-specific projects. Project-specific mailing lists, Stack Overflow groups, and issue trackers (e.g., on GitHub, www.github.com) are typically very active and provide forums for discussing problems and obtaining help, as well as a way of getting involved in developing these tools. The Python computing community also organizes yearly conferences and meet-ups at many venues around the world, such as the SciPy (<http://conference.scipy.org>) and PyData (<http://pydata.org>) conference series.

Environments for Computing with Python

Several different environments are suitable for working with Python for scientific and technical computing. This diversity has both advantages and disadvantages compared to a single endorsed environment that is common in proprietary computing products: diversity provides flexibility and dynamism that lends itself to specialization for particular uses, but on the other hand, it can also be confusing for new users, and it can be more complicated to set up a full productive environment. Here, I give an orientation of common environments for scientific computing so that their benefits can be weighed against each other and an informed decision can be reached regarding which one to use in different situations and for different purposes. The following are the three environments discussed in this chapter.

- The Python interpreter or the IPython console run code interactively. Together with a text editor for writing code, this provides a lightweight development environment.
- The Jupyter Notebook is a web application in which Python code can be written and executed through a web browser. This environment is great for numerical computing, analysis, and problem-solving because it allows us to collect the code, the output produced by the code, related technical documentation, and the analysis and interpretation all in one document.
- The Spyder Integrated Development Environment can write and interactively run Python code. An IDE like Spyder is a great tool for developing libraries and reusable Python modules.

These environments have justified uses, and it is largely a matter of personal preference for which one to use. However, I recommend exploring the Jupyter Notebook environment, because it is highly suitable for interactive and exploratory computing and data analysis, where data, code, documentation, and results are tightly connected. For the development of Python modules and packages, I recommend using the Spyder IDE because of its integration with code analysis tools and the Python debugger.

Python and the rest of the software stack required for scientific computing with Python can be installed and configured in many ways, and in general, the installation details also vary from system to system. The Appendix goes through one popular cross-platform method to install the tools and libraries required for this book.

Python

The Python programming language and the standard implementation of the Python interpreter are frequently updated and made available through new releases.² Currently, the active version of Python available for production use is the Python 3 series; this book requires Python 3.8 or greater. Note that at the time of writing, versions prior to Python 3.8 have already passed end-of-life, meaning they will no longer receive important bug fixes and security updates. Should you encounter any such legacy Python environment, it is therefore recommended that you upgrade the Python interpreter to a newer version.

Interpreter

The standard way to execute Python code is to run the program directly through the Python interpreter. On most systems, the Python interpreter is invoked using the `python` command. When a Python source file is passed as an argument to this command, the Python code in the file is executed.

²The Python language and the default Python interpreter are managed and maintained by the Python Software Foundation (www.python.org).

```
$ python hello.py
Hello from Python!
```

Here, the `hello.py` file contains a single line.

```
print("Hello from Python!")
```

To see which version of Python is installed, we can invoke the `python` command with the `--version` argument.

```
$ python --version
Python 3.11.4
```

It is common to have more than one version of Python installed on the same system. Each version of Python maintains its own set of libraries and provides its own interpreter command (so each Python environment can have different libraries installed). On many systems, specific versions of the Python interpreter are available through commands such as `python3.11`. It is also possible to set up *virtual* Python environments independent of the system-provided environments, which has many advantages. I strongly recommend becoming familiar with this way of working with Python. Appendix A describes setting up and working with these kinds of environments.

In addition to executing Python script files, a Python interpreter can be used as an interactive console (also known as a REPL (read-evaluate-print-loop)). Entering `python` at the command prompt (without any Python files as arguments) launches the Python interpreter in an interactive mode. When doing so, you are presented with a prompt.

```
$ python
Python 3.11.4 (main, Jul 5 2023, 08:41:25) [Clang 14.0.6 ] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

From here, Python code can be entered, and for each statement, the interpreter evaluates the code and prints the result to the screen. The Python interpreter itself already provides a very useful environment for interactively exploring Python code, especially since the release of Python 3.4, which includes basic facilities such as a command history and basic autocompletion.

IPython Console

Although the **interactive command-line interface** provided by the standard Python interpreter has been greatly improved in the Python interpreter itself, it is still, in certain aspects, rudimentary, and it does not provide a complete environment for interactive computing. IPython³ is an enhanced command-line REPL environment for Python, with additional interactive and exploratory computing features. For example, IPython provides improved command history browsing (also between sessions), an input and output caching system, improved auto-completion, more verbose and helpful exception tracebacks, and more. IPython is now much more than an enhanced Python command-line interface, which is explored in more detail later in this chapter and throughout the book. For instance, under the hood, IPython is a client-server application that separates the front end (user interface) from the back end (kernel) and executes the Python code. This allows multiple types of user interfaces to communicate and work with the same kernel, and a user-interface application can connect multiple kernels using IPython's framework for parallel computing.

³ See the IPython project web page, <http://ipython.org>, for more information and its official documentation.

Running the `ipython` command launches the IPython command prompt.

```
$ ipython
Python 3.11.4 (main, Jul 5 2023, 08:41:25) [Clang 14.0.6 ]
Type 'copyright', 'credits' or 'license' for more information
IPython 8.12.2 -- An enhanced Interactive Python. Type '?' for help
In [1]:
```

■ **Caution** Each IPython installation corresponds to a specific version of Python. If several versions of Python are available on your system, you may also have several versions of IPython. On many systems, IPython for Python 3 is invoked with the `ipython3` command, although the exact setup varies from system to system. Note that here, the “3” refers to the Python version, which differs from the version of IPython itself (at the time of writing it is 8.12.2).

The following sections briefly overview some of the IPython features that are most relevant to interactive computing. It is worth noting that IPython is used in many different contexts in scientific computing with Python, for example, as a kernel in the Jupyter Notebook application and in the Spyder IDE, which is covered in more detail later in this chapter. It is time well spent to get familiar with the tricks and techniques that IPython offers to improve your productivity when working with interactive computing.

Input and Output Caching

In the IPython console, the input prompt is denoted as `In [1]:` and the corresponding output is denoted as `Out [1]:`, where the numbers within the square brackets are incremented for each new input and output. These inputs and outputs are called *cells* in IPython. The input and the output of previous cells can later be accessed through the `In` and `Out` variables that IPython automatically creates. The `In` and `Out` variables are a list and a dictionary, respectively, that can be indexed with a cell number. For instance, consider the following IPython session.

```
In [1]: 3 * 3
Out[1]: 9
In [2]: In[1]
Out[2]: '3 * 3'
In [3]: Out[1]
Out[3]: 9
In [4]: In
Out[4]: ['', '3 * 3', 'In[1]', 'Out[1]', 'In']
In [5]: Out
Out[5]: {1: 9, 2: '3 * 3', 3: 9, 4: ['', '3 * 3', 'In[1]', 'Out[1]', 'In', 'Out']}
```

Here, the first input was `3 * 3`, and the result was 9, which later is available as `In[1]` and `Out[1]`. A single underscore `_` is a shorthand notation for referring to the most recent output, and a double underscore `__` refers to the output that preceded the most recent output. Input and output caching is often useful in interactive and exploratory computing since the result of a computation can be accessed even if it was not explicitly assigned to a variable.

Note that when a cell is executed, the value of the last statement in an input cell is, by default, displayed in the corresponding output cell unless the statement is an assignment or if the value is Python null value `None`. The output can be suppressed by ending the statement with a semicolon.

```

In [6]: 1 + 2
Out[6]: 3
In [7]: 1 + 2;    # output suppressed by the semicolon
In [8]: x = 1     # no output for assignments
In [9]: x = 2; x  # these are two statements. The value of 'x' is shown in
                  # the output
Out[9]: 2

```

Autocompletion and Object Introspection

In IPython, pressing the TAB key activates autocompletion, which displays a list of symbols (variables, functions, classes, etc.) with names that are valid completions of what has already been typed. The autocompletion in IPython is contextual, and it looks for matching variables and functions in the current namespace or among the attributes and methods of a class when invoked after the name of a class instance. For example, `os.<TAB>` produces a list of the variables, functions, and classes in the `os` module, and pressing TAB after typing `os.w` results in a list of symbols in the `os` module that starts with `w`.

```

In [10]: import os
In [11]: os.w<TAB>
os.wait  os.wait3  os.wait4  os.waitpid  os.walk  os.write  os.writev

```

This feature is called *object introspection*, a powerful tool for interactively exploring the properties of Python objects. Object introspection works on modules, classes, attributes, methods, functions, and arguments.

Documentation

Object introspection is convenient for exploring the API of a module, such as its member classes and functions, and together with the documentation strings or *docstrings* that are commonly provided in Python code, it provides a built-in dynamic reference manual for almost any Python module that is installed and can be imported. A Python object followed by a question mark displays the documentation string for the object. This is similar to the Python function `help`. An object can also be followed by two question marks, in this case, IPython tries to display more detailed documentation, including the Python source code, if available. For example, to display help for the `cos` function in the `math` library.

```

In [12]: import math
In [13]: math.cos?
Signature: math.cos(x, /)
Docstring: Return the cosine of x (measured in radians).
Type:      builtin_function_or_method

```

Docstrings can be specified for Python modules, functions, classes, and their attributes and methods. A well-documented module includes full API documentation in the code itself. From a developer's point of view, it is convenient to document a code together with the implementation. This encourages writing and maintaining documentation, and Python modules tend to be well-documented.

Interaction with the System Shell

IPython also provides extensions to the Python language that make interacting with the underlying system convenient. Anything that follows an exclamation mark is evaluated using the system shell (such as bash shell). For example, on a Unix-like system, such as Linux or macOS, listing files in the current directory can be done using the following.

```
In[14]: !ls
file1.py    file2.py    file3.py
```

In Microsoft Windows, the equivalent command would be `!dir`. This method for interacting with the operating system is a powerful feature that makes it easy to navigate the file system and use the IPython console as a system shell. The output generated by a command following an exclamation mark can easily be captured in a Python variable. For example, a file listing produced by `!ls` can be stored in a Python list using the following.

```
In[15]: files = !ls
In[16]: len(files)
3
In[17]: files
['file1.py', 'file2.py', 'file3.py']
```

Likewise, we can pass the values of Python variables to shell commands by prefixing the variable name with a `$` sign.

```
In[18]: file = "file1.py"
In[19]: !ls -l $file
-rw-r--r--  1 rob  staff 131 Oct 22 16:38 file1.py
```

This two-way communication with the IPython console and the system shell can be very convenient, for example, when processing data files.

IPython Extensions

IPython provides extension commands that are called *magic functions* in IPython terminology. These commands all start with one or two `%` signs.⁴ A single `%` sign is used for one-line commands, and two `%` signs are used for commands that operate on cells (multiple lines). For a complete list of available extension commands, type `%lsmagic`, and the documentation for each command can be obtained by typing the magic command followed by a question mark.

```
In[20]: %lsmagic?
Docstring: List currently available magic functions.
File:      /usr/local/lib/python3.6/site-packages/IPython/core/magics/basic.py
```

⁴When `%automagic` is activated (type `%automagic` at the IPython prompt to toggle this feature), the `%` sign that precedes the IPython commands can be omitted, unless there is a name conflict with a Python variable or function. However, for clarity, the `%` signs are explicitly shown here.

File System Navigation

In addition to the interaction with the system shell described in the previous section, IPython provides commands for navigating and exploring the file system. These commands are familiar to Unix shell users: `%ls` (list files), `%pwd` (return current working directory), `%cd` (change working directory), `%cp` (copy file), `%less` (show the content of a file in the pager), and `%writefile filename` (write content of a cell to the file `filename`). Note that autocomplete in IPython also works with the files in the current working directory, which makes IPython as convenient to explore the file system as the system shell. It is worth noting that these IPython commands are system-independent and can be used on both Unix-like operating systems and Windows.

Running Scripts from the IPython Console

The `%run` command is an important and useful extension, perhaps one of the most important features of the IPython console. This command can execute an external Python source code file within an interactive IPython session. Keeping a session active between multiple runs of a script makes it possible to explore the variables and functions defined in a script interactively after the execution of the script has finished. To demonstrate this functionality, consider a script file `fib.py` that contains the following code.

```
def fib(n):
    """
    Return a list of the first n Fibonacci numbers.
    """
    f0, f1 = 0, 1
    f = [1] * n
    for i in range(1, n):
        f[i] = f0 + f1
        f0, f1 = f1, f[i]
    return f

print(fib(10))
```

It defines a function that generates a sequence of n Fibonacci numbers and prints the result for $n = 10$ to the standard output. It can be run from the system terminal using the standard Python interpreter.

```
$ python fib.py
[1, 1, 2, 3, 5, 8, 13, 21, 34, 55]
```

It can also be run from an interactive IPython session, which produces the same output but also adds the symbols defined in the file to the local namespace so that the `fib` function is available in the interactive session after the `%run` command has been issued.

```
In [21]: %run fib.py
Out[22]: [1, 1, 2, 3, 5, 8, 13, 21, 34, 55]
In [23]: %who
fib
In [23]: fib(6)
Out[23]: [1, 1, 2, 3, 5, 8]
```

The preceding example also used the `%who` command, which lists all defined symbols (variables and functions).⁵ The `%whos` command is similar, but also gives more detailed information about the type and value of each symbol, when applicable.

Debugger

IPython includes a handy debugger mode, which can be invoked postmortem after a Python exception (error) has been raised. After the traceback of an unintercepted exception has been printed to the IPython console, it is possible to step directly into the Python debugger using the IPython command `%debug`. This possibility can eliminate the need to rerun the program from the beginning using the debugger or after employing the common debugging method of sprinkling print statements into the code. If the exception is unexpected and happens late in a time-consuming computation, this can be a big time-saver.

To see how the `%debug` command can be used, consider the following incorrect invocation of the `fib` function defined earlier. It is incorrect because a float is passed to the function while the function is implemented, assuming that the argument passed to it is an integer. On line 7 the code runs into a type error, and the Python interpreter raises an exception of `TypeError`. IPython catches the exception and prints a useful traceback of the call sequence on the console. If we are clueless about why the code on line 7 contains an error, entering the debugger by typing `%debug` in the IPython console could be useful. We then get access to the local namespace at the source of the exception, which can allow us to explore in more detail why the exception was raised.

```
In [24]: fib(1.0)
-----
TypeError                                Traceback (most recent call last)
<ipython-input-24-874ca58a3dfb> in <module>()
----> 1 fib.fib(1.0)
/Users/rob/code/fib.py in fib(n)
      5     """
      6     f0, f1 = 0, 1
----> 7     f = [1] * n
      8     for i in range(1, n):
      9         f[n] = f0 + f1
TypeError: can't multiply sequence by non-int of type 'float'
In [25]: %debug
> /Users/rob/code/fib.py(7)fib()
      6     f0, f1 = 0, 1
----> 7     f = [1] * n
      8     for i in range(1, n):
ipdb> print(n)
1.0
```

■ **Tip** Type a question mark at the debugger prompt to show a help menu that lists available commands.

```
ipdb> ?
```

More information about the Python debugger and its features is available in the Python Standard Library documentation: <http://docs.python.org/3/library/pdb.html>.

⁵The Python function `dir` provides a similar feature.

Reset

Resetting the namespace of an IPython session is often useful to ensure that a program is run in a pristine environment, uncluttered by existing variables and functions. The `%reset` command provides this functionality (use the `-f` flag to force the reset). Using this command can often eliminate the need for otherwise common exit-restart cycles of the console. Although it is necessary to reimport modules after the `%reset` command has been used, it is important to know that even if the modules have changed since the last import, a new import after a `%reset` does not import the new module but rather reenables a cached version of the module from the previous import. When developing Python modules, this is usually not the desired behavior. In that case, a reimport of a previously imported (and since updated) module can often be achieved by using the `reload` function from `IPython.lib.deepreload`. However, this method does not always work, as some libraries run code at import time that is only intended to run once. In this case, the only option might be to terminate and restart the IPython interpreter.

Timing and Profiling Code

The `%timeit` and `%time` commands provide simple benchmarking facilities useful when looking for bottlenecks and attempting to optimize code. The `%timeit` command runs a Python statement several times and estimates the runtime (use `%%timeit` to do the same for a multiline cell). The exact number of times the statement is run is determined heuristically unless explicitly set using the `-n` and `-r` flags. See `%timeit?` for details. The `%timeit` command does not return the resulting value of the expression. If the result of the computation is required, the `%time` or `%%time` (for a multiline cell) commands can be used instead, but `%time` and `%%time` only run the statement once and give a less accurate estimate of the average runtime.

The following example demonstrates a typical usage of the `%timeit` and `%time` commands.

```
In [26]: %timeit fib(100)
100000 loops, best of 3: 16.9 µs per loop
In [27]: result = %time fib(100)
CPU times: user 33 µs, sys: 0 ns, total: 33 µs
Wall time: 48.2
```

While the `%timeit` and `%time` commands are useful for measuring the elapsed runtime of a computation, they do not give detailed information about what part of the computation takes more time. Such analyses require a more sophisticated code profiler, such as the one provided by the Python standard library module `cProfile`.⁶ The Python profiler is accessible in IPython through the `%prun` (for statements) and `%run` commands with the `-p` flag (for running external script files). The output from the profiler is rather verbose and can be customized using optional flags to the `%prun` and `%run -p` commands (see `%prun?` for a detailed description of the available options).

As an example, consider a function that simulates N random walkers, each taking M steps, and then calculates the furthest distance from the starting point achieved by any of the random walkers.

```
In [28]: import numpy as np
In [29]: def random_walker_max_distance(M, N):
...:     """
...:         Simulate N random walkers taking M steps, and return the largest
...:         Distance from the starting point achieved by any of the random
...:         walkers.
```

⁶Which can, for example, be used with the standard Python interpreter to profile scripts by running `python -m cProfile script.py`

```

...: """
...:     trajectories = [np.random.randn(M).cumsum() for _ in range(N)]
...:     return np.max(np.abs(trajectories))

```

Calling this function using the profiler with `%prun` results in the following output, which includes information about how many times each function was called and a breakdown of the total and cumulative time spent in each function. From this information, we can conclude that in this simple example, the calls to the `np.random.randn` function consume the bulk of the elapsed computation time.

In [30]: `%prun random_walker_max_distance(400, 10000)`

20011 function calls in 0.285 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
10000	0.181	0.000	0.181	0.000	{method 'randn' of 'mtrand.RandomState' objects}
10000	0.053	0.000	0.053	0.000	{method 'cumsum' of 'numpy.ndarray' objects}
1	0.020	0.020	0.277	0.277	2615584822.py:3(random_walker_max_distance)
1	0.019	0.019	0.253	0.253	2615584822.py:8(<listcomp>)
1	0.008	0.008	0.285	0.285	<string>:1(<module>)
1	0.004	0.004	0.004	0.004	{method 'reduce' of 'numpy.ufunc' objects}
1	0.000	0.000	0.285	0.285	{built-in method builtins.exec}
1	0.000	0.000	0.004	0.004	fromnumeric.py:71(_wrapreduction)
1	0.000	0.000	0.004	0.004	fromnumeric.py:2692(max)
1	0.000	0.000	0.000	0.000	fromnumeric.py:72(<dictcomp>)
1	0.000	0.000	0.000	0.000	fromnumeric.py:2687(_max_dispatcher)
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
1	0.000	0.000	0.000	0.000	{method 'items' of 'dict' objects}

Interpreter and Text Editor as Development Environment

In principle, the Python or the IPython interpreter and a good text editor are all required for a fully productive Python development environment. This simple setup is, in fact, the preferred development environment for many experienced programmers. However, the following sections look at the Jupyter Notebook and Spyder's integrated development environment. These environments provide richer features that improve productivity when working with interactive and exploratory computing applications.

Jupyter

The Jupyter project⁷ is a **spin-off from the IPython project** that includes the Python independent **frontends**—most notably the notebook application, which is discussed in more detail in the following section—and the **communication framework** that enables the separation of the **frontend** from the computational **backends**,

⁷For more information about Jupyter, see <http://jupyter.org>.

known as **kernels**. Prior to the creation of the Jupyter project, the notebook application and its underlying framework were a part of the **IPython project**. However, because the notebook frontend is language **agnostic** (it can also be used with many other languages, such as R and Julia), it was spun off a separate project to better cater to the **wider computational community** and avoid a perceived **bias** toward Python. Now, the remaining role of IPython is to focus on Python-specific applications, such as the interactive Python console, and to provide a Python **kernel** for the Jupyter environment.

In the **Jupyter framework**, the front end can be connected to multiple computational backend kernels, for example, for different programming languages, versions of Python, or for different Python environments. The kernel maintains the **state of the interpreter**. It performs the actual computations, while the front end manages how code is **entered and organized** and how the results of calculations are **visualized** to the user.

This section discusses the Jupyter QtConsole and Notebook frontends. It briefly introduces some of their rich display and interactivity features and the workflow organization that the notebook provides. The Jupyter Notebook is the **Python environment for computational work** that I generally recommend in this book, and the code listings in the rest of this book are understood to be read as if they are cells in a notebook.

The Jupyter QtConsole

The Jupyter QtConsole is an **enhanced console application** that can substitute for the standard IPython console. The QtConsole is launched by passing the `qtconsole` argument to the `jupyter` command.

```
$ jupyter qtconsole
```

This opens a new **IPython application** in a console that can display **rich media objects** such as images, figures, and mathematical equations. The Jupyter QtConsole also provides a menu-based mechanism for displaying **autocompletion results**, and it shows **docstrings** for functions in a **pop-up window** when typing the opening parenthesis of a function or a method call. A screenshot of the Jupyter Qtconsole is shown in Figure 1-3.

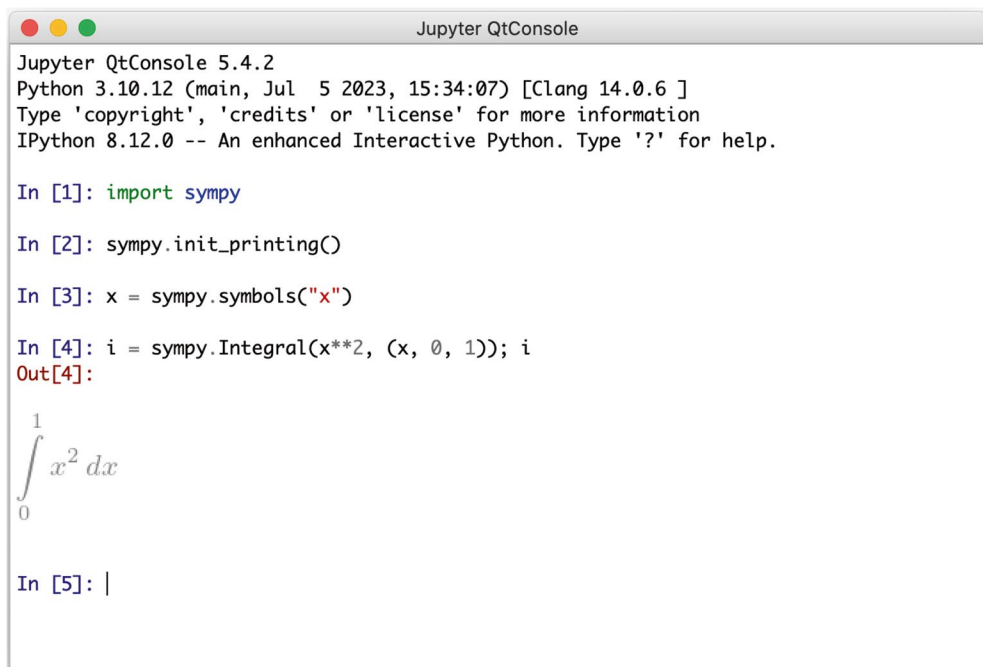


Figure 1-3. A screenshot of the Jupyter QtConsole application

The Jupyter Notebook

In addition to the **interactive console**, Jupyter also provides a **web-based notebook application** that has made it famous. The notebook offers **many** advantages over a traditional development environment when working with **data analysis and computational problem-solving**. In particular, the notebook environment allows us to write and run code, display the output produced by the code, and document and interpret the code and the results—**all in one document**. This means the **entire** analysis workflow is captured in one file, which can be saved, restored, and reused **later**. In contrast, when working with a text editor or an IDE, the code, the corresponding data files and figures, and the documentation are **spread out over multiple files** in the file system, and it takes a **significant effort and discipline** to keep such a **workflow organized**.

The Jupyter Notebook features a **rich display system** that can show media such as **equations, figures, and videos** as embedded objects in the notebook. Creating user interface (UI) elements with HTML and JavaScript is possible using **Jupyter's widget system**. These widgets can be used in interactive applications that connect the **web application with Python code executed** in the IPython kernel (on the server side). These and many other features of the Jupyter Notebook make it a great environment for **interactive and literate computing**, as we will see in examples throughout this book.

To launch the Jupyter Notebook environment, the notebook argument is passed to the jupyter command-line application.

```
$ jupyter notebook
```

This launches a notebook kernel and a web application that, by default, serves up a **web server** on port **8888** on localhost, which is accessed using the local address `http://localhost:8888/` in a web browser.⁸ By default, running `jupyter notebook` opens a web page in the default web browser. The application lists all notebooks that are available in the directory from where the Jupyter Notebook was launched, as well as a simple directory browser that can be used to **navigate subdirectories** relative to the location where the notebook server was launched and to open notebooks from therein. Figure 1-4 shows a screenshot of a web browser and the Jupyter Notebook web application.

⁸This web application is by default only accessible locally from the system where the notebook application was launched.




Figure 1-4. A screenshot of the Jupyter Notebook *web application*

Clicking the New button creates a new notebook and opens it in a new page in the browser (see Figure 1-5). A newly created notebook is named Untitled, or Untitled1, for example, depending on the availability of unused filenames. A notebook can be renamed by clicking the title field on the top of the notebook page. The Jupyter Notebook files are stored in JSON format using the `ipynb` filename extension. A Jupyter Notebook file is not pure Python code. When necessary, the Python code in a notebook can easily be extracted using either File ► Download as ► Python or the Jupyter utility `nbconvert` (see in the following section).

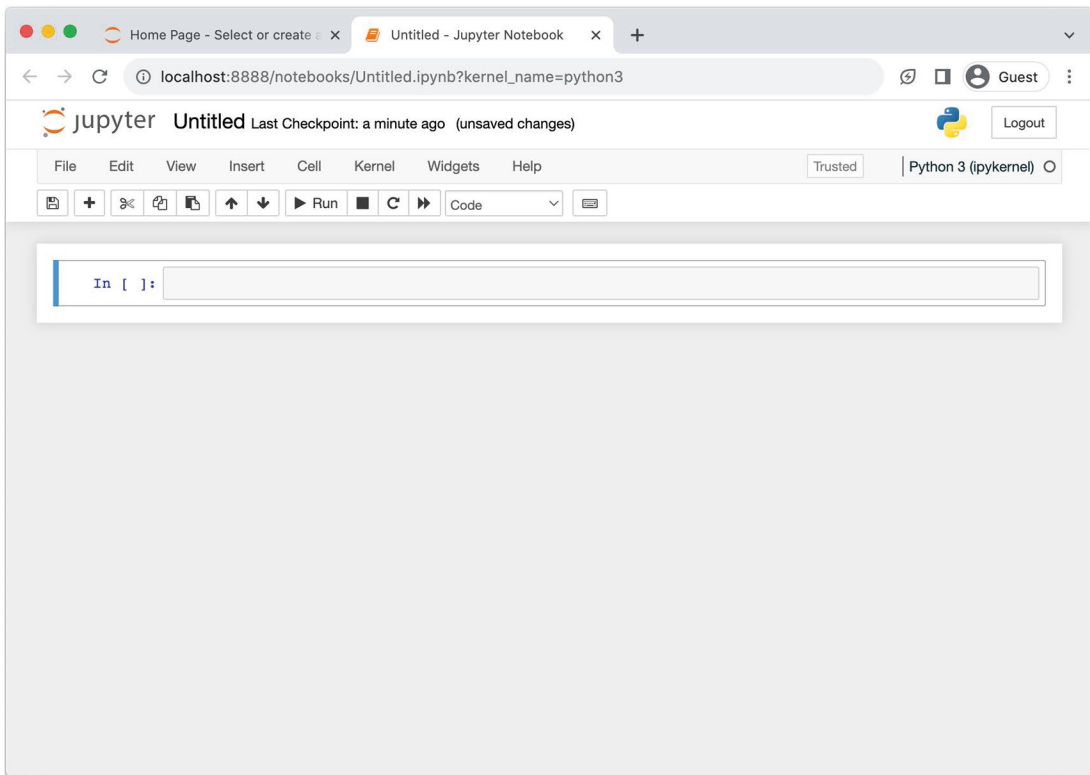


Figure 1-5. A newly created and empty Jupyter Notebook

Jupyter Lab

Jupyter Lab is an **alternative development environment** from the Jupyter project. The application can be launched by running `jupyter lab` from the command line. It combines the Jupyter Notebook interface with a file browser, text editor, shell, and IPython consoles in a web-based IDE-like environment; see Figure 1-6.

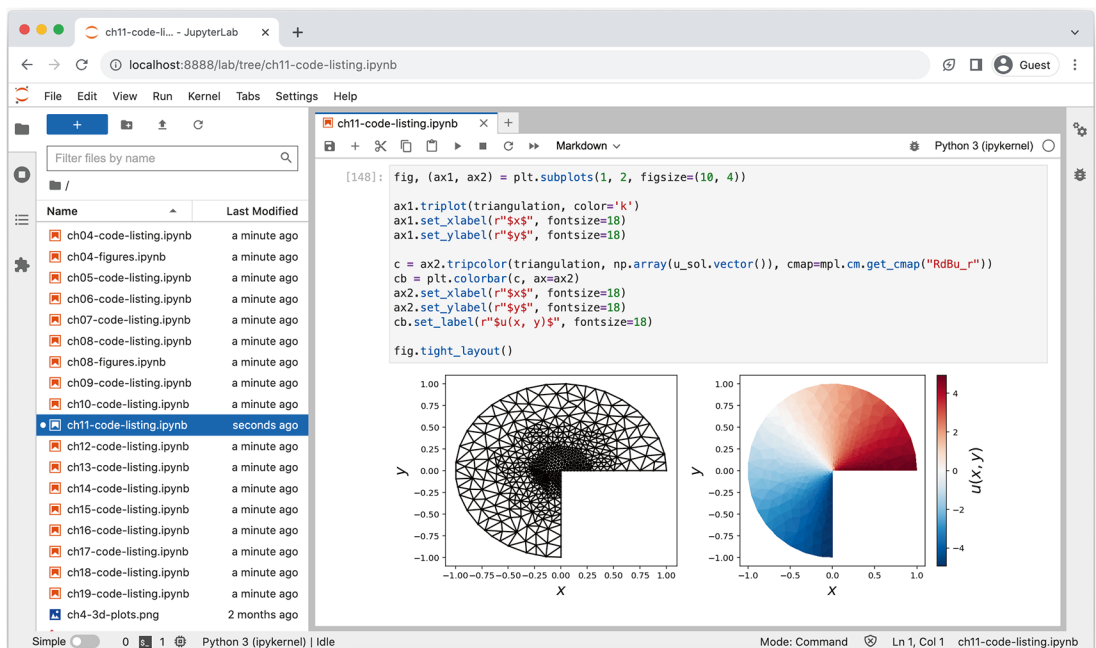


Figure 1-6. The Jupyter Lab interface includes a **file browser** (left) and **multitab notebook editor** (right). The notebook displayed here shows code and output from an example in Chapter 11

The Jupyter Lab environment consolidates the many advantages of the notebook environment and the strengths of traditional **IDEs**. Access to shell consoles and text editors all within the same web frontend is also convenient when working on a **Jupyter server** that runs on a remote system, such as a computing cluster or in the cloud.

Cell Types

The main content of a notebook, below the menu bar and the toolbar, is organized as input and output cells. The cells can be of several types, and the selected cell type can be changed using the cell-type drop-down menu in the toolbar (which initially displays “Code”). The following describes the most important types.

- **Code:** A **code cell** can contain an arbitrary amount of multiline Python code. Pressing **Shift-Enter** sends the code in the cell to the kernel process, where the kernel evaluates it using the **Python interpreter**. The result is sent back to the browser and displayed in the corresponding **output cell**.
- **Markdown:** The content of a Markdown cell can contain **marked-up plain text**, which is interpreted using the Markdown language and HTML. A Markdown cell can also contain **LaTeX formatted equations**, which are rendered in the notebook using the **JavaScript-based LaTeX engine MathJax**.
- **Raw:** A raw text cell is displayed **without any processing**.

Editing Cells

Using the menu bar and the toolbar, cells can be added, removed, moved up and down, cut and pasted, and so on. These functions are also mapped to keyboard shortcuts, which are convenient and time-saving when working with Jupyter Notebooks. The notebook uses a two-mode input interface, with an edit mode and a command mode. The edit mode can be entered by clicking a cell or by pressing the Enter key on the keyboard when a cell is in focus. Once in edit mode, the content of the input cell can be edited. Leaving the edit mode is done by pressing the ESC key or by using Shift-Enter to execute the cell. In command mode, the up and down arrows can be used to move focus between cells, and many keyboard shortcuts are mapped to the basic cell manipulation actions available through the toolbar and the menu bar. Table 1-1 summarizes the most important Jupyter Notebook keyboard shortcuts for the command mode.

Table 1-1. A Summary of Keyboard Shortcuts in the Jupyter Notebook Command Mode

Keyboard Shortcut	Description
b	Create a new cell <i>below</i> the currently selected cell.
a	Create a new cell <i>above</i> the currently selected cell.
d-d	Delete the currently selected cell.
1 to 6	Heading cell for levels 1 to 6.
x	Cut the currently selected cell.
c	Copy the currently selected cell.
v	Paste the cell from the clipboard.
m	Convert a cell to a Markdown cell.
y	Convert a cell to a code cell.
Up	Select the previous cell.
Down	Select the next cell.
Enter	Enter edit mode.
Escape	Exit edit mode.
Shift-Enter	Run the cell.
0-0	Restart the kernel.
i-i	Interrupt an executing cell.
s	Save the notebook.

While a notebook cell is being executed, the input prompt number is represented with an asterisk, In[*], and an indicator in the upper right corner of the page signals that the IPython kernel is busy. The execution of a cell can be interrupted using the menu option Kernel ► Interrupt or by typing i-i in the command mode (i.e., press the i key twice in a row).

Markdown Cells

One of the key features of the Jupyter Notebook is that code cells and output cells can be complemented with documentation contained in text cells. Text input cells are called *Markdown cells*. The input text is interpreted and reformatted using the Markdown markup language. The Markdown language is designed to be a lightweight typesetting system that allows text with simple markup rules to be converted to HTML and other formats for richer display. The markup rules are designed to be user-friendly and readable as is in plain-text format. For example, a piece of text can be made italics by surrounding it with asterisks, `*text*`. We can make it bold by surrounding it with double asterisks, `**text**`. Markdown also supports creating enumerated and bulleted lists, tables, and hyper-references. Jupyter supports an extension to Markdown that allows mathematical expressions to be typeset in LaTeX using the JavaScript LaTeX library MathJax. Generously documenting the code and the resulting output using Markdown cells and the many rich display options they provide is a good way to take advantage of Jupyter's offers. Table 1-2 introduces basic Markdown and equation formatting features that can be used in a Jupyter Notebook Markdown cell.

Table 1-2. Summary of Markdown Syntax for Jupyter Notebook Markdown Cells

Function	Syntax by Example
Italics	<code>*text*</code>
Bold	<code>**text**</code>
Strike-through	<code>~~text~~</code>
Fixed-width font	<code>`text`</code>
URL	<code>[URL text](http://www.example.com)</code>
New paragraph	Separate the text of two paragraphs with an empty line.
Verbatim	Lines that start with four blank spaces are displayed as is, without any further processing, using a fixed-width font. This is useful for code-like text segments. <pre> ___def func(x): ___ return x ** 2 </pre>
Table	<pre> A B C +---+---+---+ 1 2 3 +---+---+---+ 4 5 6 </pre>
Horizontal line	A line containing three dashes is rendered as a horizontal line separator: <pre> ---</pre>
Heading	<pre> # Level 1 heading ## Level 2 heading ### Level 3 heading ... </pre>
Block quote	Lines that start with a “>” are rendered as a block quote. <pre> > Text here is indented and offset > from the main text body. </pre>

(continued)

Table 1-2. (continued)

Function	Syntax by Example
Unordered list	* Item one * Item two * Item three
Ordered list	1. Item one 2. Item two 3. Item three
Image	![Alternative text](image-file.png) ⁹ or ![Alternative text](http://www.example.com/image.png)
Inline LaTeX equation	\LaTeX
Displayed LaTeX equation (centered and on a new line)	\LaTeX or $\begin{env} \dots \end{env}$ where env can be a LaTeX environment such as equation, eqnarray, align, etc.

Markdown cells can also contain HTML code, and the Jupyter Notebook interface displays it as rendered HTML. It is a very powerful feature for the Jupyter Notebook. Its disadvantage is that such HTML code cannot be converted to other formats, such as PDF, using the `nbconvert` tool (see later section in this chapter). Therefore, it is generally better to use Markdown formatting when possible and resort to HTML only when necessary.

More information about MathJax and Markdown is available on the projects' web pages at www.mathjax.com and <http://daringfireball.net/projects/markdown>, respectively.

Rich Output Display

The result produced by the last statement in a notebook cell is normally displayed in the corresponding output cell, just like in the standard Python interpreter or the IPython console. The default output cell formatting is a string representation of the object generated, for example, by the `__repr__` method. However, the notebook environment enables a much richer output formatting, as it, in principle, allows displaying arbitrary HTML in the output cell area. The `IPython.display` module provides several classes and functions that make it easy to programmatically render formatted output in a notebook. For example, the `Image` class provides a way to display images from the local file system or online resources in a notebook, as shown in Figure 1-7. Other useful classes from the same module are `HTML`, for rendering HTML code, and `Math`, for rendering LaTeX expressions. The `display` function can explicitly request an object to be rendered and displayed in the output area.



Figure 1-7. An example of rich Jupyter Notebook output cell formatting is where an image has been displayed in the cell output area using the `Image` class

⁹The path/filename is relative to the notebook directory.

An example of how HTML code can be rendered in the notebook using the HTML class is shown in Figure 1-8. Here, we first construct a string containing HTML code for a table with version information for a list of Python libraries. This HTML code is then rendered in the output cell area by creating an instance of the HTML class. Since this statement is the last (and only) statement in the corresponding input cell, Jupyter renders the representation of this object in the output cell area.

```
[3]: import scipy, numpy, matplotlib
modules = [numpy, matplotlib, scipy]
row = "<tr> <td>%s</td> <td>%s</td> </tr>"
rows = "\n".join([row % (module.__name__, module.__version__) for module in modules])
s = "<table> <tr><th>Library</th><th>Version</th> </tr> %s</table>" % rows
```

```
[4]: s
```

```
[4]: '<table> <tr><th>Library</th><th>Version</th> </tr> <tr> <td>numpy</td> <td>1.25.2</td> </tr>\n<tr> <td>matplotlib</td> <td>3.7.1</td> </tr>\n<tr> <td>scipy</td> <td>1.11.1</td> </tr></table>'
```

```
[5]: HTML(s)
```

```
[5]:
```

Library	Version
numpy	1.25.2
matplotlib	3.7.1
scipy	1.11.1

Figure 1-8. Example of rich Jupyter Notebook output cell formatting, where an HTML table containing module version information has been rendered and displayed using the HTML class

For an object to be displayed in an HTML formatted representation, all we need to do is to add a method called `_repr_html_` to the class definition. For example, we can easily implement your own primitive version of the HTML class and use it to render the same HTML code as in the previous example, as demonstrated in Figure 1-9.

```
[6]: class HTMLDisplay(object):
    def __init__(self, code):
        self.code = code

    def _repr_html_(self):
        return self.code
```

```
[7]: HTMLDisplay(s)
```

```
[7]:
```

Library	Version
numpy	1.25.2
matplotlib	3.7.1
scipy	1.11.1

Figure 1-9. Example of how to render HTML code in the Jupyter Notebook, using a class that implements the `_repr_html_` method

Jupyter supports a large number of representations in addition to the `_repr_html_` shown in the preceding text, for example, `_repr_png_`, `_repr_svg_`, and `_repr_latex_`, to mention a few. The former two can generate and display graphics in the notebook output cell, as used by, for example, the Matplotlib library (see the following interactive example and Chapter 4). The Math class, which uses the `_repr_latex_` method, can render mathematical formulas in the Jupyter Notebook. This is often useful in scientific and technical applications. Examples of how formulas can be rendered using the Math class and the `_repr_latex_` method are shown in Figure 1-10.

```
[8]: Math(r'\hat{H} = -\frac{1}{2}\epsilon\hat{\sigma}_z - \frac{1}{2}\delta\hat{\sigma}_x')
```

```
[8]: 
$$\hat{H} = -\frac{1}{2}\epsilon\hat{\sigma}_z - \frac{1}{2}\delta\hat{\sigma}_x$$

```

```
[9]: class QubitHamiltonian(object):
    def __init__(self, epsilon, delta):
        self.epsilon = epsilon
        self.delta = delta

    def _repr_latex_(self):
        return "$\hat{H} = -%.2f\hat{\sigma}_z - %.2f\hat{\sigma}_x$" % \
            (self.epsilon/2, self.delta/2)
```

```
[10]: QubitHamiltonian(0.5, 0.25)
```

```
[10]: 
$$\hat{H} = -0.25\hat{\sigma}_z - 0.12\hat{\sigma}_x$$

```

Figure 1-10. An example of how a LaTeX formula is rendered using the Math class and how the `_repr_latex_` method can be used to generate a LaTeX formatted representation of an object

Using the various representation methods recognized by Jupyter or the convenience classes in the IPython.display module, we have great flexibility in shaping how results are visualized in the Jupyter Notebook. However, the possibilities do not stop there: an exciting feature of the Jupyter Notebook is that interactive applications, with two-way communication between the frontend and the backend kernel, can be created using, for example, a library of widgets (UI components) or directly with JavaScript and HTML. For example, using the interact function from the ipywidgets library, we can very easily create an interactive graph that takes an input parameter that is determined from a UI slider, as shown in Figure 1-11.

```
[11]: import matplotlib.pyplot as plt
import numpy as np
from scipy import stats

def f(mu):
    X = stats.norm(loc=mu, scale=np.sqrt(mu))
    N = stats.poisson(mu)
    x = np.linspace(0, X.ppf(0.999))
    n = np.arange(0, x[-1])

    fig, ax = plt.subplots()
    ax.plot(x, X.pdf(x), color='black', lw=2, label="Normal( $\mu$ =%d,  $\sigma^2$ =%d)" % (mu, mu))
    ax.bar(n, N.pmf(n), align='edge', label=r"Poisson( $\lambda$ =%d)" % mu)
    ax.set_ylim(0, X.pdf(x).max() * 1.25)
    ax.legend(loc=2, ncol=2)
    plt.close(fig)
    return fig

[12]: from ipywidgets import interact
import ipywidgets as widgets

[13]: interact(f, mu=widgets.FloatSlider(min=1.0, max=20.0, step=1.0));
```




Figure 1-11. An example of an interactive application created using the IPython widget `interact`. The `interact` widget provides a slider UI element that allows the value of an input parameter to be changed. When the slider is dragged, the provided function is reevaluated, which, in this case, renders a new graph

The example in Figure 1-11 plots the distribution functions for the Normal distribution and the Poisson distribution, where the distributions' mean and variance are taken as input from the UI object created by the `interact` function. Moving the slider back and forth shows how the Normal and the Poisson distributions (with equal variance) approach each other as the distribution mean increases and how they behave very differently for small mean values. Interactive graphs like this are a great tool for building intuition and exploring computation problems, and the Jupyter Notebook is a fantastic enabler for this kind of investigation.¹⁰

¹⁰ For more information about how to create interactive applications using Jupyter and IPython widgets, see the documentation for the `ipywidgets` library <https://ipywidgets.readthedocs.io/en/latest>.

nbconvert

Jupyter Notebooks can be converted to many different read-only formats using the nbconvert application, which can be invoked by calling nbconvert from the command line. Supported formats include, among others, PDF and HTML. Converting Jupyter Notebooks to PDF or HTML is useful when sharing notebooks with colleagues or publishing them online when the reader does not necessarily need to run the code but primarily views the results contained in the notebooks.

HTML

In the notebook web application, the menu option File ► Download as ► HTML can generate an HTML document representing a static view of a notebook. An HTML document can also be generated from the command prompt using the nbconvert application. For example, a notebook called Notebook.ipynb can be converted to HTML using the following command.

```
$ nbconvert --to html Notebook.ipynb
```

This generates an HTML page that is self-contained in terms of style sheets and JavaScript resources (which are loaded from public CDN servers), and it can be published as is online. However, image resources using Markdown or HTML tags are not included and must be distributed with the resulting HTML file.

For public online publishing of Jupyter Notebooks, the Jupyter project provides a convenient web service called nbviewer, available at <http://nbviewer.org>. By feeding it a URL to a public notebook file, the nbviewer application automatically converts the notebook to HTML and displays the result. One of the many benefits of this method of publishing Jupyter Notebooks is that the notebook author only needs to maintain one file—the notebook file itself—and when it is updated and uploaded to its online location, the static view of the notebook provided by nbviewer is automatically updated as well. However, it requires publishing the source notebook at a publicly accessible URL, so it can only be used for public sharing.

■ **Tip** The Jupyter project maintains a Wiki page that indexes many interesting Jupyter Notebooks published online at <http://github.com/jupyter/jupyter/wiki#a-gallery-of-interesting-Jupyter-Notebooks>. These notebooks demonstrate many of IPython's and Jupyter's more advanced features. They can be a great resource for learning more about Jupyter Notebooks as well as the many topics covered by those notebooks.

PDF

Converting a Jupyter Notebook to PDF format requires first converting the notebook to LaTeX and then compiling the LaTeX document to PDF format. To do the LaTeX to PDF conversion, a LaTeX environment must be available on the system (see Appendix A for pointers on how to install these tools). The nbconvert application can do both the notebook-to-LaTeX and the LaTeX-to-PDF conversions in one go, using the --to pdf argument (the --to latex argument can be used to obtain the intermediate LaTeX source).

```
$ nbconvert --to pdf Notebook.ipynb
```

The style of the resulting document can be specified using the --template name argument, where built-in templates include base, article, and report (these templates can be found in the nbconvert/templates/latex directory where Jupyter is installed). Extending one of the existing templates makes it easy to customize

the appearance of the generated document. For example, in LaTeX it is common to include additional information about the document that is not available in Jupyter Notebooks, such as a document title (if different from the notebook filename) and the author of the document. This information can be added to a LaTeX document generated by the nbconvert application by creating a custom template. For example, the following template extends the built-in template `article` and overrides the title and author blocks.

```
((*- extends 'article.tplx' -*))
((* block title *)) \title{Document title} ((* endblock title *))
((* block author *)) \author{Author's Name} ((* endblock author *))
```

Assuming that this template is stored in a file called `custom_template.tplx`, the following command can be used to convert a notebook to PDF format using this customized template.

```
$ nbconvert --to pdf --template custom_template.tplx Notebook.ipynb
```

The result is LaTeX and PDF documents where the title and author fields are set as requested in the template.

Python

A Jupyter Notebook in its JSON-based file format can be converted to Python code using the nbconvert application and the `python` format.

```
$ nbconvert --to python Notebook.ipynb
```

This generates the `Notebook.py` file, which only contains executable Python code (or, if IPython extensions were used in the notebook, a file that is executable with `ipython`). The non-code content of the notebook is also included in the resulting Python code file in the form of comments that do not prevent the file from being interpreted by the Python interpreter. Converting a notebook to Python code is useful, for example, when using the Jupyter Notebooks to develop functions and classes that need to be imported into other Python files or notebooks.

Spyder: An Integrated Development Environment

An **integrated development environment** is an enhanced text editor with features such as code execution, documentation, and debugging from within the editor. Many free and commercial IDE environments have good support for Python-based projects. Spyder¹¹ is an **excellent free** IDE that is particularly well suited for **computing and data analysis** using Python. The rest of this section focuses on Spyder and further explores its features. However, there are also many other suitable IDEs. For example, **Eclipse**¹² is a popular and powerful multilanguage IDE, and the **PyDev**¹³ extension to Eclipse provides a good **Python environment**. **PyCharm**¹⁴ is another powerful Python IDE that has gained significant popularity among Python developers recently, and **Visual Studio Code**¹⁵ from Microsoft is yet another great option. For readers with previous experience with any of these tools, they could be a productive and familiar environment for computational work.

¹¹ <https://spyder-ide.org>

¹² <http://www.eclipse.org>

¹³ <http://pydev.org>

¹⁴ <http://www.jetbrains.com/pycharm>

¹⁵ <https://code.visualstudio.com>

However, the Spyder IDE was specifically created for Python programming, particularly for scientific computing with Python. As such, it has features useful for interactive and exploratory computing, most notably, integration with the IPython console directly in the IDE. The Spyder user interface consists of several optional panes, which can be arranged in different ways within the IDE application. The following are the **most important panes**.

- Source code editor
- Consoles for the Python and the IPython interpreters and the system shell
- Object inspector, for showing documentation for Python objects
- Variable explorer
- File explorer
- Command history
- Profiler

Each pane can be configured to be shown or hidden, depending on the user's preferences and needs, using the **View ► Panes** menu. Furthermore, panes can be organized together in tabbed groups. In the default layout, three pane groups are displayed. The left pane group contains the source code editor. The top-right pane group contains the variable explorer, the file explorer, and the object inspector. The bottom-right pane group contains Python and IPython consoles.

Running the `spyder` command at the shell prompt launches the Spyder IDE. Figure 1-12 shows a screenshot of the default layout of the Spyder application.




Figure 1-12. A screenshot of the Spyder IDE application. The code editor is in the left panel, the top-right panel shows the object inspector (help viewer), and the bottom-right panel shows an IPython console

Source Code Editor

The source code editor in Spyder supports **code highlighting**, intelligent **autocompletion**, working with multiple open files simultaneously, parenthesis matching, indentation guidance, and many other features that one would expect from a modern source code editor. The added benefit of using an IDE is that code in the editor can be run—as a whole (shortcut **F5**) or a selection (shortcut **F9**)—in attached Python or IPython consoles with persistent sessions between successive runs.

In addition, the Spyder editor has very useful support for **static code checking** with **pylint**,¹⁶ **pyflakes**,¹⁷ and **pep8**,¹⁸ which are **external tools** that analyze Python source code and report errors such as undefined symbols, syntax errors, coding style violations, and more. Such warnings and errors are shown line-by-line as a **yellow triangle** with an exclamation mark in the editor's left margin, next to the **line number**. Static code checking is extremely important in Python programming. Since Python is an interpreted and **lazily evaluated language**, simple bugs like **undefined symbols** may not be discovered until the offending code line reaches runtime. For rarely used code paths, sometimes such bugs can be very hard to discover. **Real-time static code checking** and **coding style checks** in the Spyder editor can be activated and deactivated in the preference windows (Spyder ► Preferences in macOS, and Tools ► Preferences in Linux and Windows). In the Completion and Linting section, I recommend checking “**Enable code style linting**” in the “**Code style and formatting**” tab.

■ **Tip** The Python **language** is versatile, and equivalent Python source code can be written in various styles and manners. However, a Python coding style standard, PEP8, has been implemented to encourage a uniform appearance of Python code. I strongly recommend studying the PEP8 coding style standard and complying with it in your code. The PEP8 is described at www.python.org/dev/peps/pep-0008.

Consoles in Spyder

The integrated Python and IPython consoles can be used to execute a file edited in the text editor window or run interactively typed Python code. When executing Python source code files from the editor, the namespace variables created in the script are retained in the IPython or Python session in the console. It is an important feature that makes Spyder an interactive computing environment, in addition to a traditional IDE application, since it allows exploring the values of variables after a script has finished executing. Spyder supports having multiple Python and IPython consoles opened simultaneously, and, for example, a new IPython console can be launched through the menu option Consoles ► Open an IPython console. When running a script from the editor, by pressing **F5** or the run button in the toolbar, the script is run in the most recently activated console by default. This makes it possible to maintain different consoles with independent namespaces for different scripts or projects.

Use the `%reset` command and the reload function to clear a namespace and reload updated modules when possible. If that is insufficient, it is possible to restart the IPython kernel corresponding to an IPython console or the Python interpreter via the drop-down menu for the top-right icon in the console panel. Finally, a value feature for audit trail and logging is that IPython console sessions can be exported as an HTML file by right-clicking the console window and selecting Save as HTML/XML in the pop-up menu.

¹⁶ <http://www.pylint.org>

¹⁷ <http://github.com/PyCQA/pyflakes>

¹⁸ <http://pep8.readthedocs.org>

Object Inspector

The object inspector (Help pane) is a **great aid** when writing Python code. It can display **richly formatted documentation strings** for objects defined in source code created with the editor and for symbols defined in library modules installed on the system. The **object text field** at the top of the object inspector panel can be used to **type the name** of a module, function, or class to **display the documentation string**. Modules and symbols **do not need to be imported** into the local namespace to be able to display their docstrings using the object inspector. The documentation for an object in the editor or the console can also be opened in the object inspector by selecting the object with the cursor and using the **shortcut Ctrl-i** (Cmd-i in macOS). It is even possible to **automatically display docstrings** for callable objects when its **opening left parenthesis** is typed. This gives an immediate reminder of the arguments and their order for the **callable object**, which can be a **great productivity booster**. To activate this feature, navigate to the Help page in the Preferences window and check the boxes in the **“Automatic connections”** section.

Summary

This chapter introduced the Python environment for scientific and technical computing. This environment is, in fact, an ecosystem of libraries and tools for computing, which includes not only Python software but everything from low-level number-crunching libraries up to graphical user interface applications and web applications. In this multilanguage ecosystem, Python is the language that ties it all together into a coherent and productive environment for computing. IPython is a core component of Python’s computing environment, and we briefly surveyed some of its most important features before covering the higher-level user environments provided by the Jupyter Notebook and the Spyder IDE. These are the tools in which most exploratory and interactive computing is carried out. The rest of this book focuses on computing using Python libraries, assuming that we are working within one of the environments provided by IPython, the Jupyter Notebook, or Spyder.

Further Reading

The Jupyter Notebook is a particularly rich platform for interactive computing, and it is also a very actively developed software. One of the most recent developments within the Jupyter Notebook is its widget system, which consists of user-interface components that can be used to create interactive interfaces within the browser that is displaying the notebook. This book briefly touches upon Jupyter widgets, but it is a very interesting and rapidly developing part of the Jupyter project, and I do recommend exploring their potential applications for interactive computing. The Jupyter Notebook widgets are well documented at <https://ipywidgets.readthedocs.io/en/latest>. There are also two interesting books by Cyrille Rossant on this topic that I highly recommend: *Learning IPython for Interactive Computing and Data Visualization* (Packt, 2013) and *IPython Interactive Computing and Visualization Cookbook* (Packt, 2014).