

CHAPTER 2



Vectors, Matrices, and Multidimensional Arrays

Vectors, matrices, and arrays of higher dimensions are essential tools in numerical computing. They enable us to represent and manipulate data efficiently, particularly when computations need to be repeated for a set of input values. By formulating computations in terms of array operations, we can perform vectorized¹ computing, eliminating the need for many explicit loops over array elements by applying batch operations to the array data. Vectorized computing results in concise and more maintainable code, allowing the implementation of array operations to be delegated to more efficient low-level libraries. As a result, vectorized computations can be significantly faster than sequential element-by-element computations. This is particularly important in interpreted languages like Python, where looping over arrays element-by-element incurs significant performance overhead.

In Python's scientific computing environment, the NumPy library provides efficient data structures for working with arrays. The core of NumPy is implemented in C and provides efficient functions for manipulating and processing arrays. At first glance, NumPy arrays may resemble Python's list data structure. But an important difference is that while Python lists are generic containers of objects, NumPy arrays are homogenous and typed arrays of fixed size. Homogenous means that all elements in the array have the same data type. Fixed-size means that an array cannot be resized (without creating a new array). For these and other reasons, operations and functions acting on NumPy arrays can be much more efficient than those using Python lists. Along with the data structures for arrays, NumPy offers a large collection of fundamental operators and functions that work on these data structures. It also includes submodules that implement higher-level algorithms such as linear algebra and fast Fourier transform.

This chapter first looks at the basic data structure for arrays in NumPy, including several ways to create them. Next, let's look at operations for manipulating arrays and doing computations with arrays. NumPy's multidimensional data arrays are a foundation for nearly all numerical libraries for Python. Spending time on getting familiar with NumPy and developing an understanding of how NumPy works is therefore important.

¹ Many modern processors provide instructions that operate on arrays. These are also known as vectorized operations, but here vectorized refers to high-level array-based operations, regardless of how they are implemented at the processor level.

■ **NumPy** The NumPy library provides data structures for representing a rich variety of arrays and methods and functions for operating on such arrays. NumPy provides the numerical backend for nearly every scientific or technical library for Python. It is a very important part of the scientific Python ecosystem. At the time of writing, the latest version of NumPy is 1.24.2. More information about NumPy is available at www.numpy.org.

Importing the Modules

To use the NumPy library, we need to import it into the program. By convention, the numpy module is imported under the alias np, like so.

```
In [1]: import numpy as np
```

After this, we can access functions and classes in the numpy module using the np namespace. Throughout this book, assume that the NumPy module is imported this way.

The NumPy Array Object

The core of the NumPy library is the data structures for representing multidimensional arrays of homogeneous data. Homogeneous refers to all elements in an array having the same data type.² The main data structure for multidimensional arrays in NumPy is the ndarray class. In addition to the data stored in the array, this data structure also contains important metadata, such as shape, size, data type, and other attributes. Table 2-1 presents a detailed description of these attributes. A full list of attributes with descriptions is available in the ndarray docstring, which can be accessed by calling `help(np.ndarray)` in the Python interpreter or `np.ndarray?` in an IPython console.

Table 2-1. Basic Attributes of the ndarray Class

Attribute	Description
shape	A tuple that contains the number of elements (i.e., the length) for each dimension (axis) of the array
size	The total number of elements in the array
ndim	The number of dimensions (axes)
nbytes	The number of bytes used to store the data
dtype	The data type of the elements in the array

The following example demonstrates how these attributes are accessed for an instance data of the ndarray class.

```
In [2]: data = np.array([[1, 2], [3, 4], [5, 6]])
```

²This does not necessarily need to be the case for Python lists, which therefore can be heterogenous.

```
In [3]: type(data)
Out[3]: numpy.ndarray
In [4]: data
Out[4]: array([[1, 2],
               [3, 4],
               [5, 6]])
In [5]: data.ndim
Out[5]: 2
In [6]: data.shape
Out[6]: (3, 2)
In [7]: data.size
Out[7]: 6
In [8]: data.dtype
Out[8]: dtype('int64')
In [9]: data.nbytes
Out[9]: 48
```

Here the ndarray instance `data` is created from a nested Python list using the `np.array` function. More ways to create ndarray instances from data and rules of various kinds are introduced later in this chapter. In the preceding example, the data is a two-dimensional array (`data.ndim`) of shape 3×2 , as indicated by `data.shape`, and in total, it contains six elements (`data.size`) of type `int64` (`data.dtype`), which amounts to a total size of 48 bytes (`data.nbytes`).

Data Types

In the previous section, we encountered the `dtype` attribute of the ndarray object. This attribute describes the data type of the elements within the array. As the array is homogeneous, all elements have the same data type. The basic numerical data types supported in NumPy are shown in Table 2-2. Nonnumerical data types, such as strings, objects, and user-defined compound types, are also supported.

Table 2-2. Basic Numerical Data Types Available in NumPy

dtype	Variants	Description
int	int8, int16, int32, int64	Integers
uint	uint8, uint16, uint32, uint64	Unsigned (nonnegative) integers
bool	bool	Boolean (True or False)
float	float16, float32, float64, float128	Floating-point numbers
complex	complex64, complex128, complex256	Complex-valued floating-point numbers

For numerical work, the most important data types are `int` (for integers), `float` (for floating-point numbers), and `complex` (for complex floating-point numbers). Each of these data types comes in different sizes, such as `int32` for 32-bit integers, `int64` for 64-bit integers, and so on. This offers more fine-grained control over data types than the standard Python types, which only provide one type for integers and one type for floats. When working with data types, it's typically not essential to explicitly select the bit size. However, it's often necessary to choose whether to use arrays of integers, floating-point numbers, or complex values explicitly.

The following example demonstrates how to use the `dtype` attribute to generate arrays of integer-, float-, and complex-valued elements.

```

In [10]: np.array([1, 2, 3], dtype=int)
Out[10]: array([1, 2, 3])
In [11]: np.array([1, 2, 3], dtype=float)
Out[11]: array([ 1.,  2.,  3.])
In [12]: np.array([1, 2, 3], dtype=complex)
Out[12]: array([ 1.+0.j,  2.+0.j,  3.+0.j])

```

Once a NumPy array is created, its `dtype` cannot be changed other than by creating a new copy with type-casted array values. Typecasting an array is straightforward and can be done using the `np.array` function.

```

In [13]: data = np.array([1, 2, 3], dtype=float)
In [14]: data
Out[14]: array([ 1.,  2.,  3.])
In [15]: data.dtype
Out[15]: dtype('float64')
In [16]: data = np.array(data, dtype=int)
In [17]: data.dtype
Out[17]: dtype('int64')
In [18]: data
Out[18]: array([1, 2, 3])

```

Or, it can be done using the `astype` method of the `ndarray` class.

```

In [19]: data = np.array([1, 2, 3], dtype=float)
In [20]: data
Out[20]: array([ 1.,  2.,  3.])
In [21]: data.astype(int)
Out[21]: array([1, 2, 3])

```

When computing with NumPy arrays, the data type might get converted from one type to another, if required by the operation. For example, when adding float-valued and complex-valued arrays, the result is a complex-valued array.

```

In [22]: d1 = np.array([1, 2, 3], dtype=float)
In [23]: d2 = np.array([1, 2, 3], dtype=complex)
In [24]: d1 + d2
Out[24]: array([ 2.+0.j,  4.+0.j,  6.+0.j])
In [25]: (d1 + d2).dtype
Out[25]: dtype('complex128')

```

In some cases, depending on the application and its requirements, it is essential to create arrays with data types appropriately set to, for example, `int` or `complex`. The default type is `float`. Consider the following example.

```

In [26]: np.sqrt(np.array([-1, 0, 1]))
Out[26]: RuntimeWarning: invalid value encountered in sqrt

```

```

        array([ nan,  0.,  1.])
In [27]: np.sqrt(np.array([-1, 0, 1], dtype=complex))
Out[27]: array([ 0.+1.j,  0.+0.j,  1.+0.j])

```

Here, using the `np.sqrt` function to compute the square root of each element in an array gives different results depending on the data type of the array. The square root of -1 results in the imaginary unit (denoted as `1j` in Python) only when the data type of the array is complex.

Real and Imaginary Parts

Regardless of the value of the `dtype` attribute, all NumPy array instances have the `real` and `imag` attributes for extracting the real and imaginary parts of the array, respectively.

```

In [28]: data = np.array([1, 2, 3], dtype=complex)
In [29]: data
Out[29]: array([ 1.+0.j,  2.+0.j,  3.+0.j])
In [30]: data.real
Out[30]: array([ 1.,  2.,  3.])
In [31]: data.imag
Out[31]: array([ 0.,  0.,  0.])

```

The same functionality is also provided by the functions `np.real` and `np.imag`, which also can be applied to other array-like objects, such as Python lists. Note that Python supports complex numbers, and the `imag` and `real` attributes are also available for Python scalars.

Order of Array Data in Memory

Multidimensional arrays are stored as contiguous data in memory. There is freedom of choice in arranging the array elements in this memory segment. Consider a two-dimensional array containing rows and columns: one possible way to store this array as a consecutive sequence of values is to store the rows after each other, and another equally valid approach is to store the columns one after another. The former is called row-major format, and the latter is column-major format. Whether to use row-major or column-major is a matter of conventions. For example, row-major format is used in the C programming language, and Fortran uses the column-major format. A NumPy array can be specified to be stored in row-major format, using the keyword argument `order= 'C'`, and column-major format, using the keyword argument `order= 'F'`, when the array is created or reshaped. The default format is row-major. The 'C' or 'F' ordering of NumPy array is particularly relevant when NumPy arrays are used in interfaces with software written in C and Fortran, which is often required when working with numerical computing with Python.

Row-major and column-major ordering are special cases of strategies for mapping the index used to address an element to the offset for the element in the array's memory segment. In general, the NumPy array attribute `ndarray.strides` defines exactly how this mapping is done. The `strides` attribute is a tuple of the same length as the number of axes (dimensions) of the array. Each value in `strides` is the factor by which the index for the corresponding axis is multiplied when calculating the memory offset (in bytes) for a given index expression.

For example, consider a C-order array *A* with shape $(2, 3)$, which corresponds to a two-dimensional array with two and three elements along the first and the second dimensions, respectively. If the data type is `int32`, then each element uses 4 bytes, and the total memory buffer for the array, therefore, uses $2 \times 3 \times 4 = 24$ bytes. The `strides` attribute of this array is $(4 \times 3, 4 \times 1) = (12, 4)$ because each increment of *m* in *A*[*n*, *m*] increases the memory offset with one item, or 4 bytes. Likewise, each increment of *n* increases the memory offset with three items or 12 bytes (because the second dimension of the array has length 3). If, on the other hand, the same array was stored in 'F' order, the `strides` would instead be $(4, 8)$. Using `strides` to describe the mapping of array index to array memory offset is clever because it can be used to describe different mapping strategies, and many common operations on arrays, such as the transpose, can be implemented by simply changing the `strides` attribute, which can eliminate the need for moving data around in the memory. Operations that only require changing the `strides` attribute result in new `ndarray` objects that refer to the same data as the original array. Such arrays are called *views*. For efficiency, NumPy strives to create views rather than copies when applying operations on arrays. This is generally a good thing, but it is important to be aware that some array operations result in views rather than new independent arrays because modifying their data also modifies the data of the original array. Several examples of this behavior are presented later in this chapter.

Creating Arrays

The previous section looked at NumPy's basic data structure for representing arrays, the `ndarray` class and the basic attributes of this class. This section focuses on functions from the NumPy library that can create `ndarray` instances.

Arrays can be generated in several ways, depending on their properties and the applications they are used for. For example, as we saw in the previous section, one way to initialize an `ndarray` instance is to use the `np.array` function on a Python list, which, for example, can be explicitly defined. However, this method is limited to small arrays. In many situations, it is necessary to generate arrays with elements that follow some given rule, such as filled with constant values, increasing integers, uniformly spaced numbers, random numbers, and so forth. In other cases, we might need to create arrays from data stored in a file. The requirements are many and varied, and the NumPy library provides a comprehensive set of functions for creating arrays in a variety of ways. This section looks at many of these functions in more detail. For a complete list, see the NumPy reference manual or the available docstrings by typing `help(np)` or using the autocompletion `np.<TAB>`. A summary of frequently used array-generating functions is given in Table 2-3.

Table 2-3. Summary of NumPy Functions for Generating Arrays

Function Name	Type of Array
<code>np.array</code>	Create an array for which the elements are given by an array-like object, which, for example, can be a (nested) Python list, a tuple, an iterable sequence, or another ndarray instance.
<code>np.zeros</code>	Create an array with the specified dimensions and data type that is filled with zeros.
<code>np.ones</code>	Create an array with the specified dimensions and data type that is filled with ones.
<code>np.diag</code>	Create a diagonal array with specified values along the diagonal and zeros elsewhere.
<code>np.arange</code>	Create an array with evenly spaced values between the specified start, end, and increment values.
<code>np.linspace</code>	Create an array with evenly spaced values between specified start and end values, using a specified number of elements.
<code>np.logspace</code>	Create an array with values that are logarithmically spaced between the given start and end values.
<code>np.meshgrid</code>	Generate coordinate matrices (and higher-dimensional coordinate arrays) from one-dimensional coordinate vectors.
<code>np.fromfunction</code>	Create an array and fills it with values specified by a given function, which is evaluated for each combination of indices for the given array size.
<code>np.fromfile</code>	Create an array with the data from a binary (or text) file. NumPy also provides a corresponding function <code>np.tofile</code> with which NumPy arrays can be stored to disk and later read back using <code>np.fromfile</code> .
<code>np.genfromtxt</code> , <code>np.loadtxt</code>	Create an array from data read from a text file, for example, a comma-separated value (CSV) file. The <code>np.genfromtxt</code> function also supports data files with missing values.
<code>np.random.rand</code>	Generate an array with random numbers that are uniformly distributed between 0 and 1. Other types of distributions are also available in the <code>np.random</code> module.

Arrays Created from Lists and Other Array-Like Objects

Using the `np.array` function, NumPy arrays can be constructed from explicit Python lists, iterable expressions, and other array-like objects (such as other ndarray instances). For example, to create a one-dimensional array from a Python list, we simply pass the Python list as an argument to the `np.array` function.

```
In [32]: np.array([1, 2, 3, 4])
Out[32]: array([ 1,  2,  3,  4])
In [33]: data.ndim
Out[33]: 1
In [34]: data.shape
Out[34]: (4,)
```

We can use a nested Python list to create a two-dimensional array with the same data as in the previous example.

```
In [35]: np.array([[1, 2], [3, 4]])
Out[35]: array([[1,  2],
                [3,  4]])
In [36]: data.ndim
Out[36]: 2
In [37]: data.shape
Out[37]: (2, 2)
```

Arrays Filled with Constant Values

The functions `np.zeros` and `np.ones` create and return arrays filled with zeros and ones, respectively. As first argument, they take an integer or a tuple that describes the number of elements along each dimension of the array. For example, to create a 2×3 array filled with zeros, and an array of length 4 filled with ones, we can use the following.

```
In [38]: np.zeros((2, 3))
Out[38]: array([[ 0.,  0.,  0.],
                [ 0.,  0.,  0.]])
In [39]: np.ones(4)
Out[39]: array([ 1.,  1.,  1.,  1.] )
```

Like other array-generating functions, the `np.zeros` and `np.ones` functions also accept an optional keyword argument that specifies the data type for the elements in the array. By default, the data type is `float64`, and it can be changed to the required type by explicitly specifying the `dtype` argument.

```
In [40]: data = np.ones(4)
In [41]: data.dtype
Out[41]: dtype('float64')
In [42]: data = np.ones(4, dtype=np.int64)
In [43]: data.dtype
Out[43]: dtype('int64')
```

An array filled with an arbitrary constant value can be generated by creating an array filled with ones and then multiplying the array with the desired fill value. However, NumPy also provides the `np.full` function that does this in one step. The following two ways of constructing arrays with ten elements, which are initialized to the numerical value 5.4 in this example, produce the same results. However, `np.full` is slightly more efficient since it avoids multiplication.

```
In [44]: x1 = 5.4 * np.ones(10)
In [45]: x2 = np.full(10, 5.4)
```

An existing array can also be filled with constant values using the `np.fill` function, which takes an array and a value as arguments and sets all elements in the array to the given value. The following two methods to create an array give the same results.

```
In [46]: x1 = np.empty(5)
In [47]: x1.fill(3.0)
In [48]: x1
```



```
Out[48]: array([ 3.,  3.,  3.,  3.,  3.])
In [49]: x2 = np.full(5, 3.0)
In [50]: x2
Out[50]: array([ 3.,  3.,  3.,  3.,  3.])
```

This last example also used the `np.empty` function, which generates an array with uninitialized values, of the given size. This function should only be used when the initialization of all elements can be guaranteed by other means, such as an explicit loop over the array elements or another explicit assignment. This function is described in more detail later in this chapter.

Arrays Filled with Incremental Sequences

In numerical computing, it is very common to require arrays with evenly spaced values between a starting value and an ending value. NumPy provides two similar functions to create such arrays: `np.arange` and `np.linspace`. Both functions take three arguments, where the first two arguments are the start and end values. The third argument of `np.arange` is the increment, while for `np.linspace`, it is the total number of points in the array.

For example, we could use either of the following to generate arrays with values between 0 and 10, with increment 1.

```
In [51]: np.arange(0.0, 11, 1)
Out[51]: array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.])
In [52]: np.linspace(0, 10, 11)
Out[52]: array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.])
```

However, note that `np.arange` does not include the end value (11), while by default `np.linspace` does (although this behavior can be changed using the optional `endpoint` keyword argument). Whether to use `np.arange` or `np.linspace` is mostly a matter of personal preference, but it is generally recommended to use `np.linspace` whenever the increment is a noninteger.

Arrays Filled with Logarithmic Sequences

The `np.logspace` function is similar to `np.linspace`, but the increments between the elements in the array are logarithmically distributed, and the first two arguments for the start and end values, are the powers of the optional `base` keyword argument (which defaults to 10). For example, we can use the following to generate an array with logarithmically distributed values between 1 and 100.

```
In [53]: np.logspace(0, 2, 5) # 5 data points between 10**0=1 to 10**2=100
Out[53]: array([ 1. ,  3.16227766, 10. , 31.6227766 , 100.])
```

Meshgrid Arrays

Multidimensional coordinate grids can be generated using the `np.meshgrid` function. Given two one-dimensional coordinate arrays (i.e., arrays containing a set of coordinates along a given dimension), we can generate two-dimensional coordinate arrays using the `np.meshgrid` function. An illustration of this is given in the following example.

```
In [54]: x = np.array([-1, 0, 1])
In [55]: y = np.array([-2, 0, 2])
In [56]: X, Y = np.meshgrid(x, y)
```

```
In [57]: X
Out[57]: array([[ -1,  0,  1],
                [ -1,  0,  1],
                [ -1,  0,  1]])

In [58]: Y
Out[58]: array([[ -2, -2, -2],
                [  0,  0,  0],
                [  2,  2,  2]])
```

A common use for two-dimensional coordinate arrays, like *X* and *Y* in this example, is to evaluate functions over two variables *x* and *y*. This can be used, for example, when plotting functions as colormap plots and contour plots. To evaluate the expression $(x + y)^2$ at all combinations of values from the *x* and *y* arrays in the preceding section, we can use the two-dimensional coordinate arrays *X* and *Y*.

```
In [59]: Z = (X + Y) ** 2
In [60]: Z
Out[60]: array([[ 9,  4,  1],
                [ 1,  0,  1],
                [ 1,  4,  9]])
```

It is also possible to generate higher-dimensional coordinate arrays by passing more arrays as an argument to the `np.meshgrid` function. Alternatively, the functions `np.mgrid` and `np.ogrid` can also be used to generate coordinate arrays using a slightly different syntax based on indexing and slice objects. See their docstrings or the NumPy documentation for details.

Creating Uninitialized Arrays

To create an array of specific size and data type, but without initializing the elements in the array to any particular values, we can use the `np.empty` function. The advantage of using this function, instead of, for example, `np.zeros`, which creates an array initialized with zero-valued elements, is that we can avoid the initiation step. If all elements are guaranteed to be initialized later in the code, this can save a little bit of time, especially when working with large arrays. To illustrate the use of the `np.empty` function, consider the following example.

```
In [61]: np.empty(3, dtype=float)
Out[61]: array([ 1.28822975e-231,  1.28822975e-231,  2.13677905e-314])
```

This generated a new array with three elements of type float. There is no guarantee that the elements have any particular values, and the actual values vary from time to time. For this reason, it is important that all values are explicitly assigned before the array is used; otherwise, unpredictable errors are likely to arise. Often the `np.zeros` function is a safer alternative to `np.empty`, and if the performance gain is not essential, it is better to use `np.zeros`, to minimize the likelihood of subtle and hard-to-reproduce bugs due to uninitialized values in the array returned by `np.empty`.

Creating Arrays with Properties of Other Arrays

It is often necessary to create new arrays that share properties, such as shape and data type, with another array. NumPy provides a family of functions for this purpose: `np.ones_like`, `np.zeros_like`, `np.full_like`, and `np.empty_like`. A typical use case is a function that takes arrays of unspecified type and size as arguments and requires working arrays of the same size and type. For example, a boilerplate example of this situation is given in the following function.

```
def f(x):
    y = np.ones_like(x)
    # compute with x and y
    return y
```

At the first line of the body of this function, a new array `y` is created using `np.ones_like`, which results in an array of the same size and data type as `x` and filled with ones.

Creating Matrix Arrays

Matrices, or two-dimensional arrays, are an important case in numerical computing. One of the useful features of NumPy is its ability to generate common matrices. The `np.identity` function, for example, creates a square matrix with ones on the diagonal and zeros in all other positions.

```
In [62]: np.identity(4)
Out[62]: array([[ 1.,  0.,  0.,  0.],
                [ 0.,  1.,  0.,  0.],
                [ 0.,  0.,  1.,  0.],
                [ 0.,  0.,  0.,  1.]])
```

The similar function `np.eye` generates matrices with ones on a diagonal (optionally offset). This is illustrated in the following example, which produces matrices with nonzero diagonals above and below the diagonal, respectively.

```
In [63]: np.eye(3, k=1)
Out[63]: array([[ 0.,  1.,  0.],
                [ 0.,  0.,  1.],
                [ 0.,  0.,  0.]])
In [64]: np.eye(3, k=-1)
Out[64]: array([[ 0.,  0.,  0.],
                [ 1.,  0.,  0.],
                [ 0.,  1.,  0.]])
```

To construct a matrix with an arbitrary one-dimensional array on the diagonal, we can use the `np.diag` function (which also takes the optional keyword argument `k` to specify an offset from the diagonal), as demonstrated in the following.

```
In [65]: np.diag(np.arange(0, 20, 5))
Out[65]: array([[0,  0,  0,  0],
                [0,  5,  0,  0],
                [0,  0, 10,  0],
                [0,  0,  0, 15]])
```

This gave a third argument to the `np.arange` function, which specifies the step size in the enumeration of elements in the array returned by the function. The resulting array, therefore, contains the values `[0, 5, 10, 15]`, which are inserted on the diagonal of a two-dimensional matrix by the `np.diag` function.

Indexing and Slicing

Elements and subarrays of NumPy arrays are accessed using the standard square bracket notation that is also used with Python lists. Within the square bracket, a variety of different index formats are used for different types of element selection. In general, the expression within the bracket is a tuple, where each item in the tuple is a specification of which elements to select from each axis (dimension) of the array.

One-Dimensional Arrays

Along a single axis, integers are used to select single elements, and slices are used to select ranges and sequences of elements. Positive integers are used to index elements from the beginning of the array (index starts at 0), and negative integers are used to index elements from the end of the array, where the last element is indexed with -1, the second to last element with -2, and so on.

Slices are specified using the `:` notation that is also used for Python lists. In this notation, a range of elements can be selected using an expression like `m:n`, which selects elements starting with *m* and ending with *n* - 1 (note that the *n*th element is not included). The slice `m:n` can also be written more explicitly as `m : n : 1`, where the number 1 specifies that every element between *m* and *n* should be selected. To select every second element between *m* and *n*, use `m : n : 2`, and to select every *p* element, use `m : n : p`, and so on. If *p* is negative, elements are returned in reversed order starting from *m* to *n* + 1 (which implies that *m* has to be larger than *n* in this case). Table 2-4 summarizes indexing and slicing operations for NumPy arrays.

Table 2-4. Examples of Array Indexing and Slicing Expressions

Expression	Description
<code>a[m]</code>	Select the element at index <i>m</i> , where <i>m</i> is an integer (start counting from 0).
<code>a[-m]</code>	Select the <i>n</i> th element from the end of the list, where <i>m</i> is an integer. The last element in the list is addressed as -1, the second to last element as -2, and so on.
<code>a[m:n]</code>	Select elements with index starting at <i>m</i> and ending at <i>n</i> - 1 (<i>m</i> and <i>n</i> are integers).
<code>a[:]</code>	Select all elements in the given axis.
<code>a[:n]</code>	Select elements starting with index 0 and going up to index <i>n</i> - 1 (integer).
<code>a[m:]</code>	Select elements starting with index <i>m</i> (integer) and going up to the last element in the array.
<code>a[m:n:p]</code>	Select elements with index <i>m</i> through <i>n</i> (exclusive), with increment <i>p</i> .
<code>a[::-1]</code>	Select all the elements, in reverse order.

The following examples demonstrate index and slicing operations for NumPy arrays. To begin with, consider an array with a single axis (dimension) that contains a sequence of integers between 0 and 10.

```
In [66]: a = np.arange(0, 11)
In [67]: a
Out[67]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

Note that the end value 11 is not included in the array. To select specific elements from this array, for example, the first, the last, and the fifth element, we can use integer indexing.

```
In [68]: a[0] # the first element
Out[68]: 0
In [69]: a[-1] # the last element
```

```
Out[69]: 10
In [70]: a[4] # the fifth element, at index 4
Out[70]: 4
```

To select a range of elements, say from the second to the second-to-last element, selecting every element and every second element, respectively, we can use index slices.

```
In [71]: a[1:-1]
Out[71]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
In [72]: a[1:-1:2]
Out[72]: array([1, 3, 5, 7, 9])
```

To select the first five and the last five elements from an array, we can use the slices `:5` and `-5:`, since if `m` or `n` is omitted in `m:n`, the defaults are the beginning and the end of the array, respectively.

```
In [73]: a[:5]
Out[73]: array([0, 1, 2, 3, 4])
In [74]: a[-5:]
Out[74]: array([6, 7, 8, 9, 10])
```

To reverse the array and select only every second value, we can use the slice `::-2`, as shown in the following example.

```
In [75]: a[::-2]
Out[75]: array([10, 8, 6, 4, 2, 0])
```

Multidimensional Arrays

With multidimensional arrays, element selections like those introduced in the previous section can be applied on each axis (dimension). The result is a reduced array where each element matches the given selection rules. As a specific example, consider the following two-dimensional array.

```
In [76]: f = lambda m, n: n + 10 * m
In [77]: A = np.fromfunction(f, (6, 6), dtype=int)
In [78]: A
Out[78]: array([[ 0, 10, 20, 30, 40, 50],
                [10, 11, 21, 31, 41, 51],
                [20, 21, 22, 32, 42, 52],
                [30, 31, 32, 33, 34, 35],
                [40, 41, 42, 43, 44, 45],
                [50, 51, 52, 53, 54, 55]])
```

We can extract columns and rows from this two-dimensional array using a combination of slice and integer indexing.

```
In [79]: A[:, 1] # the second column
Out[79]: array([10, 11, 21, 31, 41, 51])
In [80]: A[1, :] # the second row
Out[80]: array([10, 11, 12, 13, 14, 15])
```

By applying a slice on each of the array axes, we can extract subarrays (submatrices in this two-dimensional example).

```
In [81]: A[:3, :3] # upper half diagonal block matrix
Out[81]: array([[ 0,  1,  2],
               [10, 11, 12],
               [20, 21, 22]])

In [82]: A[3:, :3] # lower left off-diagonal block matrix
Out[82]: array([[30, 31, 32],
               [40, 41, 42],
               [50, 51, 52]])
```

With element spacing other than 1, submatrices made up from nonconsecutive elements can be extracted.

```
In [83]: A[::2, ::2] # every second element starting from 0, 0
Out[83]: array([[ 0,  2,  4],
               [20, 22, 24],
               [40, 42, 44]])

In [84]: A[1::2, 1::3] # every second and third element starting from 1, 1
Out[84]: array([[11, 14],
               [31, 34],
               [51, 54]])
```

This ability to extract subsets of data from a multidimensional array is a simple but very powerful feature with many applications in data processing.

Views

Subarrays extracted from arrays using slice operations are alternative *views* of the same underlying array data. More specifically, they are arrays that refer to the same data in the memory as the original array, but with a different `strides` configuration. When elements in a view are assigned new values, the original array's values are updated. The following is an example.

```
In [85]: B = A[1:5, 1:5]
In [86]: B
Out[86]: array([[11, 12, 13, 14],
               [21, 22, 23, 24],
               [31, 32, 33, 34],
               [41, 42, 43, 44]])

In [87]: B[:, :] = 0
In [88]: A
Out[88]: array([[ 0,  1,  2,  3,  4,  5],
               [10,  0,  0,  0,  0, 15],
               [20,  0,  0,  0,  0, 25],
               [30,  0,  0,  0,  0, 35],
               [40,  0,  0,  0,  0, 45],
               [50, 51, 52, 53, 54, 55]])
```

In assigning new values to the elements in array B, which is created from array A, we also modify the values in A (since both arrays refer to the same data in the memory). The fact that extracting subarrays results in views rather than new independent arrays eliminates the need for copying data and improves performance. When a copy rather than a view is needed, the view can be copied explicitly using the `copy` method of the `ndarray` instance.

```
In [89]: C = B[1:3, 1:3].copy()
In [90]: C
Out[90]: array([[0, 0],
               [0, 0]])
In [91]: C[:, :] = 1 # this does not affect B since C is a copy of the view B[1:3, 1:3]
In [92]: C
Out[92]: array([[1, 1],
               [1, 1]])
In [93]: B
Out[93]: array([[0, 0, 0, 0],
               [0, 0, 0, 0],
               [0, 0, 0, 0],
               [0, 0, 0, 0]])
```

In addition to the `copy` attribute of the `ndarray` class, an array can be copied using the `np.copy` function or, equivalently, using the `np.array` function with the `copy=True` keyword argument.

Fancy Indexing and Boolean-Valued Indexing

The previous section looked at indexing NumPy arrays with integers and slices to extract individual elements or ranges of elements. NumPy provides another convenient method to index arrays, called fancy indexing. With fancy indexing, an array can be indexed with another NumPy array, a Python list, or a sequence of integers whose values select elements in the indexed array. To clarify this concept, consider the following example: we first create a NumPy array with 11 floating-point numbers and then index the array with another NumPy array (and Python list) to extract element numbers 0, 2, and 4 from the original array.

```
In [94]: A = np.linspace(0, 1, 11)
Out[94]: array([ 0. ,  0.1,  0.2,  0.3,  0.4,  0.5,  0.6,  0.7,  0.8,  0.9,  1. ])
In [95]: A[np.array([0, 2, 4])]
Out[95]: array([ 0. ,  0.2,  0.4])
In [96]: A[[0, 2, 4]] # The same thing can be accomplished by indexing with a Python list
Out[96]: array([ 0. ,  0.2,  0.4])
```

This indexing method can be used along each axis (dimension) of a multidimensional NumPy array. It requires that the elements in the array or list used for indexing are integers.

Another variant of indexing NumPy arrays is to use Boolean-valued index arrays. In this case, each element (with values `True` or `False`) indicates whether or not to select the element from the list with the corresponding index. That is, if element n in the indexing array of Boolean values is `True`, then element n is selected from the indexed array. If the value is `False`, then element n is not selected. This indexing method is handy when filtering out elements from an array. For example, to select all the elements from the array A (as defined in the preceding section) that exceed the value 0.5, we can use the following combination of the comparison operator applied to a NumPy array and indexing using a Boolean-valued array.

```
In [97]: A > 0.5
Out[97]: array([False, False, False, False, False, False, True, True, True, True],
dtype=bool)
In [98]: A[A > 0.5]
Out[98]: array([ 0.6,  0.7,  0.8,  0.9,  1.  ])
```

Unlike arrays created using slices, the arrays returned using fancy indexing and Boolean-valued indexing are not views but new independent arrays. Nonetheless, it is possible to assign values to elements selected using fancy indexing.

```
In [99]: A = np.arange(10)
In [100]: indices = [2, 4, 6]
In [101]: B = A[indices]
In [102]: B[0] = -1 # this does not affect A
In [103]: A
Out[103]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
In [104]: A[indices] = -1 # this alters A
In [105]: A
Out[105]: array([ 0,  1, -1,  3, -1,  5, -1,  7,  8,  9])
```

And likewise for Boolean-valued indexing.

```
In [106]: A = np.arange(10)
In [107]: B = A[A > 5]
In [108]: B[0] = -1 # this does not affect A
In [109]: A
Out[109]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
In [110]: A[A > 5] = -1 # this alters A
In [111]: A
Out[111]: array([ 0,  1,  2,  3,  4,  5, -1, -1, -1, -1])
```

A visual summary of different methods to index NumPy arrays is given in Figure 2-1. Note that each type of indexing discussed here can be independently applied to each dimension of an array.




Figure 2-1. A visual summary of indexing methods for NumPy arrays. These diagrams represent NumPy arrays of shape (4, 4), and the highlighted elements are selected using the indexing expression shown above the block representations of the arrays

Reshaping and Resizing

When working with data in array form, it is often useful to rearrange arrays and alter the way they are interpreted. For example, an $N \times N$ matrix array could be rearranged into a vector of length N^2 , or a set of one-dimensional arrays could be concatenated together or stacked next to each other to form a matrix. NumPy provides a rich set of functions for this type of manipulation. Table 2-5 summarizes a selection of these functions.

Table 2-5. Summary of NumPy Functions for Manipulating the Dimensions and the Shape of Arrays

Function/Method	Description
<code>np.reshape</code> , <code>np.ndarray.reshape</code>	Reshape an N-dimensional array. The total number of elements must remain the same.
<code>np.ndarray.flatten</code>	Create a copy of an N-dimensional array and reinterprets it as a one-dimensional array (i.e., all dimensions are collapsed into one).
<code>np.ravel</code> , <code>np.ndarray.ravel</code>	Create a view (if possible, otherwise a copy) of an N-dimensional array in which it is interpreted as a one-dimensional array.
<code>np.squeeze</code>	Remove axes with length 1.
<code>np.expand_dims</code> , <code>np.newaxis</code>	Add a new axis (dimension) of length 1 to an array, where <code>np.newaxis</code> is used with array indexing.
<code>np.transpose</code> , <code>np.ndarray.transpose</code> , <code>np.ndarray.T</code>	Transpose the array. The transpose operation corresponds to reversing (or, more generally, permuting) the axes of the array.
<code>np.hstack</code>	Stack a list of arrays horizontally (along axis 1): for example, given a list of column vectors, it appends the columns to form a matrix.
<code>np.vstack</code>	Stack a list of arrays vertically (along axis 0): for example, given a list of row vectors, it appends the rows to form a matrix.
<code>np.dstack</code>	Stack arrays depth-wise (along axis 2).
<code>np.concatenate</code>	Create a new array by appending arrays after each other along a given axis.
<code>np.resize</code>	Resize an array. Create a new copy of the original array, with the requested size. If necessary, the original array is repeated to fill up the new array.
<code>np.append</code>	Append an element to an array. Create a new copy of the array.
<code>np.insert</code>	Insert a new element at a given position. Create a new copy of the array.
<code>np.delete</code>	Delete an element at a given position. Create a new copy of the array.

Reshaping an array does not require modifying the underlying array data; it only changes how the data is interpreted by redefining the array's `strides` attribute. An example of this type of operation is a 2×2 array (matrix) that is reinterpreted as a 1×4 array (vector). In NumPy, the `np.reshape` function or the `ndarray` class `reshape` method can reconfigure how the underlying data is interpreted. It takes an array and the new shape of the array as arguments.

```
In [112]: data = np.array([[1, 2], [3, 4]])
In [113]: np.reshape(data, (1, 4))
Out[113]: array([[1, 2, 3, 4]])
In [114]: data.reshape(4)
Out[114]: array([1, 2, 3, 4])
```

The requested new array shape must match the number of elements in the original size. However, the number of axes (dimensions) does not need to be conserved, as illustrated in the previous example, whereas in the first case, the new array has dimension 2 and shape `(1, 4)`. In contrast, in the second case, the new array has dimension 1 and shape `(4,)`. This example also demonstrates two ways of invoking the

reshape operation: the `np.reshape` function and the `ndarray.reshape` method. Note that reshaping an array produces a view of the array, and if an independent copy of the array is needed, the view has to be copied explicitly (e.g., using `np.copy`).

The `np.ravel` (and its corresponding `ndarray` method) is a special case of reshape, which collapses all dimensions of an array and returns a flattened one-dimensional array with a length corresponding to the total number of elements in the original array. The `ndarray` method `flatten` performs the same function but returns a copy instead of a view.

```
In [115]: data = np.array([[1, 2], [3, 4]])
In [116]: data
Out[116]: array([[1, 2],
                 [3, 4]])
In [117]: data.flatten()
Out[117]: array([ 1,  2,  3,  4])
In [118]: data.flatten().shape
Out[118]: (4,)
```

While `np.ravel` and `np.flatten` collapse the axes of an array into a one-dimensional array, it is also possible to introduce new axes into an array, either by using `np.reshape` or, when adding new empty axes, using indexing notation and the `np.newaxis` keyword at the place of a new axis. In the following example, the array `data` has one axis, so it should normally be indexed with a tuple with one element. However, if it is indexed with a tuple with more than one element, and if the extra indices in the tuple have the value `np.newaxis`, the corresponding new axes are added.

```
In [119]: data = np.arange(0, 5)
In [120]: column = data[:, np.newaxis]
In [121]: column
Out[121]: array([[0],
                 [1],
                 [2],
                 [3],
                 [4]])
In [122]: row = data[np.newaxis, :]
In [123]: row
Out[123]: array([[0, 1, 2, 3, 4]])
```

The `np.expand_dims` function can also be used to add new dimensions to an array, and in the preceding example, the expression `data[:, np.newaxis]` is equivalent to `np.expand_dims(data, axis=1)`, and `data[np.newaxis, :]` is equivalent to `np.expand_dims(data, axis=0)`. Here the `axis` argument specifies the location relative to the existing axes where the new axis is to be inserted.

Up to now, we have seen methods to rearrange arrays in ways that do not affect the underlying data. Earlier in this chapter, we learned how to extract subarrays using various indexing techniques. In addition to reshaping and selecting subarrays, it is often necessary to merge arrays into bigger arrays, for example, when joining separately computed or measured data series into a higher-dimensional array, such as a matrix. For this task, NumPy provides the functions `np.vstack`, for vertical stacking of, for example, rows into a matrix, and `np.hstack` for horizontal stacking of, for example, columns into a matrix. The `np.concatenate` function provides similar functionality, but it takes a keyword argument `axis` that specifies the axis along which the arrays will be concatenated.

The shape of the arrays passed to `np.hstack`, `np.vstack`, and `np.concatenate` is important to achieve the desired type of array joining. For example, consider the following case: we have one-dimensional arrays of data and want to stack them vertically to obtain a matrix where the rows are made up of one-dimensional arrays. We can use `np.vstack` to achieve this.

```
In [124]: data = np.arange(5)
In [125]: data
Out[125]: array([0, 1, 2, 3, 4])
In [126]: np.vstack((data, data, data))
Out[126]: array([[0, 1, 2, 3, 4],
                 [0, 1, 2, 3, 4],
                 [0, 1, 2, 3, 4]])
```

If we instead want to stack the arrays horizontally to obtain a matrix where the arrays are the column vectors, we might first attempt something similar using `np.hstack`.

```
In [127]: data = np.arange(5)
In [128]: data
Out[128]: array([0, 1, 2, 3, 4])
In [129]: np.hstack((data, data, data))
Out[129]: array([0, 1, 2, 3, 4, 0, 1, 2, 3, 4, 0, 1, 2, 3, 4])
```

This stacks the arrays horizontally, but not in the way intended here. To make `np.hstack` treat the input arrays as columns and stack them accordingly, we need to make the input arrays two-dimensional arrays of shape `(1, 5)` rather than one-dimensional arrays of shape `(5,)`. As discussed earlier, we can insert a new axis by indexing with `np.newaxis`.

```
In [130]: data = data[:, np.newaxis]
In [131]: np.hstack((data, data, data))
Out[131]: array([[0, 0, 0],
                 [1, 1, 1],
                 [2, 2, 2],
                 [3, 3, 3],
                 [4, 4, 4]])
```

The behavior of the functions for horizontal and vertical stacking, as well as concatenating arrays using `np.concatenate`, is clearest when the stacked arrays have the same number of dimensions as the final array and when the input arrays are stacked along an axis for which they have length 1.

The number of elements in a NumPy array cannot be changed once the array has been created. For example, to insert, append, and remove elements from a NumPy array using the `np.append`, `np.insert`, and `np.delete` functions, a new array must be created and the data copied to it. It may sometimes be tempting to use these functions to grow or shrink the size of a NumPy array. But, due to the overhead of creating new arrays and copying the data, pre-allocating arrays with sizes is usually a good idea to pre-allocate them with sizes such that they do not need to be resized later.

Vectorized Expressions

The purpose of storing numerical data in arrays is to be able to process the data with concise vectorized expressions that represent batch operations that are applied to all elements in the arrays. Efficient use of vectorized expressions eliminates the need for many explicit `for` loops. This approach makes the code more

concise, easier to maintain, and performs better. NumPy implements functions and vectorized operations corresponding to most fundamental mathematical functions and operators. Many of these functions and operations act on arrays on an elementwise basis, and binary operations require all arrays in an expression to be of compatible size. The meaning of compatible size is normally that the variables in an expression represent either scalars or arrays of the same size and shape. More generally, a binary operation involving two arrays is well defined if the arrays can be *broadcasted* into the same shape and size.

In an operation between a scalar and an array, broadcasting refers to the scalar being distributed and the operation applied to each element in the array. When an expression contains arrays of unequal sizes, the operations may still be well defined if the smaller of the array can be broadcasted (“effectively expanded”) to match the larger array according to NumPy’s broadcasting rule: an array can be broadcasted over another array if their axes on a one-by-one basis either have the same length or if either of them has length 1. If the number of axes of the two arrays is not equal, the array with fewer axes is padded with new axes of length 1 from the left until the numbers of dimensions of the two arrays agree.

Two simple examples that illustrate array broadcasting are shown in Figure 2-2: a 3×3 matrix is added to a 1×3 row vector and a 3×1 column vector, respectively, and in both cases the result is a 3×3 matrix. However, the elements in the two resulting matrices are different, because the way the elements of the row and column vectors are broadcasted to the shape of the larger array is different depending on the shape of the arrays, according to NumPy’s broadcasting rule.




Figure 2-2. Visualization of broadcasting of row and column vectors into the shape of a matrix. The highlighted elements represent true elements of the arrays, while the light blue elements describe the broadcasting of the elements of the array of smaller size

Arithmetic Operations

The standard arithmetic operations with NumPy arrays perform elementwise operations. Consider, for example, the addition, subtraction, multiplication, and division of equal-sized arrays.

```
In [132]: x = np.array([[1, 2], [3, 4]])
In [133]: y = np.array([[5, 6], [7, 8]])
In [134]: x + y
Out[134]: array([[ 6,  8],
                 [10, 12]])
In [135]: y - x
Out[135]: array([[4, 4],
                 [4, 4]])
In [136]: x * y
Out[136]: array([[ 5, 12],
                 [21, 32]])
In [137]: y / x
Out[137]: array([[ 5.,  3.],
                 [ 2.33333333,  2.]])
```

In operations between scalars and arrays, the scalar value is applied to each element in the array, as we could expect.

```
In [138]: x * 2
Out[138]: array([[2, 4],
                 [6, 8]])
In [139]: 2 ** x
Out[139]: array([[ 2,  4],
                 [ 8, 16]])
In [140]: y / 2
Out[140]: array([[ 2.5,  3.],
                 [ 3.5,  4.]])
In [141]: (y / 2).dtype
Out[141]: dtype('float64')
```

Note that the `dtype` of the resulting array for an expression can be promoted if the computation requires it, as shown in the preceding example with the division between an integer array and an integer scalar, which in that case resulted in an array with a `dtype` that is `np.float64`.

If an arithmetic operation is performed on arrays with incompatible sizes or shapes, a `ValueError` exception is raised.

```
In [142]: x = np.array([1, 2, 3, 4]).reshape(2, 2)
In [143]: z = np.array([1, 2, 3, 4])
In [144]: x / z
-----
ValueError                                Traceback (most recent call last)
<ipython-input-144-b88ced08eb6a> in <module>()
----> 1 x / z
ValueError: operands could not be broadcast together with shapes (2,2) (4,)
```

Here the array `x` has shape `(2, 2)` and the array `z` has shape `(4,)`, which cannot be broadcasted into a form that is compatible with `(2, 2)`. If, on the other hand, `z` has shape `(2,)`, `(2, 1)`, or `(1, 2)`, then it can be broadcasted to the shape `(2, 2)` by effectively repeating the array `z` along the axis with length 1. Let's first consider an example with an array `z` of shape `(1, 2)`, where the first axis (axis 0) has length 1.

```
In [145]: z = np.array([[2, 4]])
In [146]: z.shape
Out[146]: (1, 2)
```

Dividing the array `x` with array `z` is equivalent to dividing `x` with an array `zz` that is constructed by repeating (here using `np.concatenate`) the row vector `z` to obtain an array `zz` that has the same dimensions as `x`.

```
In [147]: x / z
Out[147]: array([[ 0.5,  0.5],
                 [ 1.5,  1. ]])
In [148]: zz = np.concatenate([z, z], axis=0)
In [149]: zz
Out[149]: array([[2, 4],
                 [2, 4]])
In [150]: x / zz
Out[150]: array([[ 0.5,  0.5],
                 [ 1.5,  1. ]])
```

Let's also consider the example in which the array `z` has shape `(2, 1)` and where the second axis (axis 1) has length 1.

```
In [151]: z = np.array([[2], [4]])
In [152]: z.shape
Out[152]: (2, 1)
```

In this case, dividing `x` with `z` is equivalent to dividing `x` with an array `zz` that is constructed by repeating the column vector `z` until a matrix with the same dimension as `x` is obtained.

```
In [153]: x / z
Out[153]: array([[ 0.5 ,  1.  ],
                 [ 0.75,  1. ]])
In [154]: zz = np.concatenate([z, z], axis=1)
In [155]: zz
Out[155]: array([[2, 2],
                 [4, 4]])
In [156]: x / zz
Out[156]: array([[ 0.5 ,  1.  ],
                 [ 0.75,  1. ]])
```

In summary, these examples show how arrays with shapes `(1, 2)` and `(2, 1)` are broadcasted to the shape `(2, 2)` of the array `x` when the operation `x / z` is performed. In both cases, the result of the operation `x / z` is the same as first repeating the smaller array `z` along its axis of length 1 to obtain a new array `zz` with the same shape as `x` and then performing the equal-sized array operation `x / zz`. However, the implementation of the broadcasting does not explicitly perform this expansion and the corresponding memory copies. But thinking of the array broadcasting in these terms can be helpful.

A summary of the operators for arithmetic operations with NumPy arrays is given in Table 2-6. These operators use the standard symbols used in Python. The result of an arithmetic operation with one or two arrays is a new independent array with its own data in the memory. Evaluating complicated arithmetic expressions might trigger many memory allocation and copy operations. Working with large arrays can lead to a large memory footprint and negatively impact performance. In such cases, in-place operation (see Table 2-6) can reduce the memory footprint and improve performance. As an example of in-place operators, consider the following two statements with the same effect.

```
In [157]: x = x + y
In [158]: x += y
```

Table 2-6. Operators for Elementwise Arithmetic Operation on NumPy Arrays

Operator	Operation
+, +=	Addition
-, -=	Subtraction
*, *=	Multiplication
/, /=	Division
//, //=	Integer division
**, **=	Exponentiation

The two expressions have the same effect, but in the first case, `x` is reassigned to a new array, while in the second case, the values of array `x` are updated in place. Extensive use of in-place operators tends to impair code readability, and in-place operators should be used only when necessary.

Elementwise Functions

In addition to arithmetic expressions using operators, NumPy provides vectorized functions for elementwise evaluation of many elementary mathematical functions and operations. Table 2-7 gives a summary of elementary mathematical functions in NumPy.³ Each of these functions takes a single array (of arbitrary dimension) as input and returns a new array of the same shape, where for each element, the function has been applied to the corresponding element in the input array. The data type of the output array is not necessarily the same as that of the input array.

³Note that this is not a complete list of the available elementwise functions in NumPy. See the NumPy reference documentations for comprehensive lists.

Table 2-7. Selection of NumPy Functions for Elementwise Elementary Mathematical Functions

NumPy Function	Description
<code>np.cos</code> , <code>np.sin</code> , <code>np.tan</code>	Trigonometric functions
<code>np.arccos</code> , <code>np.arcsin</code> , <code>np.arctan</code>	Inverse trigonometric functions
<code>np.cosh</code> , <code>np.sinh</code> , <code>np.tanh</code>	Hyperbolic trigonometric functions
<code>np.arccosh</code> , <code>np.arcsinh</code> , <code>np.arctanh</code>	Inverse hyperbolic trigonometric functions
<code>np.sqrt</code>	Square root
<code>np.exp</code>	Exponential
<code>np.log</code> , <code>np.log2</code> , <code>np.log10</code>	Logarithms of base e, 2, and 10, respectively

For example, the `np.sin` function (which takes only one argument) is used to compute the sine function for all values in the array.

```
In [159]: x = np.linspace(-1, 1, 11)
In [160]: x
Out[160]: array([-1. , -0.8, -0.6, -0.4, -0.2,  0. ,  0.2,  0.4,  0.6,  0.8,  1.])
In [161]: y = np.sin(np.pi * x)
In [162]: np.round(y, decimals=4)
Out[162]: array([-0., -0.5878, -0.9511, -0.9511, -0.5878,  0.,  0.5878,  0.9511,  0.9511,
  0.5878,  0.])
```

The `np.pi` constant and the `np.round` function were used to round the values of `y` to four decimals. Like the `np.sin` function, many elementary math functions take one input array and produce one output array. In contrast, many mathematical operator functions (see Table 2-8) operate on two input arrays and return one array.

```
In [163]: np.add(np.sin(x) ** 2, np.cos(x) ** 2)
Out[163]: array([ 1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.])
In [164]: np.sin(x) ** 2 + np.cos(x) ** 2
Out[164]: array([ 1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.])
```

Table 2-8. Summary of NumPy Functions for Elementwise Mathematical Operations

NumPy Function	Description
<code>np.add</code> , <code>np.subtract</code> , <code>np.multiply</code> , <code>np.divide</code>	Addition, subtraction, multiplication, and division of two NumPy arrays
<code>np.power</code>	Raises first input argument to the power of the second input argument (applied elementwise)
<code>np.remainder</code>	The remainder of the division
<code>np.reciprocal</code>	The reciprocal (inverse) of each element
<code>np.real</code> , <code>np.imag</code> , <code>np.conj</code>	The real part, imaginary part, and the complex conjugate of the elements in the input arrays
<code>np.sign</code> , <code>np.abs</code>	The sign and the absolute value
<code>np.floor</code> , <code>np.ceil</code> , <code>np rint</code>	Converts to integer values
<code>np.round</code>	Rounds to a given number of decimals

Note that in this example, `np.add` and the `+` operator are equivalent; for normal use, the operator should be used.

Occasionally, it is necessary to define new functions that operate on NumPy arrays on an element-by-element basis. A good way to implement such functions is to express it in terms of existing NumPy operators and expressions. But when this is not possible, the `np.vectorize` function can be a convenient tool. This function takes a nonvectorized function and returns a vectorized function. For example, consider the following implementation of the Heaviside step function, which works for scalar input.

```
In [165]: def heaviside(x):
...:     return 1 if x > 0 else 0
In [166]: heaviside(-1)
Out[166]: 0
In [167]: heaviside(1.5)
Out[167]: 1
```

However, unfortunately, this function does not work for NumPy array input.

```
In [168]: x = np.linspace(-5, 5, 11)
In [169]: heaviside(x)
...
ValueError: The truth value of an array with more than one element is ambiguous. Use a.any()
or a.all()
```

Using `np.vectorize` the scalar Heaviside function can be converted into a vectorized function that works with NumPy arrays as input.

```
In [170]: heaviside = np.vectorize(heaviside)
In [171]: heaviside(x)
Out[171]: array([0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
```

Although the function returned by `np.vectorize` works with arrays, it will be relatively slow since the original function must be called for each element in the array. There are much better ways to implement this function using arithmetic with Boolean-valued arrays, as discussed later in this chapter.

```
In [172]: def heaviside(x):
...:     return 1.0 * (x > 0)
```

Nonetheless, `np.vectorize` can often be a quick and convenient way to vectorize a function written for scalar input.

In addition to NumPy's functions for elementary mathematical functions, as summarized in Table 2-7, there are numerous functions in NumPy for mathematical operations. A summary of a selection of these functions is given in Table 2-8.

Aggregate Functions

NumPy provides another set of functions for calculating aggregates for NumPy arrays, which take an array as input and, by default, return a scalar as output. For example, statistics such as averages, standard deviations, and variances of the values in the input array and functions for calculating the sum and the product of elements in an array are all aggregate functions.

A summary of aggregate functions is given in Table 2-9. All of these functions are also available as methods in the `ndarray` class. For example, `np.mean(data)` and `data.mean()` in the following example are equivalent.

```
In [173]: data = np.random.normal(size=(15,15))
In [174]: np.mean(data)
Out[174]: -0.032423651106794522
In [175]: data.mean()
Out[175]: -0.032423651106794522
```

Table 2-9. NumPy Functions for Calculating Aggregates of NumPy Arrays

NumPy Function	Description
<code>np.mean</code>	The average of all values in the array
<code>np.std</code>	Standard deviation
<code>np.var</code>	Variance
<code>np.sum</code>	The sum of all elements
<code>np.prod</code>	The product of all elements
<code>np.cumsum</code>	The cumulative sum of all elements
<code>np.cumprod</code>	The cumulative product of all elements
<code>np.min</code> , <code>np.max</code>	The minimum/maximum value in an array
<code>np.argmin</code> , <code>np.argmax</code>	The index of the minimum/maximum value in an array
<code>np.all</code>	Returns True if all elements in the argument array are nonzero
<code>np.any</code>	Returns True if any of the elements in the argument array is nonzero

By default, the functions in Table 2-9 aggregate over the entire input array. Using the `axis` keyword argument with these functions, and their corresponding ndarray methods, it is possible to control which axis in the array aggregation is carried out over. The `axis` argument can be an integer, which specifies the axis to aggregate values over. In many cases, the `axis` argument can also be a tuple of integers, which specifies multiple axes to aggregate over. The following example demonstrates how calling the aggregate function `np.sum` on the array of shape (5, 10, 15) reduces the dimensionality of the array depending on the values of the `axis` argument.

```
In [176]: data = np.random.normal(size=(5, 10, 15))
In [177]: data.sum(axis=0).shape
Out[177]: (10, 15)
In [178]: data.sum(axis=(0, 2)).shape
Out[178]: (10,)
In [179]: data.sum()
Out[179]: -31.983793284860798
```

A visual illustration of how aggregation over all elements, over the first axis, and over the second axis of a 3×3 array is shown in Figure 2-3. In this example, the data array is filled with integers between 1 and 9.

```
In [180]: data = np.arange(1,10).reshape(3,3)
In [181]: data
Out[181]: array([[1, 2, 3],
                 [4, 5, 6],
                 [7, 8, 9]])
```




Figure 2-3. Illustration of array aggregation functions along all axes (left), the first axis (center), and the second axis (right) of a two-dimensional array of shape 3×3

We compute the aggregate sum of the entire array, over axis 0, and over axis 1, respectively.

```
In [182]: data.sum()
Out[182]: 45
In [183]: data.sum(axis=0)
Out[183]: array([12, 15, 18])
In [184]: data.sum(axis=1)
Out[184]: array([ 6, 15, 24])
```

Boolean Arrays and Conditional Expressions

When computing with NumPy arrays, there is often a need to compare elements in different arrays and perform conditional computations based on the results of such comparisons. Like with arithmetic operators, NumPy arrays can be used with the usual comparison operators, for example, `>`, `<`, `>=`, `<=`, `==`, and `!=`, and the comparisons are made on an element-by-element basis. The broadcasting rules also apply to comparison operators. If two operators have compatible shapes and sizes, the result of the comparison is a new array with Boolean values (with dtype as `np.bool`) that gives the result of the comparison for each element.

```
In [185]: a = np.array([1, 2, 3, 4])
In [186]: b = np.array([4, 3, 2, 1])
In [187]: a < b
Out[187]: array([ True,  True, False, False], dtype=bool)
```

To use the result of a comparison between arrays in, for example, an if statement, we need to aggregate the Boolean values of the resulting arrays in some suitable fashion, to obtain a single True or False value. A common use is to apply the `np.all` or `np.any` aggregation functions, depending on the situation.

```
In [188]: np.all(a < b)
Out[188]: False
In [189]: np.any(a < b)
Out[189]: True
In [190]: if np.all(a < b):
...:     print("All elements in a are smaller than their corresponding element in b")
...: elif np.any(a < b):
...:     print("Some elements in a are smaller than their corresponding element in b")
...: else:
...:     print("All elements in b are smaller than their corresponding element in a")
Some elements in a are smaller than their corresponding element in b
```

However, the advantage of Boolean-valued arrays is that they often make it possible to avoid conditional if statements altogether. Using Boolean-valued arrays in arithmetic expressions makes writing conditional computations in vectorized form possible. When appearing in an arithmetic expression with a scalar number, or another NumPy array with a numerical data type, a Boolean array is converted to a numerical-valued array with values 0 and 1 in place of False and True, respectively.

```
In [191]: x = np.array([-2, -1, 0, 1, 2])
In [192]: x > 0
Out[192]: array([False, False, False,  True,  True], dtype=bool)
In [193]: 1 * (x > 0)
Out[193]: array([0, 0, 0, 1, 1])
In [194]: x * (x > 0)
Out[194]: array([0, 0, 0, 1, 2])
```

This is a useful property for conditional computing, such as defining piecewise functions. For example, if we need to define a function describing a pulse of a given height, width, and position, we can implement this function by multiplying the height (a scalar variable) with two Boolean-valued arrays for the spatial extension of the pulse.

```
In [195]: def pulse(x, position, height, width):
...:     return height * (x >= position) * (x <= (position + width))
In [196]: x = np.linspace(-5, 5, 11)
In [197]: pulse(x, position=-2, height=1, width=5)
Out[197]: array([0, 0, 0, 1, 1, 1, 1, 1, 0, 0])
In [198]: pulse(x, position=1, height=1, width=5)
Out[198]: array([0, 0, 0, 0, 0, 0, 1, 1, 1, 1])
```

In this example, the expression `(x >= position) * (x <= (position + width))` is a multiplication of two Boolean-valued arrays. For this case, the multiplication operator acts as an elementwise AND operator. The pulse function could also be implemented using NumPy's `np.logical_and` function for elementwise AND operations.

```
In [199]: def pulse(x, position, height, width):
...:     return height * np.logical_and(x >= position, x <= (position + width))
```

There are functions for other logical operations, such as NOT, OR, and XOR, and functions for selectively picking values from different arrays depending on a given `np.where` condition, a list of `np.select` conditions, and an array of `np.choose` indices. Table 2-10 summarizes such functions, and the following examples demonstrate their basic usage. The `np.where` function selects elements from two arrays (second and third arguments), given a Boolean-valued array condition (the first argument). For elements where the condition is True, the corresponding values from the array given as second argument are selected. If the condition is False, values from the array given as third argument are selected.

```
In [200]: x = np.linspace(-4, 4, 9)
In [201]: np.where(x < 0, x**2, x**3)
Out[201]: array([ 16.,  9.,  4.,  1.,  0.,  1.,  8., 27., 64.])
```

Table 2-10. NumPy Functions for Conditional and Logical Expressions

Function	Description
<code>np.where</code>	Chooses values from two arrays depending on the value of a condition array
<code>np.choose</code>	Chooses values from a list of arrays depending on the values of a given index array
<code>np.select</code>	Chooses values from a list of arrays depending on a list of conditions
<code>np.nonzero</code>	Returns an array with indices of nonzero elements
<code>np.logical_and</code>	Performs an elementwise AND operation
<code>np.logical_or</code> , <code>np.logical_xor</code>	Elementwise OR/XOR operations
<code>np.logical_not</code>	Elementwise NOT operation (inverting)

The `np.select` function works similarly, but instead of a Boolean-valued condition array, it expects a list of Boolean-valued condition arrays and a corresponding list of value arrays.

```
In [202]: np.select([x < -1, x < 2, x >= 2],
...:                [x**2, x**3, x**4])
Out[202]: array([ 16.,  9.,  4., -1.,  0.,  1., 16., 81., 256.])
```

The `np.choose` takes as a first argument a list or an array with indices that determine from which array in a given list of arrays an element is picked.

```
In [203]: np.choose([0, 0, 0, 1, 1, 1, 2, 2, 2],
...:               [x**2, x**3, x**4])
Out[203]: array([ 16.,   9.,   4., -1.,   0.,   1.,  16.,  81., 256.])
```

The `np.nonzero` function returns a tuple of indices that can be used to index the array (e.g., the one that the condition was based on). This has the same results as indexing the array directly with `abs(x) > 2`. But it uses fancy indexing with the indices returned by `np.nonzero` rather than Boolean-valued array indexing.

```
In [204]: np.nonzero(abs(x) > 2)
Out[204]: (array([0, 1, 7, 8]),)
In [205]: x[np.nonzero(abs(x) > 2)]
Out[205]: array([-4., -3.,  3.,  4.])
In [206]: x[abs(x) > 2]
Out[206]: array([-4., -3.,  3.,  4.])
```

Set Operations

The Python language provides a convenient *set* data structure for managing unordered collections of unique objects. The NumPy array class `ndarray` can also describe such sets, and NumPy contains functions for operating on sets stored as NumPy arrays. These functions are summarized in Table 2-11. Using NumPy arrays to describe and operate on sets allows expressing certain operations in vectorized form. For example, testing if the values in a NumPy array are included in a set can be done using the `np.in1d` function, which tests for the existence of each element of its first argument in the array passed as the second argument. To see how this works, consider the following example: first, to ensure that a NumPy array is a proper set, we can use the `np.unique` function, which returns a new array with unique values.

```
In [207]: a = np.unique([1, 2, 3, 3])
In [208]: b = np.unique([2, 3, 4, 4, 5, 6, 5])
In [209]: np.in1d(a, b)
Out[209]: array([False,  True,  True], dtype=bool)
```

Table 2-11. NumPy Functions for Operating on Sets

Function	Description
<code>np.unique</code>	Creates a new array with unique elements, where each value only appears once
<code>np.in1d</code>	Tests for the existence of an array of elements in another array
<code>np.intersect1d</code>	Returns an array with elements that are contained in two given arrays
<code>np.setdiff1d</code>	Returns an array with elements that are contained in one, but not the other, of two given arrays
<code>np.union1d</code>	Returns an array with elements that are contained in either, or both, of two given arrays

Here, the existence of each element in `a` in the set `b` was tested, and the result is a Boolean-valued array. Note that we can use the `in` keyword to test for the existence of single elements in a set represented as a NumPy array.

```
In [210]: 1 in a
Out[210]: True
In [211]: 1 in b
Out[211]: False
```

To test if *a* is a subset of *b*, we can use the `np.in1d`, as in the previous example, together with the aggregation function `np.all` (or the corresponding `ndarray` method).

```
In [212]: np.all(np.in1d(a, b))
Out[212]: False
```

The fundamental set operations union (the set of elements included in either or both sets), intersection (elements included in both sets), and difference (elements included in one of the sets but not the other) is provided by `np.union1d`, `np.intersect1d`, and `np.setdiff1d`, respectively.

```
In [213]: np.union1d(a, b)
Out[213]: array([1, 2, 3, 4, 5, 6])
In [214]: np.intersect1d(a, b)
Out[214]: array([2, 3])
In [215]: np.setdiff1d(a, b)
Out[215]: array([1])
In [216]: np.setdiff1d(b, a)
Out[216]: array([4, 5, 6])
```

Operations on Arrays

In addition to elementwise and aggregation functions, some operations act on arrays as a whole and produce a transformed array of the same size. An example of this type of operation is the transpose, which flips the order of the axes of an array. For the special case of a two-dimensional array (i.e., a matrix, the transpose simply exchanges rows and columns).

```
In [217]: data = np.arange(9).reshape(3, 3)
In [218]: data
Out[218]: array([[0, 1, 2],
                 [3, 4, 5],
                 [6, 7, 8]])
In [219]: np.transpose(data)
Out[219]: array([[0, 3, 6],
                 [1, 4, 7],
                 [2, 5, 8]])
```

The transpose function `np.transpose` also exists as a method in `ndarray` and as the special method name `ndarray.T`. For an arbitrary *N*-dimensional array, the transpose operation reverses all the axes, as seen from the following example (note that the `shape` attribute is used here to display the number of values along each axis of the array).

```
In [220]: data = np.random.randn(1, 2, 3, 4, 5)
In [221]: data.shape
Out[221]: (1, 2, 3, 4, 5)
In [222]: data.T.shape
Out[222]: (5, 4, 3, 2, 1)
```


The `np.fliplr` (flip left-right) and `np.flipud` (flip up-down) functions perform operations that are similar to the transpose: they reshuffle the elements of an array so that the elements in rows (`np.fliplr`) or columns (`np.flipud`) are reversed, and the shape of the output array is the same as the input. The `np.rot90` function rotates the elements in the first two axes in an array by 90 degrees, and like the transpose function, it can change the shape of the array. Table 2-12 summarizes NumPy functions for common array operations.

Table 2-12. Summary of NumPy Functions for Array Operations

Function	Description
<code>np.transpose</code> , <code>np.ndarray.transpose</code> , <code>np.ndarray.T</code>	Transpose (reverse axes) an array.
<code>np.fliplr/np.flipud</code>	Reverse the elements in each row/column.
<code>np.rot90</code>	Rotate the elements along the first two axes by 90 degrees.
<code>np.sort</code> , <code>np.ndarray.sort</code>	Sort an array's elements along a specified axis (which defaults to the last axis of the array). The <code>np.ndarray</code> method <code>sort</code> performs the sorting in place, modifying the input array.

Matrix and Vector Operations

So far, I have discussed general N-dimensional arrays. One of the main applications of such arrays is to represent the mathematical concepts of vectors, matrices, and tensors. In this case, we also frequently need to calculate vector and matrix operations such as scalar (inner) products, dot (matrix) products, and tensor (outer) products. Table 2-13 summarizes NumPy's functions for matrix operations.

Table 2-13. Summary of NumPy Functions for Matrix Operations

NumPy Function	Description
<code>np.dot</code>	Matrix multiplication (dot product) between two given arrays representing vectors, arrays, or tensors
<code>np.inner</code>	Scalar multiplication (inner product) between two arrays representing vectors
<code>np.cross</code>	The cross product between two arrays that represent vectors
<code>np.tensordot</code>	Dot product along specified axes of multidimensional arrays
<code>np.outer</code>	Outer product (tensor product of vectors) between two arrays representing vectors
<code>np.kron</code>	Kronecker product (tensor product of matrices) between arrays representing matrices and higher-dimensional arrays
<code>np.einsum</code>	Evaluates Einstein's summation convention for multidimensional arrays

In NumPy, the `*` operator is used for elementwise multiplication. For two two-dimensional arrays `A` and `B`, the expression `A * B` does not compute a matrix product (in contrast to many other computing environments). Originally, there was no operator for denoting matrix multiplication in Python. Instead, the NumPy function `np.dot` has been used for this purpose, together with the corresponding method in the

ndarray class. However, the `@` operator has recently been introduced for matrix multiplication⁴, and while it is supported by NumPy, this operator is still not widely used. It is, therefore, still important to be familiar with the explicit methods in NumPy for matrix multiplication.

To compute the product of two matrices, A and B , of size $N \times M$ and $M \times P$, which results in a matrix of size $N \times P$, we can use `np.dot`.

```
In [223]: A = np.arange(1, 7).reshape(2, 3)
In [224]: A
Out[224]: array([[1, 2, 3],
                [4, 5, 6]])
In [225]: B = np.arange(1, 7).reshape(3, 2)
In [226]: B
Out[226]: array([[1, 2],
                [3, 4],
                [5, 6]])
In [227]: np.dot(A, B) # or equivalently: A @ B
Out[227]: array([[22, 28],
                [49, 64]])
In [228]: np.dot(B, A) # or equivalently: B @ A
Out[228]: array([[ 9, 12, 15],
                [19, 26, 33],
                [29, 40, 51]])
```

The `np.dot` function can also be used for matrix-vector multiplication (i.e., multiplication of a two-dimensional array, which represents a matrix, with a one-dimensional array representing a vector). The following is an example.

```
In [229]: A = np.arange(9).reshape(3, 3)
In [230]: A
Out[230]: array([[0, 1, 2],
                [3, 4, 5],
                [6, 7, 8]])
In [231]: x = np.arange(3)
In [232]: x
Out[232]: array([0, 1, 2])
In [233]: np.dot(A, x)
Out[233]: array([5, 14, 23])
```

In this example, `x` can be either a two-dimensional array of shape `(1, 3)` or a one-dimensional array with shape `(3,)`. In addition to the `np.dot` function, there is a corresponding `dot` method in `ndarray`, which is used in the following example.

```
In [234]: A.dot(x)
Out[234]: array([5, 14, 23])
```

⁴As of Python 3.5 the `@` symbol has been adopted for denoting matrix multiplication. See <http://legacy.python.org/dev/peps/pep-0465> for details.

Unfortunately, nontrivial matrix multiplication expressions can often become complex and hard to read when using either `np.dot` or `np.ndarray.dot`. For example, even a relatively simple matrix expression like the one for a similarity transform, $A' = BAB^{-1}$, must be represented with relatively cryptic nested expressions,⁵ such as

```
In [235]: A = np.random.rand(3,3)
In [236]: B = np.random.rand(3,3)
In [237]: Ap = np.dot(B, np.dot(A, np.linalg.inv(B)))
or
In [238]: Ap = B.dot(A.dot(np.linalg.inv(B)))
```

The `@` operator was introduced partly to address this type of readability problem. Also, since long before this operator was available, NumPy has provided an alternative to `ndarray` data structure named `matrix`, for which expressions like `A * B` are implemented as matrix multiplication. It also provides some convenient special attributes, like `matrix.I` for the inverse matrix and `matrix.H` for the complex conjugate transpose of a matrix. With instances of this `matrix` class, we can use a vastly more readable expression.

```
In [239]: A = np.matrix(A)
In [240]: B = np.matrix(B)
In [241]: Ap = B * A * B.I
```

This may seem like a practical compromise, but unfortunately, using the `matrix` class has a few disadvantages, and its use is often discouraged. The main objection against using `matrix` is that an expression like `A * B` is context-dependent: that is, it is not immediately clear if `A * B` denotes elementwise or matrix multiplication because it depends on the type of `A` and `B`, creating another code-readability problem. This can be a particularly relevant issue if `A` and `B` are user-supplied arguments to a function, in which case it would be necessary to cast all input arrays explicitly to `matrix` instances, using, for example, `np.asmatrix` or the `np.matrix` function (since there would be no guarantee that the user calls the function with arguments of type `matrix` rather than `ndarray`). The `np.asmatrix` function creates a view of the original array in the form of an `np.matrix` instance. This does not add much to computational costs, but explicitly casting arrays back and forth between `ndarray` and `matrix` does offset much of the benefits of the improved readability of matrix expressions. A related issue is that some functions that operate on arrays and matrices might not respect the type of the input and may return an `ndarray` even though it was called with an input argument of type `matrix`. This way, a matrix of type `matrix` might be unintentionally converted to `ndarray`, which would change the behavior of expressions like `A * B`. This behavior is not likely to occur when using NumPy's array and matrix functions, but it is likely to happen when using functions from other packages. However, despite all the arguments for not using matrix matrices too extensively, I believe that using `matrix` class instances for complicated matrix expressions is important. In these cases, it might be a good idea to explicitly cast arrays to matrices before the computation and explicitly cast the result back to the `ndarray` type, following the pattern.

```
In [242]: A = np.asmatrix(A)
In [243]: B = np.asmatrix(B)
In [244]: Ap = B * A * B.I
In [245]: Ap = np.asarray(Ap)
```

⁵With the new infix matrix multiplication operator, this same expression can be expressed as the considerably more readable: `Ap = B @ A @ np.linalg.inv(B)`.

The inner product (scalar product) between two arrays representing vectors can be computed using the `np.inner` function.

```
In [246]: np.inner(x, x)
Out[246]: 5
```

Or, equivalently, use `np.dot`.

```
In [247]: np.dot(x, x)
Out[247]: 5
```

The main difference is that `np.inner` expects two input arguments with the same dimension, while `np.dot` can take input vectors of shape $1 \times N$ and $N \times 1$, respectively.

```
In [248]: y = x[:, np.newaxis]
In [249]: y
Out[249]: array([[0],
                 [1],
                 [2]])
In [250]: np.dot(y.T, y)
Out[250]: array([[5]])
```

While the inner product maps two vectors to a scalar, the outer product performs the complementary operation of mapping two vectors to a matrix.

```
In [251]: x = np.array([1, 2, 3])
In [252]: np.outer(x, x)
Out[252]: array([[1, 2, 3],
                 [2, 4, 6],
                 [3, 6, 9]])
```

The outer product can also be calculated using the Kronecker product using the `np.kron` function, which, in contrast to `np.outer`, produces an output array of shape $(M \times P, N \times Q)$ if the input arrays have shapes (M, N) and (P, Q) , respectively. Thus, for two one-dimensional arrays of length M and P , the resulting array has shape $(M \times P,)$.

```
In [253]: np.kron(x, x)
Out[253]: array([1, 2, 3, 2, 4, 6, 3, 6, 9])
```

To obtain the result that corresponds to `np.outer(x, x)`, the input array `x` must be expanded to shape $(N, 1)$ and $(1, N)$, in the first and second argument to `np.kron`, respectively.

```
In [254]: np.kron(x[:, np.newaxis], x[np.newaxis, :])
Out[254]: array([[1, 2, 3],
                 [2, 4, 6],
                 [3, 6, 9]])
```

In general, while the `np.outer` function is primarily intended for vectors as input, the `np.kron` function can be used for computing tensor products of arrays of arbitrary dimension (but both inputs must have the same number of axes). For example, we can use the following to compute the tensor product of two 2×2 matrices.

```

In [255]: np.kron(np.ones((2,2)), np.identity(2))
Out[255]: array([[ 1.,  0.,  1.,  0.],
                  [ 0.,  1.,  0.,  1.],
                  [ 1.,  0.,  1.,  0.],
                  [ 0.,  1.,  0.,  1.]])
In [256]: np.kron(np.identity(2), np.ones((2,2)))
Out[256]: array([[ 1.,  1.,  0.,  0.],
                  [ 1.,  1.,  0.,  0.],
                  [ 0.,  0.,  1.,  1.],
                  [ 0.,  0.,  1.,  1.]])

```

When working with multidimensional arrays, it is often possible to express common array operations concisely using Einstein's summation convention, in which an implicit summation is assumed over each index that occurs multiple times in an expression. For example, the scalar product between two vectors x and y is compactly expressed as $x_n y_n$, and the matrix multiplication of two matrices, A and B , is expressed as $A_{mk} B_{kn}$. NumPy provides the `np.einsum` function for carrying out Einstein summations. Its first argument is an index expression, followed by an arbitrary number of arrays included in the expression. The index expression is a string with comma-separated indices, where each comma separates the indices of each array. Each array can have any number of indices. For example, the scalar product expression $x_n y_n$ can be evaluated with `np.einsum` using the index expression "n,n", that is using `np.einsum("n,n", x, y)`.

```

In [257]: x = np.array([1, 2, 3, 4])
In [258]: y = np.array([5, 6, 7, 8])
In [259]: np.einsum("n,n", x, y)
Out[259]: 70
In [260]: np.inner(x, y)
Out[260]: 70

```

Similarly, the matrix multiplication $A_{mk} B_{kn}$ can be evaluated using `np.einsum` and the index expression "mk,kn".

```

In [261]: A = np.arange(9).reshape(3, 3)
In [262]: B = A.T
In [263]: np.einsum("mk,kn", A, B)
Out[263]: array([[ 5, 14, 23],
                  [14, 50, 86],
                  [23, 86, 149]])
In [264]: np.alltrue(np.einsum("mk,kn", A, B) == np.dot(A, B))
Out[264]: True

```

The Einstein summation convention can be particularly convenient when dealing with multidimensional arrays since the index expression defining the operation makes it explicit which operation is carried out and which axes it is performed. An equivalent computation using, for example, `np.tensordot` might require giving the axes along which the dot product will be evaluated.

Summary

This chapter provided a brief introduction to array-based programming with the NumPy library that can serve as a reference for the upcoming chapters in this book. NumPy is a core library for computing with Python that provides a foundation for nearly all computational libraries for Python. To be familiar with the NumPy library and its patterns of application is a fundamental skill for using Python for scientific and

technical computing. I started with introducing NumPy's data structure for N-dimensional arrays—the `ndarray` object—and continued by discussing functions for creating and manipulating arrays, including indexing and slicing for extracting elements from arrays. I also discussed functions and operators for performing computations with `ndarray` objects, emphasizing vectorized expressions and operators for efficient computation with arrays. Throughout the rest of this book, we will see examples of higher-level libraries for specific fields in scientific computing that use the array framework provided by NumPy.

Further Reading

The NumPy library is the topic of several books, including the *Guide to NumPy*, by creator Travis E. Oliphant, available at <http://web.mit.edu/dvp/Public/numpybook.pdf>, and a series of books by Ivan Idris: *NumPy Beginner's Guide* (Packt, 2015), *NumPy Cookbook* (Packt, 2012), and *Learning NumPy Array* (Packt, 2014). NumPy is also covered in fair detail in *Python for Data Analysis* by Wes McKinney (O'Reilly, 2013).