

CHAPTER 3



Symbolic Computing

Symbolic computing is an entirely different paradigm than the numerical array-based computing introduced in the previous chapter. In symbolic computing software, also known as a computer algebra system (CAS), representations of mathematical objects and expressions are manipulated and transformed analytically. Symbolic computing is mainly about using computers to automate analytical computations that can be done by hand with pen and paper. However, by automating the bookkeeping and the manipulations of mathematical expressions using a computer algebra system, it is possible to take analytical computing much further than can realistically be done by hand. Symbolic computing is a great tool for checking and debugging analytical calculations done by hand, but more importantly, it enables carrying out analytical analysis that may not otherwise be possible.

Analytical and symbolic computing is a key part of the scientific and technical computing landscape. Even for problems that can only be solved numerically (which is common because analytical methods are not feasible in many practical problems), pushing the limits for what can be done analytically can make a big difference before resorting to numerical techniques. It can, for example, reduce the complexity or size of the numerical problem that finally needs to be solved. In other words, instead of tackling a problem in its original form directly using numerical methods, it may be possible to use analytical methods to simplify the problem first.

In the scientific Python environment, the main library for symbolic computing is SymPy (Symbolic Python). SymPy is written in Python and provides tools for a wide range of analytical and symbolic problems. This chapter explores how SymPy can be used for symbolic computing with Python.

SymPy The Symbolic Python (SymPy) library aims to provide a full-featured computer algebra system. However, SymPy is primarily a library rather than a full environment. This makes SymPy well-suited for integration in applications and computations using other Python libraries. At the time of writing, the latest version is 1.12. More information about SymPy is available at www.sympy.org and <https://github.com/sympy/sympy/wiki/Faq>.

Importing SymPy

The SymPy library provides the Python module named `sympy`. It is common to import all symbols from this module when working with SymPy, using `from sympy import *`. But in the interest of clarity and to avoid namespace conflicts between functions and variables from SymPy and other packages such NumPy and

SciPy (see later chapters), here we import the library in its entirety as `sympy`. For the rest of this book, assume that SymPy is imported this way.

```
In [1]: import sympy
In [2]: sympy.init_printing()
```

This also called the `sympy.init_printing` function, which configures SymPy's printing system to display nicely formatted renditions of mathematical expressions, as shown in examples later in this chapter. In the Jupyter Notebook, this sets up printing so that the MathJax JavaScript library renders SymPy expressions, and the results are displayed on the browser page of the notebook.

For convenience and readability of the example codes in this chapter, assume that the following frequently used symbols are explicitly imported from SymPy into the local namespace.

```
In [3]: from sympy import I, pi, oo
```

Caution NumPy and SymPy, as well as many other libraries, provide many functions and variables with the same name. But these symbols are rarely interchangeable. For example, `numpy.pi` is a numerical approximation of the mathematical symbol π , while `sympy.pi` is a symbolic representation of π . It is important not to mix them up and use, for instance, `numpy.pi` in place of `sympy.pi` when doing symbolic computations, or vice versa. The same holds for many fundamental mathematical functions, such as `numpy.sin` vs. `sympy.sin`. Therefore, it is important to consistently use namespaces when using more than one package in computing with Python.

Symbols

A core functionality in SymPy is to represent mathematical symbols as Python objects. In the SymPy library, the `sympy.Symbol` class can be used for this purpose. A `sympy.Symbol` instance has a name and set of attributes describing its properties and methods for querying those properties and for operating on the symbol object. A symbol by itself is not of much practical use, but symbols are used as nodes in trees to represent algebraic expressions (see next section). Among the first steps in analyzing a problem with SymPy is to create symbols for the various mathematical variables and quantities required to describe the problem.

The symbol name is a string, which optionally can contain LaTeX-like markup to make the symbol name display well in, for example, IPython's and Jupyter's rich display system. The name of a `sympy.Symbol` object is set when it is created. Symbols can be created in a few different ways in SymPy, for example, using `sympy.Symbol`, `sympy.symbols`, and `sympy.var`. Normally, it is desirable to associate SymPy symbols with Python variables with the same name or a name that closely corresponds to the symbol name. For example, to create a symbol named `x` and bind it to the Python variable with the same name, we can use the constructor of the `Symbol` class and pass a string containing the symbol name as the first argument.

```
In [4]: x = sympy.Symbol("x")
```

The variable `x` now represents an abstract mathematical symbol x , of which very little information is known by default. At this point, `x` could represent, for example, a real number, an integer, a complex number, a function, and many other possibilities. In many cases, it is sufficient to represent a mathematical symbol with this abstract, unspecified `Symbol` object. Yet, sometimes, it is necessary to give the SymPy library more hints about exactly what type of symbol a `Symbol` object represents. This may help SymPy to manipulate

or `simplify` analytical expressions more efficiently. We can add various assumptions that narrow down the possible properties of a symbol by adding optional keyword arguments to the symbol-creating functions, such as `Symbol`. Table 3-1 summarizes frequently used assumptions associated with a `Symbol` class instance. For example, if we have a mathematical variable `y` known as a real number, we can use the `real=True` keyword argument when creating the corresponding `symbol instance`. We can verify that SymPy recognizes that the symbol is real by using the `is_real` attribute of the `Symbol` class.

```
In [5]: y = sympy.Symbol("y", real=True)
In [6]: y.is_real
Out[6]: True
```

If, on the other hand, we were to use `is_real` to query the previously defined symbol `x`, which was not explicitly specified as real and can represent both real and nonreal variables, we get `None` as a result.

```
In [7]: x.is_real is None
Out[7]: True
```

Note that the `is_real` returns `True` if the symbol is known to be real, `False` if the symbol is known not to be real, and `None` if it is not known if the symbol is real. Other attributes (see Table 3-1) for querying assumptions on `Symbol` objects work in the same way. Consider the following example demonstrating a symbol for which the `is_real` attribute is `False`.

```
In [8]: sympy.Symbol("z", imaginary=True).is_real
Out[8]: False
```

Table 3-1. Selected Assumptions and Their Corresponding Keyword Arguments for `Symbol` Objects*

Assumption Keyword Arguments	Attributes	Description
<code>real, imaginary</code>	<code>is_real, is_imaginary</code>	Specify that a symbol represents a real or imaginary number.
<code>positive, negative</code>	<code>is_positive, is_negative</code>	Specify that a symbol is positive or negative.
<code>Integer</code>	<code>is_integer</code>	The symbol represents an integer.
<code>odd, even</code>	<code>is_odd, is_even</code>	The symbol represents an odd or even integer.
<code>Prime</code>	<code>is_prime</code>	The symbol is a prime number and also an integer.
<code>finite, infinite</code>	<code>is_finite, is_infinite</code>	The symbol represents a quantity that is finite or infinite.

*For a complete list, see the docstring for `sympy.Symbol`.

Among the assumptions in Table 3-1, the most important ones to explicitly specify when creating new symbols are `real` and `positive`. When applicable, adding these assumptions to symbols can frequently help SymPy simplify various expressions further than otherwise. Consider the following simple example.

```
In [9]: x = sympy.Symbol("x")
In [10]: y = sympy.Symbol("y", positive=True)
In [11]: sympy.sqrt(x ** 2)
```

```
Out[11]:  $\sqrt{x^2}$ 
In [12]: sympy.sqrt(y ** 2)
Out[12]: y
```

This created two symbols, x and y , and computed the square root of the square of that symbol using the SymPy function `sympy.sqrt`. If nothing is known about the symbol in the computation, then no simplification can be done. If, on the other hand, the symbol is known to be representing a positive number, then $\sqrt{y^2} = y$ and SymPy correctly recognizes this in the latter example.

When working with mathematical symbols that represent integers rather than real numbers, it is also useful to explicitly specify this when creating the corresponding SymPy symbols, using, for example, the `integer=True`, `even=True`, or `odd=True`, if applicable. This may also allow SymPy to analytically simplify certain expressions and function evaluations, such as in the following example.

```
In [13]: n1 = sympy.Symbol("n")
In [13]: n2 = sympy.Symbol("n", integer=True)
In [13]: n3 = sympy.Symbol("n", odd=True)
In [14]: sympy.cos(n1 * pi)
Out[14]: cos(pi*n)
In [15]: sympy.cos(n2 * pi)
Out[15]:  $(-1)^n$ 
In [16]: sympy.cos(n3 * pi)
Out[16]: -1
```

To formulate a nontrivial mathematical problem, it is often necessary to define a large number of symbols. Using `sympy.Symbol` to specify each symbol one by one may become tedious, and for convenience, SymPy provides a `sympy.symbols` function for creating multiple symbols in one function call. This function takes a comma-separated string of symbol names and an arbitrary set of keyword arguments (which apply to all the symbols), and it returns a tuple of newly created symbols. Using Python's tuple unpacking syntax and a call to `sympy.symbols` is a convenient way to create symbols.

```
In [17]: a, b, c = sympy.symbols("a, b, c", negative=True)
In [18]: d, e, f = sympy.symbols("d, e, f", positive=True)
```

Numbers

The purpose of creating Python objects for mathematical symbols is to use them to represent mathematical expressions. To be able to do this, we also need to represent other mathematical objects, such as numbers, functions, and constants. This section looks at SymPy's classes for representing number objects. These classes have many methods and attributes shared with instances of `sympy.Symbol`, which allows us to treat symbols and numbers on equal footing when representing expressions.

For example, in the previous section, we saw that `sympy.Symbol` instances have attributes for querying properties of symbol objects, such as `is_real`. We need to be able to use the same attributes for all types of objects, including for example numbers such as integers and floating-point numbers, when manipulating symbolic expressions in SymPy. For this reason, we cannot directly use the built-in Python objects for integers, `int`, floating-point numbers, `float`, and so on. Instead, SymPy provides the `sympy.Integer` and `sympy.Float` classes for representing integers and floating-point numbers within the SymPy framework. This distinction is important to be aware of when working with SymPy, but fortunately, we rarely need to be concerned with creating objects of type `sympy.Integer` and `sympy.Float` to represent specific numbers, since SymPy automatically promotes Python numbers to instances of these classes when they occur in

SymPy **expressions**. However, to demonstrate this difference between Python's built-in number types and the corresponding types in SymPy, the following example **explicitly** creates instances of `sympy.Integer` and `sympy.Float` and use some of their attributes to query their properties.

```
In [19]: i = sympy.Integer(19)
In [20]: type(i)
Out[20]: sympy.core.numbers.Integer
In [21]: i.is_Integer, i.is_real, i.is_odd
Out[21]: (True, True, True)
In [22]: f = sympy.Float(2.3)
In [23]: type(f)
Out[23]: sympy.core.numbers.Float
In [24]: f.is_Integer, f.is_real, f.is_odd
Out[24]: (False, True, False)
```

Tip We can cast instances of `sympy.Integer` and `sympy.Float` back to Python built-in types using the standard type casting `int(i)` and `float(f)`.

To create a SymPy representation of a number, or in general, an arbitrary expression, we can also use the `sympy.sympify` function. This function takes a wide range of inputs and derives a SymPy-compatible expression, and it **eliminates the need** for specifying **explicitly** what types of objects are to be created. For the simple case of number input, we can use the following.

```
In [25]: i, f = sympy.sympify(19), sympy.sympify(2.3)
In [26]: type(i), type(f)
Out[26]: (sympy.core.numbers.Integer, sympy.core.numbers.Float)
```

Integer

The previous section used the `Integer` class to represent integers. It's worth pointing out that there is a difference between a `Symbol` instance with the assumption `integer=True` and an instance of `Integer`. While the `Symbol` with `integer=True` represents **some** integer, the `Integer` instance represents a **specific** integer. For both cases, the `is_integer` attribute is `True`, but there is also an `is_Integer` attribute (note the capital I), which is **only** `True` for `Integer` instances. In general, attributes with names in the form `is_Name` indicate if the object is of type `Name`, and attributes with names in the form `is_name` indicate if the object is known to satisfy the condition `name`. Thus, there is also an `is_Symbol` attribute that is `True` for `Symbol` instances.

```
In [27]: n = sympy.Symbol("n", integer=True)
In [28]: n.is_integer, n.is_Integer, n.is_positive, n.is_Symbol
Out[28]: (True, False, None, True)
In [29]: i = sympy.Integer(19)
In [30]: i.is_integer, i.is_Integer, i.is_positive, i.is_Symbol
Out[30]: (True, True, True, False)
```

Integers in SymPy are of arbitrary precision, allowing for a limitless range without fixed lower or upper bounds. This sets them apart from representing integers with a specific bit size, such as in NumPy. Consequently, SymPy enables working with extremely large numbers, as demonstrated in the following examples.

Float

We have also already encountered the type `sympy.Float` in the previous sections. Like `Integer`, `Float` is arbitrary precision, unlike Python's built-in `float` type and the float types in NumPy. This means that a `Float` can represent a real number with an arbitrary number of decimals. When a `Float` instance is created using its constructor, there are two arguments: the first argument is a Python `float` or a string representing a floating-point number, and the second (optional) argument is the precision (number of significant decimal digits) of the `Float` object. For example, it is well known that the real number 0.3 cannot be represented exactly as a regular fixed bit-size floating-point number, and when printing 0.3 to 20 significant digits, it is displayed as 0.29999999999999888977698. The SymPy `Float` object can represent the real number 0.3 without the limitations of floating-point numbers.

```
In [33]: "%.25f" % 0.3 # Create a string representation with 25 decimals
Out[33]: '0.299999999999999888977698'
In [34]: sympy.Float(0.3, 25)
Out[34]: 0.299999999999999888977698
In [35]: sympy.Float('0.3', 25)
Out[35]: 0.3
```

However, note that to correctly represent 0.3 as a `Float` object, it is necessary to initialize it from a string '0.3 rather than the Python float 0.3, which already contains a floating-point error.

Rational

A rational number is a fraction p/q of two integers, the numerator p and the denominator q . SymPy represents this type of number using the `sympy.Rational` class. Rational numbers can be created explicitly using `sympy.Rational` and the numerator and denominator as arguments.

```
In [36]: sympy.Rational(11, 13)
Out[36]: 11
           -
          13
```

Or they can be a result of a simplification carried out by SymPy. In either case, arithmetic operations between rational and integers remain rational.

```
In [37]: r1 = sympy.Rational(2, 3)
In [38]: r2 = sympy.Rational(4, 5)
In [39]: r1 * r2
Out[39]: 8
                  -
                 15
```

In [40]: `r1 / r2`

Out[40]: $\frac{5}{6}$

Constants and Special Symbols

Sympy provides predefined symbols for mathematical constants and special objects, such as the imaginary unit i and infinity. These are summarized in Table 3-2, together with their corresponding symbols in SymPy. Note that the imaginary unit is written as I in SymPy.

Table 3-2. Selected Mathematical Constants and Special Symbols and Their Corresponding Symbols in SymPy

Mathematical Symbol	Sympy Symbol	Description
π	<code>sympy.pi</code>	Ratio of the circumference to the diameter of a circle
E	<code>sympy.E</code>	The base of the natural logarithm, $e = \exp(1)$
γ	<code>sympy.EulerGamma</code>	Euler's constant
I	<code>sympy.I</code>	The imaginary unit
∞	<code>sympy.oo</code>	Infinity

Functions

In SymPy, objects that represent functions can be created with `sympy.Function`. Like `Symbol`, the `Function` object takes a name as the first argument. SymPy distinguishes between defined and undefined functions and between applied and unapplied functions. Creating a function with `Function` results in an undefined (abstract) and unapplied function, which has a name but cannot be evaluated because its expression, or body, is not defined. Such a function can represent an arbitrary function of arbitrary numbers of input variables since it also has not yet been applied to any particular symbols. An unapplied function can be applied to a set of input symbols that represent the domain of the function by calling the function instance with those symbols as arguments.¹ The result is still an unevaluated function but one that has been applied to the specified input variables and, therefore, has a set of dependent variables. As an example of these concepts, consider the following code listing where we create an undefined function f , which we apply to the symbol x , and another g function, which we directly apply to the set of symbols x , y , z .

```
In [41]: x, y, z = sympy.symbols("x, y, z")
In [42]: f = sympy.Function("f")
In [43]: type(f)
Out[43]: sympy.core.function.UndefinedFunction
In [44]: f(x)
Out[44]: f(x)
In [45]: g = sympy.Function("g")(x, y, z)
In [46]: g
```

¹Here it is important to keep in mind the distinction between a Python function, or callable Python object such as `sympy.Function`, and the symbolic function that a `sympy.Function` class instance represents.

```
Out[46]: g(x,y,z)
In [47]: g.free_symbols
Out[47]: {x,y,z}
```

This used the `free_symbols` property, which returns a set of unique symbols contained in an expression (in this case, the applied undefined `g` function), to demonstrate that an applied function is associated with a specific set of input symbols. This will be important later in this chapter, for example, when considering derivatives of abstract functions. One important application of undefined functions is for specifying differential equations or, in other words, when an equation for the function is known, but the function itself is unknown.

In contrast to undefined functions, a defined function has a specific implementation and can be numerically evaluated for all valid input parameters. It is possible to define this type of function, for example, by subclassing `sympy.Function`, but in most cases, it is sufficient to use the mathematical functions provided by SymPy. There are built-in functions for many standard mathematical functions in the global SymPy namespace. See the module documentation for `sympy.functions.elementary`, `sympy.functions.combinatorial`, and `sympy.functions.special` and their submodules for comprehensive lists of the numerous available functions, using the Python help function. For example, the SymPy function for the sine function is available as `sympy.sin` (with our import convention). Note that this is not a function in the Python sense of the word (it is, in fact, a subclass of `sympy.Function`), and it represents an unevaluated sine function that can be applied to a numerical value, a symbol, or an expression.

```
In [48]: sympy.sin
Out[48]: sympy.functions.elementary.trigonometric.sin
In [49]: sympy.sin(x)
Out[49]: sin(x)
In [50]: sympy.sin(pi * 1.5)
Out[50]: -1
```

When applied to an abstract symbol, such as `x`, the `sin` function remains unevaluated. But when possible, it is evaluated to a numerical value, for example, when applied to a number or, in some cases, when applied to expressions with certain properties, as in the following example.

```
In [51]: n = sympy.Symbol("n", integer=True)
In [52]: sympy.sin(pi * n)
Out[52]: 0
```

A third type of function in SymPy is lambda functions, or anonymous functions, which do not have names associated with them but have a specific function body that can be evaluated. Lambda functions can be created with `sympy.Lambda`.

```
In [53]: h = sympy.Lambda(x, x**2)
In [54]: h
Out[54]: (x  $\mapsto$  x2)
In [55]: h(5)
Out[55]: 25
In [56]: h(1 + x)
Out[56]: (1 + x)2
```

Expressions

The symbols introduced in the previous sections are the fundamental building blocks required to express mathematical expressions. In SymPy, mathematical expressions are represented as trees, where leaves are symbols and nodes are class instances that represent mathematical operations. These classes include Add, Mul, and Pow for basic arithmetic operators and Sum, Product, Integral, and Derivative for analytical mathematical operations. In addition, there are many other classes for mathematical operations, which are demonstrated in examples later in this chapter.

Consider, for example, the mathematical expression $1 + 2x^2 + 3x^3$. To represent this in SymPy, we only need to create the symbol x and write the expression as Python code.

```
In [54]: x = sympy.Symbol("x")
In [55]: expr = 1 + 2 * x**2 + 3 * x**3
In [56]: expr
Out[56]: 3x3 + 2x2 + 1
```

Here `expr` is an instance of `Add`, with the subexpressions `1`, `2*x**2`, and `3*x**3`. The entire expression tree for `expr` is visualized in Figure 3-1. Note that we do not need to explicitly construct the expression tree, since it is automatically built up from the expression with symbols and operators. Nevertheless, to understand how SymPy works, it is important to know how expressions are represented.

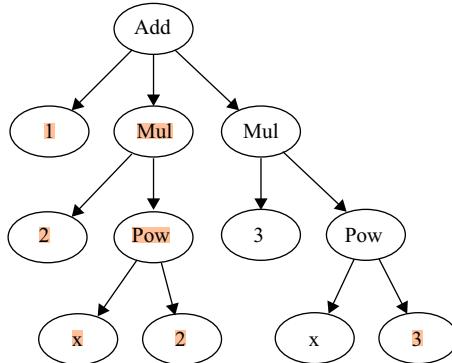


Figure 3-1. Visualization of the expression tree for $1 + 2x^2 + 3x^3$

The expression tree can be traversed explicitly using the `args` attribute, which all SymPy operations and symbols provide. For an operator, the `args` attribute is a tuple of subexpressions combined with the rule implemented by the operator class. For symbols, the `args` attribute is an empty tuple, which signifies that it is a leaf in the expression tree. The following example demonstrates how the expression tree can be explicitly accessed.

```
In [57]: expr.args
Out[57]: (1, 2x2, 3x3)
In [58]: expr.args[1]
Out[58]: 2x2
In [59]: expr.args[1].args[1]
Out[59]: x2
```

```
In [60]: expr.args[1].args[1].args[0]
Out[60]: x
In [61]: expr.args[1].args[1].args[0].args
Out[61]: ()
```

In the basic use of SymPy, explicitly manipulating expression trees is rarely necessary. But when the methods for manipulating expressions that are introduced in the following section are not sufficient, it is useful to be able to implement functions of your own that traverse and manipulate the expression tree using the `args` attribute.

Manipulating Expressions

Manipulating expression trees is one of the main jobs for SymPy, and numerous functions are provided for different types of transformations. The general idea is that expression trees can be transformed between mathematically equivalent forms using simplification and rewrite functions. These functions generally do not change the expressions passed to the functions but rather create a new expression corresponding to the modified expression. Expressions in SymPy should thus be considered immutable objects (that cannot or should not be changed in place). All the functions in this section treat SymPy expressions as immutable objects and return new expression trees rather than modify expressions in place.

Simplification

The most desirable manipulation of a mathematical expression is to simplify it. This is perhaps also the most ambiguous operation since it is nontrivial to determine algorithmically if one expression appears simpler than another to a human being, and in general, it is not obvious which methods should be employed to arrive at a simpler expression. Nonetheless, black-box simplification is an important part of any CAS, and SymPy includes the `sympy.simplify` function, which attempts to simplify a given expression using a variety of methods and approaches. As illustrated in the following example, the simplification function can also be invoked through the `simplify` method.

```
In [67]: expr = 2 * (x**2 - x) - x * (x + 1)
In [68]: expr
Out[68]: 2x2 - x(x+1)-2x
In [69]: sympy.simplify(expr)
Out[69]: x(x-3)
In [70]: expr.simplify()
Out[70]: x(x-3)
In [71]: expr
Out[71]: 2x2 - x(x+1)-2x
```

Note that here, both `sympy.simplify(expr)` and `expr.simplify()` return new expression trees, and leave the expression `expr` untouched, as mentioned earlier. In this example, the expression `expr` can be simplified by expanding the products, canceling terms, and then factoring the expression again. In general, `sympy.simplify` attempts a variety of different strategies and also simplifies, for example, trigonometric and power expressions. This is exemplified in the following example.

```
In [72]: expr = 2 * sympy.cos(x) * sympy.sin(x)
In [73]: expr
Out[73]: 2 sin(x)cos(x)
```

```
In [74]: sympy.simplify(expr)
Out[74]: sin(2x)
```

The following is another example.

```
In [75]: expr = sympy.exp(x) * sympy.exp(y)
In [76]: expr
Out[76]: exp(x)exp(y)
In [77]: sympy.simplify(expr)
Out[77]: exp(x+y)
```

Each specific type of simplification can also be carried out with more specialized functions, such as `sympy.trigsimp` and `sympy.powsimp`, for trigonometric and power simplifications, respectively. These functions only perform the simplification that their names indicate and leave other parts of an expression in their original form. A summary of simplification functions is given in Table 3-3. When the exact simplification steps are known, it is generally better to rely on the more specific simplification functions since their actions are more well-defined and less likely to change in future versions of SymPy. The `sympy.simplify` function, on the other hand, relies on heuristic approaches that may change in the future and, as a consequence, produce different results for a particular input expression.

Table 3-3. Summary of Selected SymPy Functions for Simplifying Expressions

Function	Description
<code>sympy.simplify</code>	Attempt various methods and approaches to obtain a simpler form of a given expression.
<code>sympy.trigsimp</code>	Attempt to simplify an expression using trigonometric identities.
<code>sympy.powsimp</code>	Attempt to simplify an expression using laws of powers.
<code>sympy.compsimp</code>	Simplify combinatorial expressions.
<code>sympy.ratsimp</code>	Simplify an expression by writing on a common denominator.

Expand

When the black-box simplification provided by `sympy.simplify` does not produce satisfying results, it is often possible to make progress by manually guiding SymPy using more specific algebraic operations. An important tool in this process is to expand expression in various ways. The `sympy.expand` function performs a variety of expansions depending on the values of optional keyword arguments. By default, the function distributes products over additions into a fully expanded expression. For example, a product of the type $(x + 1)(x + 2)$ can be expanded to $x^2 + 3x + 2$ using the following.

```
In [78]: expr = (x + 1) * (x + 2)
In [79]: sympy.expand(expr)
Out[79]: x2 + 3x + 2
```

Some of the available keyword arguments are `mul=True` for expanding products (as in the preceding example) and `trig=True` for trigonometric expansions.

```
In [80]: sympy.sin(x + y).expand(trig=True)
Out[80]: sin(x)cos(y) + sin(y)cos(x)
```

`log=True` is for expanding logarithms.

```
In [81]: a, b = sympy.symbols("a, b", positive=True)
In [82]: sympy.log(a * b).expand(log=True)
Out[82]: log(a) + log(b)
```

`complex=True` is for separating real and imaginary parts of an expression.

```
In [83]: sympy.exp(I*a + b).expand(complex=True)
Out[83]: ieb sin(a) + eb cos(a)
```

And `power_base=True` and `power_exp=True` for expanding the base and the exponent of a power expression, respectively.

```
In [84]: sympy.expand((a * b)**x, power_base=True)
Out[84]: axbx
In [85]: sympy.expand((a-b)**x).expand(power_exp=True)
Out[85]: eaxe-bx
```

Calling the `sympy.expand` function with these keyword arguments set to `True` is equivalent to calling the more specific functions `sympy.expand_mul`, `sympy.expand_trig`, `sympy.expand_log`, `sympy.expand_complex`, `sympy.expand_power_base`, and `sympy.expand_power_exp`, respectively, but an advantage of the `sympy.expand` function is that several types of expansions can be performed in a single function call.

Factor, Collect, and Combine

A common use pattern for the `sympy.expand` function is to expand an expression, let SymPy cancel terms or factors, and then factor or combine the expression again. The `sympy.factor` function attempts to factor an expression as far as possible and is in some sense the opposite to `sympy.expand` with `mul=True`. It can be used to factor algebraic expressions, such as the following.

```
In [86]: sympy.factor(x**2 - 1)
Out[86]: (x - 1)(x + 1)
In [87]: sympy.factor(x * sympy.cos(y) + sympy.sin(z) * x)
Out[87]: x(sin(x) + cos(y))
```

The inverse of the other types of expansions in the previous section can be carried out using `sympy.trigsimp`, `sympy.powsimp`, and `sympy.logcombine`, as shown in the following example.

```
In [90]: sympy.logcombine(sympy.log(a) - sympy.log(b))
Out[90]: log( $\frac{a}{b}$ )
```

When working with mathematical expressions, fine-grained control over factoring is often necessary. The SymPy function `sympy.collect` factors terms that contain a given symbol or list of symbols. For example, $x + y + xyz$ cannot be completely factorized, but we can partially factor terms containing x or y .

```
In [89]: expr = x + y + x * y * z
In [90]: expr.collect(x)
Out[90]: x(yz + 1) + y
In [91]: expr.collect(y)
Out[91]: x + y(xz + 1)
```

By passing a list of symbols or expressions to the `sympy.collect` function, or to the corresponding `collect` method, we can collect multiple symbols in one function call. Also, when using the `collect` method, which returns the new expression, it is possible to chain multiple methods calls in the following way.

```
In [93]: expr = sympy.cos(x + y) + sympy.sin(x - y)
In [94]: expr.expand(trig=True).collect(
    ...:     [sympy.cos(x), sympy.sin(x)])
    ...: ).collect(
    ...:     sympy.cos(y) - sympy.sin(y)
    ...: )
Out[95]: (sin(x) + cos(x))(-sin(y) + cos(y))
```

Apart, Together, and Cancel

The final type of mathematical simplification to consider here is the rewriting of fractions. The functions `sympy.apart` and `sympy.together`, which rewrite a fraction as a partial fraction and combine partial fractions to a single fraction, can be used in the following way.

```
In [95]: sympy.apart(1/(x**2 + 3*x + 2), x)
Out[95]: - $\frac{1}{x+2} + \frac{1}{x+1}$ 

In [96]: sympy.together(1 / (y * x + y) + 1 / (1+x))
Out[96]:  $\frac{y+1}{y(x+1)}$ 

In [97]: sympy.cancel(y / (y * x + y))
Out[97]:  $\frac{1}{x+1}$ 
```

The first example used `sympy.apart` to rewrite the expression $(x^2+3x+2)^{-1}$ as the partial fraction $-\frac{1}{x+2} + \frac{1}{x+1}$ and used `sympy.together` to combine the sum of fractions $1/(yx+y) + 1/(1+x)$ into an expression in the form of a single fraction. This example also used the `sympy.cancel` function to cancel shared factors between numerator and the denominator in the expression $y/(yx+y)$.

Substitutions

The previous sections have been concerned with rewriting expressions using various mathematical identities. Another frequently used form of manipulation of mathematical expressions is the substitution of symbols or subexpressions within an expression. For example, we may want to perform a variable substitution and replace the variable x with y or replace a symbol with another expression. In SymPy, there are two methods for carrying out substitutions: `subs` and `replace`. Usually, `subs` is the most suitable alternative, but in some cases `replace` provides a more powerful tool, which, for example, can make replacements based on wildcard expressions (see docstring for `sympy.Symbol.replace` for details).

In the most basic use of `subs`, the method is invoked from an expression, and the symbol or expression that is to be replaced (`x`) is given as the first argument, and the new symbol or the expression (`y`) is given as the second argument. The result is that all occurrences of `x` in the expression are replaced with `y`.

```
In [98]: (x + y).subs(x, y)
Out[98]: 2y
In [99]: sympy.sin(x * sympy.exp(x)).subs(x, y)
Out[99]: sin(yey)
```

Instead of chaining multiple `subs` calls when several substitutions are required, we can alternatively pass a dictionary that maps old symbols or expressions to new symbols or expressions as the first and only argument to `subs`.

```
In [100]: sympy.sin(x * z).subs({z: sympy.exp(y), x: y, sympy.sin: sympy.cos})
Out[100]: cos(yey)
```

A common application of the `subs` method is to substitute numerical values instead of symbols for numerical evaluation (see the following section for more details). A convenient way of doing this is to define a dictionary that translates the symbols to numerical values and pass this dictionary as the argument to the `subs` method. For example, consider the following.

```
In [101]: expr = x * y + z**2 *x
In [102]: values = {x: 1.25, y: 0.4, z: 3.2}
In [103]: expr.subs(values)
Out[103]: 13.3
```

Numerical Evaluation

Even when working with symbolic mathematics, it is almost invariably sooner or later required to evaluate the symbolic expressions numerically, for example, when producing plots or concrete numerical results. A SymPy expression can be evaluated using either the `sympy.N` function or the `evalf` method of SymPy expression instances.

```
In [104]: sympy.N(1 + pi)
Out[104]: 4.14159265358979
In [105]: sympy.N(pi, 50)
Out[105]: 3.1415926535897932384626433832795028841971693993751
In [106]: (x + 1/pi).evalf(10)
Out[106]: x + 0.3183098862
```

Both `sympy.N` and the `evalf` method take an optional argument that specifies the number of significant digits to which the expression is to be evaluated, as shown in the previous example where SymPy's arbitrary precision float capabilities were leveraged to evaluate the value of π up to 50 digits.

When we need to evaluate an expression numerically for a range of input values, we could loop over the values and perform successive `evalf` calls, as shown in the following example.

```
In [114]: expr = sympy.sin(pi * x * sympy.exp(x))
In [115]: [expr.subs(x, xx).evalf(3) for xx in range(0, 10)]
Out[115]: [0, 0.774, 0.642, 0.722, 0.944, 0.205, 0.974, 0.977, -0.870, -0.695]
```

However, this method is rather slow, and SymPy provides a more efficient method for doing this operation using the `sympy.lambdify` function. This function takes a set of free symbols and an expression as arguments and generates a function that efficiently evaluates the numerical value of the expression. The produced function takes the same number of arguments as the number of free symbols passed as the first argument to `sympy.lambdify`.

```
In [109]: expr_func = sympy.lambdify(x, expr)
In [110]: expr_func(1.0)
Out[110]: 0.773942685266709
```

Note that the `expr_func` function expects numerical (scalar) values as arguments, so we cannot, for example, pass a symbol as an argument to this function; it is strictly for numerical evaluation. The `expr_func` created in the previous example is a scalar function and is not directly compatible with vectorized input in the form of NumPy arrays, as discussed in Chapter 2. However, SymPy is also able to generate functions that are NumPy-array aware: passing 'numpy' as the optional third argument to `sympy.lambdify` creates a vectorized function that accepts NumPy arrays as input. This is often an efficient way to numerically evaluate symbolic expressions² for a large number of input parameters. The following code exemplifies how the SymPy expression `expr` is converted into a NumPy-array aware vectorized function that can be efficiently evaluated:

```
In [111]: expr_func = sympy.lambdify(x, expr, 'numpy')
In [112]: import numpy as np
In [113]: xvalues = np.arange(0, 10)
In [114]: expr_func(xvalues)
Out[114]: array([ 0.          ,  0.77394269,  0.64198244,  0.72163867,  0.94361635,
       0.20523391,  0.97398794,  0.97734066, -0.87034418, -0.69512687])
```

This method for generating data from SymPy expressions is useful for plotting and many other data-oriented applications.

Calculus

So far, we have looked at how to represent mathematical expressions in SymPy and how to perform basic simplification and transformation of such expressions. With this framework in place, we are now ready to explore symbolic calculus, or analysis, which is a cornerstone in applied mathematics and has many applications throughout science and engineering. The central concept in calculus is the change of functions as input variables are varied, as quantified with derivatives and differentials, and accumulations of functions over ranges of input, as quantified by integrals. This section looks at how to compute derivatives and integrals of functions in SymPy.

²See also the `ufuncify` from the `sympy.utilities.autowrap` module and the `aesara_function` from the `sympy.printing.aesaracode` module. These provide similar functionality as `sympy.lambdify`, but use different computational backends.

Derivatives

A function's derivative describes its change rate at a given point. In SymPy, we can calculate the derivative of a function using `sympy.diff` or the `diff` method of SymPy expression instances. The argument to these functions is a symbol, or several symbols, with respect to which the function or the expression is to be derived. To represent the first-order derivative of an abstract function $f(x)$ with respect to x , we can do the following.

```
In [119]: f = sympy.Function('f')(x)
In [120]: sympy.diff(f, x) # equivalent to f.diff(x)
```

```
Out[120]:  $\frac{d}{dx} f(x)$ 
```

To represent higher-order derivatives, we must repeat the symbol x in the argument list in the call to `sympy.diff` or specify an integer as an argument following a symbol, which defines the number of times the expression should be derived with respect to that symbol.

```
In [117]: sympy.diff(f, x, x)
```

```
Out[117]:  $\frac{d^2}{dx^2} f(x)$ 
```

```
In [118]: sympy.diff(f, x, 3) # equivalent to sympy.diff(f, x, x, x)
```

```
Out[118]:  $\frac{d^3}{dx^3} f(x)$ 
```

This method is readily extended to multivariate functions.

```
In [119]: g = sympy.Function('g')(x, y)
In [120]: g.diff(x, y) # equivalent to sympy.diff(g, x, y)
```

```
Out[120]:  $\frac{\partial^2}{\partial x \partial y} g(x, y)$ 
```

```
In [121]: g.diff(x, 3, y, 2) # equivalent to sympy.diff(g, x, x, x, y, y)
```

```
Out[121]:  $\frac{\partial^5}{\partial x^3 \partial y^2} g(x, y)$ 
```

These examples so far only involve formal derivatives of undefined functions. Naturally, we can also evaluate the derivatives of defined functions and expressions, which result in new expressions that correspond to the evaluated derivatives. For example, `sympy.diff` lets us easily evaluate derivatives of arbitrary mathematical expressions, such as polynomials.

```
In [122]: expr = x**4 + x**3 + x**2 + x + 1
In [123]: expr.diff(x)
Out[123]:  $4x^3 + 3x^2 + 2x + 1$ 
In [124]: expr.diff(x, x)
Out[124]:  $2(6x^2 + 3x + 1)$ 
In [125]: expr = (x + 1)**3 * y ** 2 * (z - 1)
In [126]: expr.diff(x, y, z)
Out[126]:  $6y(x + 1)^2$ 
```

We can also easily evaluate trigonometric and other more complicated mathematical expressions.

```
In [127]: expr = sympy.sin(x * y) * sympy.cos(x / 2)
In [128]: expr.diff(x)
```

```
Out[128]:  $y \cos\left(\frac{x}{2}\right) \cos(xy) - \frac{1}{2} \sin\left(\frac{x}{2}\right) \sin(xy)$ 
```

```
In [129]: expr = sympy.functions.special.polynomials.hermite(x, 0)
In [130]: expr.diff(x).doit()
```

```
Out[130]:  $\frac{2^x \sqrt{\pi} \operatorname{polygamma}\left(0, -\frac{x}{2} + \frac{1}{2}\right)}{2\Gamma\left(-\frac{x}{2} + \frac{1}{2}\right)} + \frac{2^x \sqrt{\pi} \log(2)}{\Gamma\left(-\frac{x}{2} + \frac{1}{2}\right)}$ 
```

Derivatives are usually relatively easy to compute, and `sympy.diff` should be able to evaluate the derivative of most standard mathematical functions defined in SymPy.

Note that in these examples, calling `sympy.diff` on an expression directly results in a new expression. If we want to symbolically represent the derivative of a definite expression, we can create an instance of the `sympy.Derivative` class, passing the expression as the first argument, followed by the symbols with respect to which the derivative is to be computed.

```
In [131]: d = sympy.Derivative(sympy.exp(sympy.cos(x)), x)
In [132]: d
```

```
Out[132]:  $\frac{d}{dx} e^{\cos(x)}$ 
```

This formal representation of a derivative can then be evaluated by calling the `doit` method on the `sympy.Derivative` instance.

```
In [133]: d.doit()
Out[133]:  $-e^{\cos(x)} \sin(x)$ 
```

This pattern of delayed evaluation is reoccurring throughout SymPy, and full control of when a formal expression is evaluated to a specific result is useful in many situations, particularly with expressions that can be simplified or manipulated while represented as a formal expression rather than after it has been evaluated.

Integrals

SymPy evaluates integrals using the `sympy.integrate` function, and formal integrals can be represented using `sympy.Integral` (which, as in the case with `sympy.Derivative`, can be explicitly evaluated by calling the `doit` method). Integrals come in two basic forms: definite and indefinite, where a definite integral has specified integration limits and can be interpreted as an area or volume. In contrast, an indefinite integral does not have integration limits and denotes the antiderivative (inverse of the derivative of a function). SymPy handles both indefinite and definite integrals using the `sympy.integrate` function.

If the `sympy.integrate` function is called with only an expression as an argument, the indefinite integral is computed. On the other hand, a definite integral is computed if the `sympy.integrate` function additionally is passed a tuple in the form `(x, a, b)`, where `x` is the integration variable and `a` and `b` are the integration limits. For a single-variable $f(x)$ function, the indefinite and definite integrals are computed using the following.

```
In [135]: a, b, x, y = sympy.symbols("a, b, x, y")
...: f = sympy.Function("f")(x)
In [136]: sympy.integrate(f)
Out[136]:  $\int f(x)dx$ 
In [137]: sympy.integrate(f, (x, a, b))
Out[137]:  $\int_a^b f(x)dx$ 
```

When these methods are applied to explicit functions, the integrals are evaluated accordingly.

```
In [138]: sympy.integrate(sympy.sin(x))
Out[138]: -cos(x)
In [139]: sympy.integrate(sympy.sin(x), (x, a, b))
Out[139]: cos(a) - cos(b)
```

Definite integrals can also include limits that extend from negative infinity or to positive infinite, using SymPy's symbol for infinity `oo`.

```
In [139]: sympy.integrate(sympy.exp(-x**2), (x, 0, oo))
Out[139]:  $\frac{\sqrt{\pi}}{2}$ 
In [140]: a, b, c = sympy.symbols("a, b, c", positive=True)
In [141]: sympy.integrate(a * sympy.exp(-((x-b)/c)**2), (x, -oo, oo))
Out[141]:  $\sqrt{\pi}ac$ 
```

Computing integrals symbolically is generally a difficult problem, and SymPy cannot give symbolic results for any integral we can come up with. When SymPy fails to evaluate an integral, an instance of `sympy.Integral`, representing the formal integral, is returned instead.

```
In [142]: sympy.integrate(sympy.sin(x * sympy.cos(x)))
Out[142]:  $\int \sin(x \cos(x))dx$ 
```

Multivariable expressions can also be integrated with `sympy.integrate`. In an indefinite integral of a multivariable expression, the integration variable must be specified explicitly.

```
In [140]: expr = sympy.sin(x*sympy.exp(y))
In [141]: sympy.integrate(expr, x)
Out[141]: -e-ycos(xey)
In [142]: expr = (x + y)**2
In [143]: sympy.integrate(expr, x)
Out[143]:  $\frac{x^3}{3} + x^2y + xy^2$ 
```

We can carry out multiple integrations by passing more than one symbol, or multiple tuples that contain symbols and their integration limits.

```
In [144]: sympy.integrate(expr, x, y)
Out[144]:
In [145]: sympy.integrate(expr, (x, 0, 1), (y, 0, 1))
Out[145]:  $\frac{7}{6}$ 
```

Series

Series expansions are an important tool in many disciplines in computing. With a series expansion, an arbitrary function can be written as a polynomial, with coefficients given by the function's derivatives at the point around which the series expansion is made. The nth-order approximation of the function is obtained by truncating the series expansion at some order n. In SymPy, the series expansion of a function or an expression can be computed using the function `sympy.series` or the `series` method available in SymPy expression instances. The first argument to `sympy.series` is a function or expression to be expanded, followed by a symbol with respect to which the expansion is to be computed (it can be omitted for single-variable expressions and functions). In addition, it is also possible to request a particular point around which the series expansions are to be performed (using the `x0` keyword argument, with default `x0=0`), specifying the order of the expansion (using the `n` keyword argument, with default `n=6`) and specifying the direction from which the series is computed (i.e., from below or above `x0` using the `dir` keyword argument, which defaults to `dir='+'`).

For an undefined $f(x)$ function, the expansion up to the sixth order around $x0=0$ is computed as follows.

```
In [147]: x, y = sympy.symbols("x, y")
In [148]: f = sympy.Function("f")(x)
In [149]: sympy.series(f, x)
Out[149]: 
$$\begin{aligned} & f(0) + x \left. \frac{d}{dx} f(x) \right|_{x=0} + \frac{x^2}{2} \left. \frac{d^2}{dx^2} f(x) \right|_{x=0} + \frac{x^3}{6} \left. \frac{d^3}{dx^3} f(x) \right|_{x=0} \\ & + \frac{x^4}{24} \left. \frac{d^4}{dx^4} f(x) \right|_{x=0} + \frac{x^5}{120} \left. \frac{d^5}{dx^5} f(x) \right|_{x=0} + \mathcal{O}(x^6) \end{aligned}$$

```

To change the point around which the function is expanded, we specify the `x0` argument as in the following example.

```
In [147]: x0 = sympy.Symbol("{x_0}")
In [151]: f.series(x, x0, n=2)
Out[151]: 
$$f(x_0) + (x - x_0) \left. \frac{d}{d\xi_1} f(\xi_1) \right|_{\xi_1=x_0} + \mathcal{O}((x - x_0)^2; x \rightarrow x_0)$$

```

This also specified `n=2`, to request a series expansion with only terms up to and including the second order. Note that the order object represents the errors due to the truncated terms $\mathcal{O}(\dots)$. The order object is useful for keeping track of the order of an expression when computing with series expansions, such as multiplying or adding different expansions. However, for concrete numerical evolution, removing the order term from the expression is necessary, which can be done using the `remove0` method.

In [152]: `f.series(x, x0, n=2).remove0()`

Out[152]:
$$f(x_0) + (x - x_0) \frac{d}{d\xi_1} f(\xi_1) \Big|_{\xi_1=x_0}$$

While the expansions shown in the preceding text were computed for an unspecified $f(x)$ function, we can naturally also compute the series expansions of specific functions and expressions. In those cases, we obtain specific evaluated results. For example, we can easily generate the well-known expansions of many standard mathematical functions.

In [153]: `sympy.cos(x).series()`

Out[153]:
$$1 - \frac{x^2}{2} + \frac{x^4}{24} + \mathcal{O}(x^6)$$

In [154]: `sympy.sin(x).series()`

Out[154]:
$$x - \frac{x^3}{6} + \frac{x^5}{120} + \mathcal{O}(x^6)$$

In [155]: `sympy.exp(x).series()`

Out[155]:
$$1 + x + \frac{x^2}{2} + \frac{x^3}{6} + \frac{x^4}{24} + \frac{x^5}{120} + \mathcal{O}(x^6)$$

In [156]: `(1/(1+x)).series()`

Out[156]:
$$1 - x + x^2 - x^3 + x^4 - x^5 + \mathcal{O}(x^6)$$

Arbitrary expressions of symbols and functions can generally be multivariable functions.

In [157]: `expr = sympy.cos(x) / (1 + sympy.sin(x * y))`

In [158]: `expr.series(x, n=4)`

Out[158]:
$$1 - xy + x^2 \left(y^2 - \frac{1}{2} \right) + x^3 \left(-\frac{5y^3}{6} + \frac{y}{2} \right) + \mathcal{O}(x^4)$$

In [159]: `expr.series(y, n=4)`

Out[159]:
$$\cos(x) - xy \cos(x) + x^2 y^2 \cos(x) - \frac{5x^3 y^3 \cos(x)}{6} + \mathcal{O}(y^4)$$

Limits

Another important tool in calculus is limits, which denotes the value of a function as one of its dependent variables approaches a specific value or as the variable's value approaches negative or positive infinity. An example of a limit is one of the definitions of the derivative.

$$\frac{d}{dx} f(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

While limits are more of a theoretical tool and do not have as many practical applications as series expansions, it is still useful to compute limits using SymPy. In SymPy, limits can be evaluated using the `sympy.limit` function, which takes an expression, a symbol it depends on, and the value that the symbol

approaches in the limit. For example, to compute the limit of the $\sin(x)/x$ function, as the variable x goes to zero, that is, $\lim_{x \rightarrow 0} \sin(x)/x$, we can use the following.

```
In [161]: sympy.limit(sympy.sin(x) / x, x, 0)
Out[161]: 1
```

Here we obtained the well-known answer 1 for this limit. We can also use `sympy.limit` to compute symbolic limits, which can be illustrated by computing derivatives using the previous definition (although it is more efficient to use `sympy.diff`).

```
In [162]: f = sympy.Function('f')
...: x, h = sympy.symbols("x, h")
In [163]: diff_limit = (f(x + h) - f(x))/h
In [164]: sympy.limit(diff_limit.subs(f, sympy.cos), h, 0)
Out[164]: -sin(x)
In [165]: sympy.limit(diff_limit.subs(f, sympy.sin), h, 0)
Out[165]: cos(x)
```

A more practical example of using limits is to find the asymptotic behavior as a function, for example, as its dependent variable approaches infinity. As an example, consider the $f(x) = (x^2 - 3x)/(2x - 2)$ function, and suppose we are interested in the large- x dependence of this function. It is in the form $f(x) \rightarrow px + q$, and we can compute p and q using `sympy.limit` as in the following.

```
In [166]: expr = (x**2 - 3*x) / (2*x - 2)
In [167]: p = sympy.limit(expr/x, x, sympy.oo)
In [168]: q = sympy.limit(expr - p*x, x, sympy.oo)
In [169]: p, q
Out[169]:  $\left(\frac{1}{2}, -1\right)$ 
```

Thus, the asymptotic behavior of $f(x)$ as x becomes large is the linear function $f(x) \rightarrow x/2 - 1$.

Sums and Products

Sums and products can be symbolically represented using the SymPy classes `sympy.Sum` and `sympy.Product`. They both take an expression as their first argument and as a second argument, they take a tuple of the form $(n, n1, n2)$, where n is a symbol and $n1$ and $n2$ are the lower and upper limits for the symbol n , in the sum or product, respectively. After `sympy.Sum` or `sympy.Product` objects have been created, they can be evaluated using the `doit` method.

```
In [171]: n = sympy.symbols("n", integer=True)
In [172]: x = sympy.Sum(1/(n**2), (n, 1, oo))
In [173]: x
Out[173]:  $\sum_{n=1}^{\infty} \frac{1}{n^2}$ 
In [174]: x.doit()
Out[174]:  $\frac{\pi^2}{6}$ 
```

```
In [175]: x = sympy.Product(n, (n, 1, 7))
```

```
In [176]: x
```

```
Out[176]:  $\prod_{n=1}^7 n$ 
```

```
In [177]: x.doit()
```

```
Out[177]: 5040
```

Note that the sum in the previous example was specified with an upper limit of infinity. Therefore, this sum was not evaluated by explicit summation but was computed analytically. SymPy can evaluate many summations of this type, including when the summand contains symbolic variables other than the summation index, such as in the following example.

```
In [178]: x = sympy.Symbol("x")
```

```
In [179]: sympy.Sum((x)**n/(sympy.factorial(n)), (n, 1, oo)).doit().simplify()
```

```
Out[179]:  $e^x - 1$ 
```

Equations

Equation solving is a fundamental part of mathematics with applications in nearly every branch of science and technology, and it is immensely important. SymPy can solve a wide variety of equations symbolically, although many equations cannot be solved analytically, even in principle. If an equation, or a system of equations, can be solved analytically, there is a good chance that SymPy can find the solution. If not, numerical methods might be the only option.

In its simplest form, equation solving involves a single equation with a single unknown variable and no additional parameters: for example, finding the value of x that satisfies the second-degree polynomial equation $x^2 + 2x - 3 = 0$. This equation is easy to solve, but in SymPy, we can use the `sympy.solve` function to find the solutions of x that satisfy this equation using the following.

```
In [170]: x = sympy.Symbol("x")
```

```
In [171]: sympy.solve(x**2 + 2*x - 3)
```

```
Out[171]: [-3, 1]
```

That is, the solutions are $x=-3$ and $x=1$. The argument to the `sympy.solve` function is an expression that is solved under the assumption that it equals zero. When this expression contains more than one symbol, the variable to be solved must be given as a second argument. The following is an example.

```
In [172]: a, b, c = sympy.symbols("a, b, c")
```

```
In [173]: sympy.solve(a * x**2 + b * x + c, x)
```

```
Out[173]:  $\left[ \frac{1}{2a} \left( -b + \sqrt{-4ac + b^2} \right), -\frac{1}{2a} \left( b + \sqrt{-4ac + b^2} \right) \right]$ 
```

In this case, the resulting solutions are expressions that depend on the symbols representing the parameters in the equation.

The `sympy.solve` function can also solve other types of equations, including trigonometric expressions.

```
In [174]: sympy.solve(sympy.sin(x) - sympy.cos(x), x)
```

```
Out[174]:  $\left[ -\frac{3\pi}{4}, \right]$ 
```

The `sympy.solve` function can also solve equations whose solution can be expressed in terms of special functions.

```
In [180]: sympy.solve(sympy.exp(x) + 2 * x, x)
```

```
Out[180]: [-LambertW(1/2)]
```

However, when dealing with general equations, even for a univariate case, it is not uncommon to encounter equations that are not solvable algebraically or that SymPy cannot solve. In these cases, SymPy returns a formal solution, which can be numerically evaluated if needed or raise an error if no method is available for that particular type of equation.

```
In [176]: sympy.solve(x**5 - x**2 + 1, x)
```

```
Out[176]: [RootOf(x^5 - x^2 + 1, 0), RootOf(x^5 - x^2 + 1, 1), RootOf(x^5 - x^2 + 1, 2), RootOf(x^5 - x^2 + 1, 3), RootOf(x^5 - x^2 + 1, 4)]
```

```
In [177]: sympy.solve(sympy.tan(x) + x, x)
```

```
NotImplementedError
```

```
Traceback (most recent call last)
```

```
...
```

```
NotImplementedError: multiple generators [x, tan(x)] No algorithms are implemented to solve
equation x + tan(x)
```

Solving a system of equations for more than one unknown variable in SymPy is a straightforward generalization of the procedure used for univariate equations. Instead of passing a single expression as the first argument to `sympy.solve`, a list of expressions representing the system of equations is used, and in this case, the second argument should be a list of symbols to solve for. For example, the following two examples demonstrate how to solve two systems that are linear and nonlinear equations in x and y , respectively.

```
In [178]: eq1 = x + 2 * y - 1
```

```
...: eq2 = x - y + 1
```

```
In [179]: sympy.solve([eq1, eq2], [x, y], dict=True)
```

```
Out[179]: [{x: -1/3, y: 2/3}]
```

```
In [180]: eq1 = x**2 - y
```

```
...: eq2 = y**2 - x
```

```
In [181]: sols = sympy.solve([eq1, eq2], [x, y], dict=True)
```

```
In [182]: sols
```

```
Out[182]: [{x: 0, y: 0}, {x: 1, y: 1}, {x: -1/2 + sqrt(3)*I/2, y: -1/2 - sqrt(3)*I/2}, {x: (1 - sqrt(3)*I)**2/4, y: -1/2 + sqrt(3)/2}]
```

Note that in both these examples, the `sympy.solve` function returns a list where each element represents a solution to the equation system. The optional keyword argument `dict=True` was also used to request that each solution be returned in dictionary format, which maps the symbols that have been solved to their values. This dictionary can conveniently be used in, for example, calls to `subs`, as in the following code that verifies that each solution indeed satisfies the two equations.

```
In [183]: [eq1.subs(sol).simplify() == 0 and eq2.subs(sol).simplify() == 0 for sol in sols]
```

```
Out[183]: [True, True, True, True]
```

Linear Algebra

Linear algebra is another fundamental branch of mathematics with important applications throughout scientific and technical computing. It concerns vectors, vector spaces, and linear mappings between vector spaces, which can be represented as matrices. In SymPy, we can represent vectors and matrices symbolically using the `sympy.Matrix` class, whose elements can be represented by numbers, symbols, or arbitrary symbolic expressions. To create a matrix with numerical entries—as in the case of NumPy arrays in Chapter 2, pass a Python list to `sympy.Matrix`.

```
In [184]: sympy.Matrix([1, 2])
Out[184]: 
$$\begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

In [185]: sympy.Matrix([[1, 2]])
Out[185]: [1 2]
In [186]: sympy.Matrix([[1, 2], [3, 4]])
Out[186]: 
$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

```

As this example demonstrates, a single list generates a column vector, while a matrix requires a nested list of values. Note that unlike the multidimensional arrays in NumPy discussed in Chapter 2, the `sympy.Matrix` object in SymPy can only have one or two dimensions: vectors and matrices.

Another way of creating new `sympy.Matrix` objects is to pass as arguments the number of rows, the number of columns, and a function that takes the row and column index as arguments and returns the value of the corresponding element.

```
In [187]: sympy.Matrix(3, 4, lambda m, n: 10 * m + n)
Out[187]: 
$$\begin{bmatrix} 0 & 1 & 2 & 3 \\ 10 & 11 & 12 & 13 \\ 20 & 21 & 22 & 23 \end{bmatrix}$$

```

The most powerful features of SymPy's matrix objects, which distinguish it from, for example, NumPy arrays, are that the elements can be symbolic expressions. For example, an arbitrary 2×2 matrix can be represented with a symbolic variable for each element.

```
In [188]: a, b, c, d = sympy.symbols("a, b, c, d")
In [189]: M = sympy.Matrix([[a, b], [c, d]])
In [190]: M
Out[190]: 
$$\begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

```

Such matrices can also be used in computations, which then remain parameterized with the symbolic values of the elements. The usual arithmetic operators are implemented for matrix objects, but note that multiplication operator `*` denotes matrix multiplication.

```
In [191]: M * M
Out[191]: 
$$\begin{bmatrix} a^2 + bc & ab + bd \\ ac + cd & bc + d^2 \end{bmatrix}$$

In [192]: x = sympy.Matrix(sympy.symbols("x_1, x_2"))
In [194]: M * x
Out[194]: 
$$\begin{bmatrix} ax_1 + bx_2 \\ cx_1 + dx_2 \end{bmatrix}$$

```

In addition to arithmetic operations, many standard linear algebra operations on vectors and matrices are implemented as SymPy functions and methods of the `sympy.Matrix` class. Table 3-4 summarizes the frequently used linear algebra-related functions (see the docstring for `sympy.Matrix` for a complete list). SymPy matrices can also be used in an element-oriented fashion with indexing and slicing operations that closely resemble those discussed for NumPy arrays in Chapter 2.

Table 3-4. Selected Functions and Methods for Operating on SymPy Matrices

Function/Method	Description
<code>transpose/T</code>	Compute the transpose of a matrix.
<code>adjoint/H</code>	Compute the adjoint of a matrix.
<code>Trace</code>	Compute the trace (sum of diagonal elements) of a matrix.
<code>Det</code>	Compute the determinant of a matrix.
<code>Inv</code>	Compute the inverse of a matrix.
<code>LUdecomposition</code>	Compute the LU decomposition of a matrix.
<code>LUsolve</code>	Solve a linear system of equations in the form $Mx = b$, for the unknown vector x , using LU factorization.
<code>QRdecomposition</code>	Compute the QR decomposition of a matrix.
<code>QRsolve</code>	Solve a linear system of equations in the form $Mx = b$, for the unknown vector x , using QR factorization.
<code>diagonalize</code>	Diagonalize matrix M , such that it can be written in the form $D = P^{-1}MP$, where D is diagonal.
<code>Norm</code>	Compute the norm of a matrix.
<code>nullspace</code>	Compute a set of vectors that span the null space of a Matrix.
<code>Rank</code>	Compute the rank of a matrix.
<code>singular_values</code>	Compute the singular values of a matrix.
<code>Solve</code>	Solve a linear system of equations in the form $Mx = b$.

Consider the following parameterized linear equation system as an example of a problem that can be solved with symbolic linear algebra using SymPy, but which is not directly solvable with purely numerical approaches:

$$x + p \cdot y = b_1,$$

$$q \cdot x + y = b_2,$$

which we would like to solve for the unknown variables x and y . Here p , q , b_1 , and b_2 are unspecified parameters. In matrix form, we can write these two equations as follows.

$$\begin{pmatrix} 1 & p \\ q & 1 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}.$$

With purely numerical methods, we would have to choose values of the parameters p and q before we could begin to solve this problem, for example, using an LU factorization (or by computing the inverse) of the matrix on the left-hand side of the equation. With a symbolic computing approach, on the other hand, we can directly proceed with computing the solution as if we carried out the calculation analytically by hand. With SymPy, we can simply define symbols for the unknown variables and parameters and set up the required matrix objects.

```
In [195]: p, q = sympy.symbols("p, q")
In [196]: M = sympy.Matrix([[1, p], [q, 1]])
In [203]: M
Out[203]:  $\begin{bmatrix} 1 & p \\ q & 1 \end{bmatrix}$ 
In [197]: b = sympy.Matrix(sympy.symbols("b_1, b_2"))
In [198]: b
Out[198]:  $[b_1 \ b_2]$ 
```

We can then use, for example, the `LUsolve` method to solve the linear equation system.

```
In [199]: x = M.LUsolve(b)
In [200]: x
Out[200]:  $\begin{bmatrix} b_1 - \frac{p(-b_1q + b_2)}{-pq + 1} \\ \frac{-b_1q + b_2}{-pq + 1} \end{bmatrix}$ 
```

Alternatively, we could also directly compute the inverse of the matrix M and multiply it with the vector b .

```
In [201]: x = M.inv() * b
In [202]: x
Out[202]:  $\begin{bmatrix} b_1 \left( \frac{pq}{-pq + 1} + 1 \right) - \frac{b_2 p}{-pq + 1} \\ -\frac{b_1 q}{-pq + 1} + \frac{b_2}{-pq + 1} \end{bmatrix}$ 
```

However, computing the inverse of a matrix is more difficult than performing the LU factorization, so if solving the equation $Mx = b$ is the objective, as it was here, then using LU factorization is more efficient. This becomes particularly noticeable for larger equation systems. With both methods considered here, we obtain a symbolic expression for the solution that can be evaluated for any parameter values without recomputing the solution. This is the strength of symbolic computing and an example of how it sometimes can excel over direct numerical computing. The example considered here could also be solved easily by hand. However, as the number of equations and unspecified parameters grows, analytical treatment by hand quickly becomes prohibitively lengthy and tedious. With the help of a computer algebra system such as SymPy, we can push the limits of which problems can be treated analytically.

Summary

This chapter introduced computer-assisted symbolic computing using Python and the SymPy library. Although analytical and numerical techniques are often considered separately, it is a fact that analytical methods underpin everything in computing and are essential in developing algorithms and numerical methods. Whether analytical mathematics is carried out by hand or using a computer algebra system such as SymPy, it is an essential tool for computational work.

I would like to encourage the following approach. Analytical and numerical methods are closely intertwined, and it is often worthwhile to start analyzing a computational problem with analytical and symbolic methods. When such methods are unfeasible, it is time to use numerical methods. Furthermore, by directly applying numerical methods to a problem before analyzing it analytically, one likely ends up solving a more difficult computational problem than is necessary.

Further Reading

Instant SymPy Starter by Ronan Lamy (Packt, 2013) is a short introduction to SymPy. The official SymPy documentation also provides a great tutorial for starting SymPy; it is available at <http://docs.sympy.org/latest/tutorial/index.html>.