

CHAPTER 5



Equation Solving

The previous chapters discussed general methodologies and techniques, namely, array-based numerical computing, symbolic computing, and visualization. These methods are the cornerstones of scientific computing that make up a fundamental toolset we have at our disposal when attacking computational problems.

This chapter begins our exploration of solving problems from different domains of applied mathematics and computational sciences using the basic techniques. The topic of this chapter is algebraic equation solving. This is a broad topic that requires the application of theory and techniques from multiple fields of mathematics. When discussing equation solving, we must distinguish between univariate and multivariate equations (i.e., equations containing one unknown variable or many unknown variables). In addition, we need to distinguish between linear and nonlinear equations. This classification is useful because solving equations of these types requires applying different mathematical methods and approaches.

Let's begin with linear equation systems, which are tremendously useful and have important applications in every field of science. Linear algebra theory enables us to solve linear equations straightforwardly, whereas nonlinear equations are generally difficult to solve and typically require more complicated and computationally demanding methods. Important reasons why linear systems are so universal are because they are readily solvable and can be used in local approximations of more complicated nonlinear systems. For example, by considering small variations from an expansion point, a nonlinear system can often be approximated by a linear system in the local vicinity of the expansion point. However, linearization can only describe local properties, and other techniques are required for the global analysis of nonlinear problems. Such methods typically employ iterative approaches for gradually constructing an increasingly accurate estimate of the solution.

This chapter uses SymPy to solve equations symbolically when possible and uses the linear algebra module from the SciPy library to solve linear equation systems numerically. The root-finding functions in the optimize module of SciPy are used for tackling nonlinear problems.

■ **SciPy** SciPy is a Python library, the collective name of the scientific computing environment for Python, and the umbrella organization for many of the core libraries for scientific computing with Python. The SciPy library is rather a collection of libraries for high-level scientific computing, which are more or less independent of each other. The SciPy library is built on top of NumPy, which provides the basic array data structures and fundamental operations on such arrays. The modules in SciPy provide domain-specific high-level computation methods, such as routines for linear algebra, optimization, interpolation, integration, and much more. At the time of writing, the most recent version of SciPy is 1.11.1. See www.scipy.org for more information.

Importing Modules

The SciPy module `scipy` should be considered a collection of selectively imported modules when required. This chapter uses the `scipy.linalg` module for solving linear systems of equations, and the `scipy.optimize` module for solving nonlinear equations. Assume that these SciPy modules are imported as follows.

```
In [1]: from scipy import linalg as la
In [2]: from scipy import optimize
```

This chapter also uses the NumPy, SymPy, and Matplotlib libraries introduced in earlier chapters. Assume that those libraries are imported using the following convention.

```
In [3]: import sympy
In [4]: sympy.init_printing()
In [5]: import numpy as np
In [6]: import matplotlib.pyplot as plt
```

Linear Equation Systems

An important application of linear algebra is solving systems of linear equations. We have already encountered linear algebra functionality in the SymPy library in Chapter 3. There are also linear algebra modules in the NumPy and SciPy libraries, `numpy.linalg` and `scipy.linalg`, which together provide linear algebra routines for numerical problems that are completely specified in terms of numerical factors and parameters.

A linear equation system can generally be written in the form

$$a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n = b_1,$$

$$a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n = b_2,$$

...

$$a_{m1}x_1 + a_{m2}x_2 + \dots + a_{mn}x_n = b_m.$$

This is a linear system of m equations in n unknown variables $\{x_1, x_2, \dots, x_n\}$, where a_{mn} and b_m are known parameters or constant values. When working with linear equation systems, it is convenient to write them in matrix form:

$$\begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_m \end{pmatrix},$$

or simply $Ax = b$, where A is a $m \times n$ matrix, b is a $m \times 1$ matrix (or m -vector), and x is the unknown $n \times 1$ solution matrix (or n -vector). Depending on the properties of the matrix A , the solution vector x may or may not exist, and if a solution does exist, it is not necessarily unique. However, if a solution exists, then vector b can be interpreted as a linear combination of the columns of matrix A , where the coefficients are given by the elements in the solution vector x .

A system for which $n < m$ is said to be underdetermined because it has fewer equations than unknown and therefore cannot completely determine a unique solution. If, on the other hand, $m > n$, then the equations are said to be overdetermined. This generally leads to conflicting constraints, and no solution exists to the equation system.

Square Systems

Square systems with $m = n$ is an important special case. It corresponds to the situation where the number of equations equals the number of unknown variables and can potentially have a unique solution. For a unique solution to exist, the matrix A must be *nonsingular*, in which case the inverse of A exists, and the solution can be written as $x = A^{-1}b$. If the matrix A is singular, that is, the rank of the matrix is less than n , $\text{rank}(A) < n$, or, equivalently, if its determinant is zero, $\det A = 0$, then the equation $Ax = b$ can either have no solution or infinitely many solutions, depending on the right-hand side vector b . For a matrix with rank deficiency, $\text{rank}(A) < n$, some columns or rows can be expressed as linear combinations of other columns or vectors. Therefore, they correspond to equations without new constraints, and the system is underdetermined. Computing the rank of the matrix A that defines a linear equation system is a useful method that can tell us whether the matrix is singular or not and whether a solution exists.

When A has full rank, the solution is guaranteed to exist. However, it may be impossible to accurately compute the solution. The *condition number* of the matrix, $\text{cond}(A)$, measures how well or poorly conditioned a linear equation system is. If the conditioning number is close to 1, if the system is said to be *well-conditioned* (condition number 1 is ideal), and if the condition number is large, the system is *ill-conditioned*. The solution to an equation system that is ill-conditioned can have large errors. A simple error analysis can provide an intuitive interpretation of the condition number. Assume a linear equation system in the form of $Ax = b$, where x is the solution vector. Now consider a small variation of b , say δb , which gives a corresponding change in the solution, δx , given by $A(x + \delta x) = b + \delta b$. Because of the linearity of the equation, we have $A\delta x = \delta b$. An important question to consider now is: how large is the relative change in x compared to the relative change in b ? Mathematically, we can formulate this question in terms of the ratios of the norms of these vectors. Specifically, we are interested in comparing $\|\delta x\|/\|x\|$ and $\|\delta b\|/\|b\|$, where $\|x\|$ denotes the norm of x . Using the matrix norm relation $\|Ax\| \leq \|A\| \cdot \|x\|$, we can write the following.

$$\frac{\|\delta x\|}{\|x\|} = \frac{\|A^{-1}\delta b\|}{\|x\|} \leq \frac{\|A^{-1}\| \cdot \|\delta b\|}{\|x\|} = \frac{\|A^{-1}\| \cdot \|b\|}{\|x\|} \cdot \frac{\|\delta b\|}{\|b\|} \leq \|A^{-1}\| \cdot \|A\| \cdot \frac{\|\delta b\|}{\|b\|}$$

A bound of the relative error in the solution x , given a relative error in the b vector, is therefore given by $\text{cond}(A) \equiv \|A^{-1}\| \cdot \|A\|$, which by definition is the condition number of the matrix A . This means that for linear equation systems characterized by an ill-conditioned matrix A , even a small perturbation in the b vector can give large errors in the solution vector x . This is particularly relevant in numerical solutions using floating-point numbers, which are only approximations of real numbers. When numerically solving a system of linear equations, it is therefore important to look at the condition number to estimate the accuracy of the solution.

The rank, condition number, and norm of a symbolic matrix can be computed in SymPy using the Matrix methods `rank`, `condition_number`, and `norm`, and for numerical problems, we can use the NumPy functions `np.linalg.matrix_rank`, `np.linalg.cond`, and `np.linalg.norm`. For example, consider the following system of two linear equations.

$$2x_1 + 3x_2 = 4$$

$$5x_1 + 4x_2 = 3$$

These two equations correspond to lines in the (x_1, x_2) plane, and their intersection is the solution to the equation system. As shown in Figure 5-1, which graphs the lines corresponding to the two equations, the lines intersect at $(-1, 2)$.




Figure 5-1. Graphical representation of a system of two linear equations

We can define this problem in SymPy by creating matrix objects for A and b and computing the rank, condition number, and norm of the matrix A using the following.

```
In [8]: A = sympy.Matrix([[2, 3], [5, 4]])
In [9]: b = sympy.Matrix([4, 3])
In [10]: A.rank()
Out[10]: 2
In [11]: A.condition_number()
Out[11]:  $\frac{\sqrt{27+2\sqrt{170}}}{\sqrt{27-2\sqrt{170}}}$ 
In [12]: sympy.N(_)
Out[12]: 7.58240137440151
In [13]: A.norm()
Out[13]:  $3\sqrt{6}$ 
```

We can do the same thing in NumPy/SciPy using NumPy arrays for A and b and functions from the `np.linalg` and `scipy.linalg` modules.

```
In [14]: A = np.array([[2, 3], [5, 4]])
In [15]: b = np.array([4, 3])
In [16]: np.linalg.matrix_rank(A)
Out[16]: 2
In [17]: np.linalg.cond(A)
Out[17]: 7.5824013744
In [18]: np.linalg.norm(A)
Out[18]: 7.34846922835
```

A direct approach to solving the linear problem is to compute the inverse of the matrix A and multiply it with the vector b , as used, for example, in the previous analytical discussions. However, this is not the most efficient computational method to find the solution vector x . A better method is LU factorization of the matrix A , such that $A = LU$ where L is a lower triangular matrix and U is an upper triangular matrix.

Given L and U , the solution vector x can be efficiently constructed by first solving $Ly = b$ with forward substitution and then solving $Ux = y$ with backward substitution. Since L and U are triangular matrices, these two procedures are computationally efficient.

In SymPy, we can perform a symbolic LU factorization by using the `LUdecomposition` method of the `sympy.Matrix` class. This method returns new `Matrix` objects for the L and U matrices, as well as a row swap matrix. When we are interested in solving an equation system $Ax = b$, we do not explicitly need to calculate the L and U matrices. But we can use the `LUsolve` method, which performs the LU factorization internally and solves the equation system using those factors. Returning to the previous example, we can compute the L and U factors and solve the equation system using the following.

```
In [19]: A = sympy.Matrix([[2, 3], [5, 4]])
In [20]: b = sympy.Matrix([4, 3])
In [21]: L, U, _ = A.LUdecomposition()
In [22]: L
Out[22]:  $\begin{bmatrix} 1 & 0 \\ 5/2 & 1 \end{bmatrix}$ 
In [23]: U
Out[23]:  $\begin{bmatrix} 2 & 3 \\ 0 & -7/2 \end{bmatrix}$ 
In [24]: L * U
Out[24]:  $\begin{bmatrix} 2 & 3 \\ 5 & 4 \end{bmatrix}$ 
In [25]: x = A.solve(b); x # equivalent to A.LUsolve(b)
Out[25]:  $\begin{bmatrix} -1 \\ 2 \end{bmatrix}$ 
```

For numerical problems, we can use the `la.lu` function from SciPy's linear algebra module. It returns a permutation matrix P and the L and U matrices, such that $A = PLU$. Like with SymPy, we can solve the linear system $Ax = b$ without explicitly calculating the L and U matrices using the `la.solve` function, which takes the A matrix and the b vector as arguments. This is generally the preferred method for solving numerical linear equation systems using SciPy.

```
In [26]: A = np.array([[2, 3], [5, 4]])
In [27]: b = np.array([4, 3])
In [28]: P, L, U = la.lu(A)
In [29]: L
Out[29]: array([[ 1. ,  0. ],
               [ 0.4,  1. ]])
In [30]: U
Out[30]: array([[ 5. ,  4. ],
               [ 0. ,  1.4]])
In [31]: P.dot(L.dot(U))
Out[31]: array([[ 2.,  3.],
               [ 5.,  4.]])
In [32]: la.solve(A, b)
Out[32]: array([-1.,  2.])
```

The advantage of using SymPy is that we may obtain exact results, and we can also include symbolic variables in the matrices. However, not all problems are solvable symbolically or some may give exceedingly lengthy results. On the other hand, the advantage of using a numerical approach with NumPy/SciPy is that we are guaranteed to obtain a result. However, it will be an approximate solution due to floating-point errors. See the following code (In [38]) for an example that illustrates the differences between the symbolic and numerical approaches and that numerical approaches can be sensitive for equation systems with large condition numbers. This example solves the equation system

$$\begin{pmatrix} 1 & \sqrt{p} \\ 1 & \frac{1}{\sqrt{p}} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$

which for $p = 1$ is singular and for p in the vicinity of one is ill-conditioned. Using SymPy, the solution is easily found to be as follows.

```
In [33]: p = sympy.symbols("p", positive=True)
In [34]: A = sympy.Matrix([[1, sympy.sqrt(p)], [1, 1/sympy.sqrt(p)]])
In [35]: b = sympy.Matrix([1, 2])
In [36]: x = A.solve(b)
In [37]: x
```

```
Out[37]: 
$$\begin{pmatrix} \frac{2p-1}{p-1} \\ -\frac{\sqrt{p}}{p-1} \end{pmatrix}$$

```

A comparison of this symbolic solution and the numerical solution is shown in Figure 5-2. Here, the errors in the numerical solution are due to numerical floatingpoint errors, and the numerical errors are significantly larger in the vicinity of $p = 1$, where the system has a large condition number. Also, if there are other sources of errors in either A or b , the corresponding errors in x can be even more severe.

```
In [38]: # Symbolic problem specification
...: p = sympy.symbols("p", positive=True)
...: A = sympy.Matrix([[1, sympy.sqrt(p)], [1, 1/sympy.sqrt(p)]])
...: b = sympy.Matrix([1, 2])
...:
...: # Solve symbolically
...: x_sym_sol = A.solve(b)
...: Acond = A.condition_number().simplify()
...:
...: # Numerical problem specification
...: AA = lambda p: np.array([[1, np.sqrt(p)], [1, 1/np.sqrt(p)]])
...: bb = np.array([1, 2])
...: x_num_sol = lambda p: np.linalg.solve(AA(p), bb)
...:
...: # Graph the difference between the symbolic (exact) and numerical results
...: fig, axes = plt.subplots(1, 2, figsize=(12, 4))
...:
...: p_vec = np.linspace(0.9, 1.1, 200)
...: for n in range(2):
...:     x_sym = np.array([x_sym_sol[n].subs(p, pp).evalf() for pp in p_vec])
```

```

...:     x_num = np.array([x_num_sol(pp)[n] for pp in p_vec])
...:     axes[0].plot(p_vec, (x_num - x_sym)/x_sym, 'k')
...: axes[0].set_title("Error in solution\n(numerical - symbolic)/symbolic")
...: axes[0].set_xlabel(r'$p$', fontsize=18)
...:
...: axes[1].plot(p_vec, [Acond.subs(p, pp).evalf() for pp in p_vec])
...: axes[1].set_title("Condition number")
...: axes[1].set_xlabel(r'$p$', fontsize=18)

```




Figure 5-2. Graph of the relative numerical errors (left) and condition number (right) as a function of the parameter p

Rectangular Systems

Rectangular systems, with $m \neq n$, can be either underdetermined or overdetermined. Underdetermined systems have more variables than equations, so the solution cannot be fully determined. Therefore, for such a system, the solution must be given in terms of the remaining free variables. This makes it difficult to treat this type of problem numerically, but a symbolic approach can often be used instead.

For example, consider the underdetermined linear equation system

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 7 \\ 8 \end{pmatrix}.$$

There are three unknown variables, but only two equations impose constraints on the relations between these variables. By writing this equation as $Ax - b = 0$, we can use the SymPy `sympy.solve` function to obtain a solution for x_1 and x_2 parameterized by the remaining free variable x_3 .

```

In [39]: x_vars = sympy.symbols("x_1, x_2, x_3")
In [40]: A = sympy.Matrix([[1, 2, 3], [4, 5, 6]])
In [41]: x = sympy.Matrix(x_vars)
In [42]: b = sympy.Matrix([7, 8])
In [43]: sympy.solve(A*x - b, x_vars)
Out[43]: {x_1 = x_3 - 19/3, 0.5x_2 = -2x_3 + 20/3}

```

This obtained the symbolic solution $x_1 = x_3 - 19/3$ and $x_2 = -2x_3 + 20/3$, which defines a line in the three-dimensional space spanned by $\{x_1, x_2, x_3\}$. Any point on this line, therefore, satisfies this underdetermined equation system.

On the other hand, if the system is overdetermined and has more equations than unknown variables, $m > n$, then we may have more constraints than degrees of freedom, and in general, there is no exact solution to such a system. However, finding an approximate solution to an overdetermined system is often interesting. An example of when this situation arises is data fitting: let's say we have a model where a variable y is a quadratic polynomial in the variable x so that $y = A + Bx + Cx^2$. We would like to fit this model to experimental data. Here, y is nonlinear in x , but y is linear in the three unknown coefficients A , B , and C , and this fact can be used to write the model as a linear equation system. If we collect data for m pairs $\{(x_i, y_i)\}_{i=1}^m$ of the variables x and y , we can write the model as an $m \times 3$ equation system:

$$\begin{pmatrix} 1 & x_1 & x_1^2 \\ \vdots & \vdots & \vdots \\ 1 & x_m & x_m^2 \end{pmatrix} \begin{pmatrix} A \\ B \\ C \end{pmatrix} = \begin{pmatrix} y_1 \\ \vdots \\ y_m \end{pmatrix}.$$

If $m = 3$, we can solve for the unknown model parameters A , B , and C , assuming the system matrix is nonsingular. However, it is intuitively clear that if the data is noisy and if we were to use more than three data points, we should be able to get a more accurate estimate of the model parameters.

However, for $m > 3$, there is generally no exact solution, and we need to introduce an approximate solution that best fits the overdetermined system $Ax \approx b$. A natural definition of the best fit for this system is to minimize the sum of square errors, $\min_x \sum_{i=1}^m (r_i)^2$, where $r = b - Ax$ is the residual vector. This leads to the

least square solution of the problem $Ax \approx b$, which minimizes the distances between the data points and the linear solution. In SymPy, we can solve for the least square solution of an overdetermined system using the `solve_least_squares` method; for numerical problems, we can use the SciPy function `la.lstsq`.

The following code demonstrates how the SciPy `la.lstsq` method can be used to fit the example model considered in the preceding section, and the result is shown in Figure 5-3. First, we define the true parameters of the model, and then we simulate measured data by adding random noise to the true model relation. The least-square problem is then solved using the `la.lstsq` function, which in addition to the solution vector x also returns the total sum of square errors (the residual r), the rank rank , and the singular values sv of the matrix A . However, the following example only uses the solution vector x .

```
In [44]: # define true model parameters
...: x = np.linspace(-1, 1, 100)
...: a, b, c = 1, 2, 3
...: y_exact = a + b * x + c * x**2
...:
...: # simulate noisy data
...: m = 100
...: X = 1 - 2 * np.random.rand(m)
...: Y = a + b * X + c * X**2 + np.random.randn(m)
...:
...: # fit the data to the model using linear least square
...: A = np.vstack([X**0, X**1, X**2]) # see np.vander for alternative
...: sol, r, rank, sv = la.lstsq(A.T, Y)
...:
...: y_fit = sol[0] + sol[1] * x + sol[2] * x**2
...: fig, ax = plt.subplots(figsize=(12, 4))
...:
```



```

...: ax.plot(X, Y, 'go', alpha=0.5, label='Simulated data')
...: ax.plot(x, y_exact, 'k', lw=2, label='True value $y = 1 + 2x + 3x^2$')
...: ax.plot(x, y_fit, 'b', lw=2, label='Least square fit')
...: ax.set_xlabel(r"$x$", fontsize=18)
...: ax.set_ylabel(r"$y$", fontsize=18)
...: ax.legend(loc=2)

```




Figure 5-3. Linear least square fit

A good data fit to a model requires the model used to describe the data to correspond well to the underlying process that produced the data. The following example (In [45]) and Figure 5-4 fit the same data used in the previous example to a linear model and to a higher-order polynomial model (up to order 15). The former case corresponds to underfitting, where a too-simple model is used for the data, and the latter case corresponds to overfitting, where a too-complex model is used for the data, and thus fits the model not only to the trend and relevant patterns but also to the measurement noise. Using an appropriate model is an important and delicate aspect of data fitting.

```

In [45]: # fit the data to the model using linear least square:
...: # 1st order polynomial
...: A = np.vstack([X**n for n in range(2)])
...: sol, r, rank, sv = la.lstsq(A.T, Y)
...: y_fit1 = sum([s * x**n for n, s in enumerate(sol)])
...:
...: # 15th order polynomial
...: A = np.vstack([X**n for n in range(16)])
...: sol, r, rank, sv = la.lstsq(A.T, Y)
...: y_fit15 = sum([s * x**n for n, s in enumerate(sol)])
...:
...: fig, ax = plt.subplots(figsize=(12, 4))
...: ax.plot(X, Y, 'go', alpha=0.5, label='Simulated data')
...: ax.plot(x, y_exact, 'k', lw=2, label='True value $y = 1 + 2x + 3x^2$')
...: ax.plot(x, y_fit1, 'b', lw=2, label='Least square fit [1st order]')
...: ax.plot(x, y_fit15, 'm', lw=2, label='Least square fit [15th order]')
...: ax.set_xlabel(r"$x$", fontsize=18)
...: ax.set_ylabel(r"$y$", fontsize=18)
...: ax.legend(loc=2)

```




Figure 5-4. Graph demonstrating underfitting and overfitting of data using the linear least square method

Eigenvalue Problems

A special system of equations of great theoretical and practical importance is the eigenvalue equation $Ax = \lambda x$, where A is an $N \times N$ square matrix, x is an unknown vector, and λ is an unknown scalar. Here x is an eigenvector and λ an eigenvalue of the matrix A . The eigenvalue equation $Ax = \lambda x$ closely resembles the linear equation system $Ax = b$, but note that both x and λ are unknown here, so we cannot directly apply the same techniques to solve this equation. A standard approach to solve this eigenvalue problem is to rewrite the equation as $(A - I\lambda)x = 0$ and note that for there to exist a nontrivial solution, $x \neq 0$, the matrix $A - I\lambda$ must be singular, and its determinant must be zero, $\det(A - I\lambda) = 0$. This gives a polynomial equation (the characteristic polynomial) of N th order whose N roots give the N eigenvalues, $\{\lambda_n\}_{n=1}^N$. Once the eigenvalues are known, standard forward substitution can solve the equation $(A - I\lambda_n)x_n = 0$ for the N th eigenvector x_n .

SymPy and the linear algebra package in SciPy contain solvers for eigenvalue problems. In SymPy, we can use the `eigenvals` and `eigenvecs` methods of the `Matrix` class, which can compute the eigenvalues and eigenvectors of some matrices with elements that are symbolic expressions. For example, we can use the following to compute the eigenvalues and eigenvectors of a symmetric 2×2 matrix with symbolic elements.

```
In [46]: eps, delta = sympy.symbols("epsilon, Delta")
In [47]: H = sympy.Matrix([[eps, delta], [delta, -eps]])
In [48]: H
Out[48]: 
$$\begin{pmatrix} \varepsilon & \Delta \\ \Delta & -\varepsilon \end{pmatrix}$$

In [49]: H.eigenvals()
Out[49]:  $\{-\sqrt{\varepsilon^2 + \Delta^2}:1, \sqrt{\varepsilon^2 + \Delta^2}:1\}$ 
In [50]: H.eigenvecs()
Out[50]:  $\left( \left( -\sqrt{\varepsilon^2 + \Delta^2}, 1, \left[ \begin{array}{c} -\frac{1}{\varepsilon + \sqrt{\varepsilon^2 + \Delta^2}} \\ 1 \end{array} \right] \right), \left( \sqrt{\varepsilon^2 + \Delta^2}, 1, \left[ \begin{array}{c} -\frac{1}{\varepsilon - \sqrt{\varepsilon^2 + \Delta^2}} \\ 1 \end{array} \right] \right) \right)$ 
```

The return value of the `eigenvals` method is a dictionary where each eigenvalue is a key, and the corresponding value is the multiplicity of that particular eigenvalue. Here the eigenvalues are $-\sqrt{\varepsilon^2 + \Delta^2}$ and $\sqrt{\varepsilon^2 + \Delta^2}$, each with multiplicity one. The return value of `eigenvecs` is more involved: a list is returned where each element is a tuple containing an eigenvalue, the multiplicity of the eigenvalue, and a list of eigenvectors. The number of eigenvectors for each eigenvalue equals the multiplicity. For the current example, we can unpack the value returned by `eigenvecs` and verify that the two eigenvectors are orthogonal using, for example, as follows.

```
In [51]: (eval1, _, vec1), (eval2, _, vec2) = H.eigenvecs()
In [52]: sympy.simplify(vec1[0].T * vec2[0])
Out[52]: [0]
```

Obtaining analytical expressions for eigenvalues and eigenvectors using these methods is often very desirable, but unfortunately, it only works for small matrices. For anything larger than a 3×3 , the analytical expression typically becomes extremely lengthy and cumbersome to work with, even using a computer algebra system such as SymPy. Therefore, we must resort to a fully numerical approach for larger systems. We can use the `la.eigvals` and `la.eig` functions in the SciPy linear algebra package. Matrices that are either Hermitian or real and symmetric have real-valued eigenvalues. For such matrices, it is advantageous to use the `la.eigvalsh` and `la.eigh` functions instead, which guarantees that the eigenvalues returned by the function are stored in a NumPy array with real values. For example, to solve a numerical eigenvalue problem with `la.eig`, we can use the following.

```
In [53]: A = np.array([[1, 3, 5], [3, 5, 3], [5, 3, 9]])
In [54]: evals, vecs = la.eig(A)
In [55]: evals
Out[55]: array([ 13.35310908+0.j, -1.75902942+0.j,  3.40592034+0.j])
In [56]: vecs
Out[56]: array([[ 0.42663918,  0.90353276, -0.04009445],
                [ 0.43751227, -0.24498225, -0.8651975 ],
                [ 0.79155671, -0.35158534,  0.49982569]])
In [57]: la.eigvalsh(A)
Out[57]: array([-1.75902942,  3.40592034, 13.35310908])
```

Since the matrix in this example is symmetric, we could use `la.eigh` and `la.eigvalsh`, giving real-valued eigenvalue arrays, as shown in the cell Out[57] in the preceding code listing.

Nonlinear Equations

This section considers *nonlinear* equations. Systems of linear equations, as considered in the prior sections, are of fundamental importance in scientific computing because they are easily solved and can be used as important building blocks in many computational methods and techniques. However, in natural sciences and engineering disciplines, many, if not most, systems are intrinsically nonlinear.

A linear function $f(x)$, by definition, satisfies additivity $f(x + y) = f(x) + f(y)$ and homogeneity $f(ax) = af(x)$, which can be written together as the superposition principle $f(ax + \beta y) = af(x) + \beta f(y)$. This gives a precise definition of linearity. A *nonlinear* function, in contrast, is a function that does not satisfy these conditions. Nonlinearity is a much broader concept; a function can be nonlinear in many ways. However, an expression that contains a variable with a power greater than one is nonlinear. For example, $x^2 + x + 1$ is nonlinear because of the x^2 term.

A nonlinear equation can always be written in the form $f(x) = 0$, where $f(x)$ is a nonlinear function that seeks the value of x (which can be a scalar or a vector) such that $f(x)$ is zero. This x is the root of the $f(x)$ function, and equation-solving is often called *root finding*. In contrast to the previous section of this chapter, here we need to distinguish between univariate equation solving and multivariate equations, in addition to single equations and systems of equations.

Univariate Equations

A univariate function $f(x)$ is a function that depends only on a single scalar variable x , and the corresponding univariate equation is $f(x) = 0$. Typical examples of this type of equation are polynomials, such as $x^2 - x + 1 = 0$, and expressions containing elementary functions, such as $x^3 - 3 \sin(x) = 0$ and $\exp(x) - 2 = 0$. Unlike linear systems, there are no general methods for determining if a nonlinear equation has a solution or multiple solutions or if a given solution is unique. This can be understood intuitively from the fact that graphs of nonlinear functions correspond to curves that can intersect $y = 0$ in an arbitrary number of ways.

Because of the vast number of possible situations, it is difficult to develop a completely automatic approach to solving nonlinear equations. Analytically, only equations on special forms can be solved exactly. For example, polynomials of up to 4th order, and in some special cases also higher orders, can be solved analytically, and some equations containing trigonometric and other elementary functions may be solvable analytically. In SymPy, we can solve many solvable univariate and nonlinear equations using the `sympy.solve` function. For example, to solve the standard quadratic equation $a + bx + cx^2 = 0$, define an expression for the equation, and pass it to the `sympy.solve` function.

```
In [58]: x, a, b, c = sympy.symbols("x, a, b, c")
In [59]: sympy.solve(a + b*x + c*x**2, x)
Out[59]: [(-b + sqrt(-4*a*c + b**2))/(2*c), -(b + sqrt(-4*a*c + b**2))/(2*c)]
```

The solution is the well-known formula for the solution of this equation. The same method can be used to solve some trigonometric equations.

```
In [60]: sympy.solve(a * sympy.cos(x) - b * sympy.sin(x), x)
Out[60]: [-2*atan((b - sqrt(a**2 + b**2))/a), -2*atan((b + sqrt(a**2 + b**2))/a)]
```

However, in general, nonlinear equations are typically not solvable analytically. For example, equations that contain both polynomial expressions and elementary functions, such as $\sin(x) = x$, are often transcendental and do not have an algebraic solution. If we attempt to solve such an equation using SymPy, we obtain an error in the form of an exception.

```
In [61]: sympy.solve(sympy.sin(x)-x, x)
...
NotImplementedError: multiple generators [x, sin(x)]
No algorithms are implemented to solve equation -x + sin(x)
```

In this situation, we must resort to various numerical techniques. As a first step, graphing the function is often very useful. This can give important clues about the number of solutions to the equation and their approximate locations. This information is often necessary when applying numerical techniques to find good approximations to the roots of the equations. For example, consider the following example (In [62]), which plots four examples of nonlinear functions, as shown in Figure 5-5. From these graphs, we can immediately conclude that the plotted functions, from left to right, have two, three, one, and a large number of roots (at least within the interval being graphed).

```

In [62]: x = np.linspace(-2, 2, 1000)
...: # four examples of nonlinear functions
...: f1 = x**2 - x - 1
...: f2 = x**3 - 3 * np.sin(x)
...: f3 = np.exp(x) - 2
...: f4 = 1 - x**2 + np.sin(50 / (1 + x**2))
...:
...: # plot each function
...: fig, axes = plt.subplots(1, 4, figsize=(12, 3), sharey=True)
...:
...: for n, f in enumerate([f1, f2, f3, f4]):
...:     axes[n].plot(x, f, lw=1.5)
...:     axes[n].axhline(0, ls=':', color='k')
...:     axes[n].set_ylim(-5, 5)
...:     axes[n].set_xticks([-2, -1, 0, 1, 2])
...:     axes[n].set_xlabel(r'$x$', fontsize=18)
...:
...: axes[0].set_ylabel(r'$f(x)$', fontsize=18)
...:
...: titles = [r'$f(x)=x^2-x-1$', r'$f(x)=x^3-3\sin(x)$',
...:           r'$f(x)=\exp(x)-2$',
...:           r'$f(x)=\sin\left(50/(1+x^2)\right)+1-x^2$']
...: for n, title in enumerate(titles):
...:     axes[n].set_title(title)

```




Figure 5-5. Graphs of four examples of nonlinear functions

To find the approximate location of a root to an equation, we can apply one of the many techniques for numerical root finding, which typically applies an iterative scheme where the function is evaluated at successive points until the algorithm has narrowed in on the solution to the desired accuracy. Two standard methods that illustrate the basic idea of how many numerical root-finding methods work are the bisection method and the Newton method.

The bisection method requires a starting interval $[a, b]$ such that $f(a)$ and $f(b)$ have different signs. This guarantees that there is at least one root within this interval. In each iteration, the function is evaluated in the middle point m between a and b , and the sign of the function is different at a and m , and then the new interval $[a, b = m]$ is chosen for the next iteration. Otherwise, the interval $[a = m, b]$ is chosen for the next iteration. This guarantees that in each iteration, the function has a different sign at the two endpoints of the interval. In each iteration, the interval is halved and converges toward the root of the equation. The following code example demonstrates a simple implementation of the bisection method with a graphical visualization of each step, as shown in Figure 5-6.

```

In [63]: # define a function, desired tolerance and starting interval [a, b]
...: f = lambda x: np.exp(x) - 2
...: tol = 0.1
...: a, b = -2, 2
...: x = np.linspace(-2.1, 2.1, 1000)
...:
...: # graph the function f
...: fig, ax = plt.subplots(1, 1, figsize=(12, 4))
...:
...: ax.plot(x, f(x), lw=1.5)
...: ax.axhline(0, ls=':', color='k')
...: ax.set_xticks([-2, -1, 0, 1, 2])
...: ax.set_xlabel(r'$x$', fontsize=18)
...: ax.set_ylabel(r'$f(x)$', fontsize=18)
...:
...: # find the root using the bisection method and visualize
...: # the steps in the method in the graph
...: fa, fb = f(a), f(b)
...:
...: ax.plot(a, fa, 'ko')
...: ax.plot(b, fb, 'ko')
...: ax.text(a, fa + 0.5, r"$a$", ha='center', fontsize=18)
...: ax.text(b, fb + 0.5, r"$b$", ha='center', fontsize=18)
...:
...: n = 1
...: while b - a > tol:
...:     m = a + (b - a)/2
...:     fm = f(m)
...:
...:     ax.plot(m, fm, 'ko')
...:     ax.text(m, fm - 0.5, r"$m_{%d}$" % n, ha='center')
...:     n += 1
...:
...:     if np.sign(fa) == np.sign(fm):
...:         a, fa = m, fm
...:     else:
...:         b, fb = m, fm
...:
...: ax.plot(m, fm, 'r*', markersize=10)
...: ax.annotate(
...:     "Root approximately at %.3f" % m,
...:     fontsize=14, family="serif", xy=(a, fm), xycoords='data',
...:     xytext=(-150, +50), textcoords='offset points',
...:     arrowprops=dict(arrowstyle="->", connectionstyle="arc3, rad=-.5"))
...:
...: ax.set_title("Bisection method")

```




Figure 5-6. Graphical visualization of how the bisection method works

Another standard root-finding method is Newton's method, which converges faster than the bisection method discussed in the previous paragraph. While the bisection method only uses the sign of the function at each point, Newton's method uses the actual function values to obtain a more accurate approximation of the nonlinear function. It approximates the $f(x)$ function with its first-order Taylor expansion $f(x + dx) = f(x) + dx f'(x)$, which is a linear function whose root is easily found to be $x - f(x)/f'(x)$. Of course, this does not need to be a root of the $f(x)$ function, but in many cases, it is a good approximation for getting closer to a root of $f(x)$. By iterating this scheme, $x_{k+1} = x_k - f(x_k)/f'(x_k)$, we may approach the root of the function. A potential problem with this method is that it fails if $f'(x_k)$ is zero at some point x_k . This special case would have to be dealt with in a real implementation of this method. The following example (In [64]) demonstrates how this method can be used to solve for the root of the equation $\exp(x) - 2 = 0$, using SymPy to evaluate the derivative of the $f(x)$ function, and Figure 5-7 visualizes the steps in this root-finding process.

```
In [64]: # define a function, desired tolerance and starting point xk
...: tol = 0.01
...: xk = 2
...:
...: s_x = sympy.symbols("x")
...: s_f = sympy.exp(s_x) - 2
...:
...: f = lambda x: sympy.lambdify(s_x, s_f, 'numpy')(x)
...: fp = lambda x: sympy.lambdify(s_x, sympy.diff(s_f, s_x), 'numpy')(x)
...:
...: x = np.linspace(-1, 2.1, 1000)
...:
...: # setup a graph for visualizing the root finding steps
...: fig, ax = plt.subplots(1, 1, figsize=(12, 4))
...: ax.plot(x, f(x))
...: ax.axhline(0, ls=':', color='k')
...:
...: # iterate Newton's method until convergence to the desired tolerance
...: # has been reached
...: n = 0
...: while f(xk) > tol:
...:     xk_new = xk - f(xk) / fp(xk)
...:     n += 1
...:     xk = xk_new
```

```

....: ax.plot([xk, xk], [0, f(xk)], color='k', ls=':')
....: ax.plot(xk, f(xk), 'ko')
....: ax.text(xk, -.5, r'$x_{%d}$' % n, ha='center')
....: ax.plot([xk, xk_new], [f(xk), 0], 'k-')
....:
....: xk = xk_new
....: n += 1
....:
....: ax.plot(xk, f(xk), 'r*', markersize=15)
....: ax.annotate(
....:     "Root approximately at %.3f" % xk,
....:     fontsize=14, family="serif", xy=(xk, f(xk)), xycoords='data',
....:     xytext=(-150, +50), textcoords='offset points',
....:     arrowprops=dict(arrowstyle="->", connectionstyle="arc3, rad=-.5"))
....:
....: ax.set_title("Newtown's method")
....: ax.set_xticks([-1, 0, 1, 2])

```




Figure 5-7. Visualization of the root-finding steps in Newton's method for the equation $\exp(x) - 2 = 0$

A potential issue with Newton's method is that it requires both the function values and the values of the function's derivative in each iteration. The previous example used SymPy to compute the derivatives symbolically. In an all-numerical implementation, this is not possible, and a numerical approximation of the derivative would be necessary, which would require further function evaluations. A variant of Newton's method that bypasses the requirement to evaluate function derivatives is the secant method, which uses two previous function evaluations to obtain a linear approximation of the function, which can be used to compute a new root estimate. The iteration formula for the secant method is

$$x_{k+1} = x_k - f(x_k) \frac{x_k - x_{k-1}}{f(x_k) - f(x_{k-1})}$$

This is only one example of the many variants and possible refinements to the basic idea of Newton's method. State-of-the-art implementations of numerical root-finding functions typically use either the bisection method of Newton's method or a combination of both but also use various additional strategies, such as higher-order interpolations of the function, to achieve faster convergence.

The SciPy optimize module provides multiple functions for numerical root finding. The `optimize.bisect` and `optimize.newton` functions implement variants of bisection and Newton methods. The `optimize.bisect` takes three arguments: first a Python function (e.g., a lambda function) that represents the mathematical function for the equation for which a root is to be calculated, and the second and third arguments are the lower and upper values of the interval for which to perform the bisection method.

Note that the sign of the function must be different at points a and b for the bisection method to work, as discussed earlier. Using the `optimize.bisect` function, we can calculate the root of the equation $\exp(x) - 2 = 0$ from previous examples.

```
In [65]: optimize.bisect(lambda x: np.exp(x) - 2, -2, 2)
Out[65]: 0.6931471805592082
```

As long as $f(a)$ and $f(b)$ have different signs, this is guaranteed to give a root within the interval $[a, b]$. In contrast, the `optimize.newton` function for Newton's method takes a function as the first argument and an initial guess for the root of the function as the second argument. Optionally, it also takes an argument for specifying the function's derivative using the `fprime` keyword argument. If `fprime` is given, Newton's method is used; otherwise, the secant method is used instead. To find the root of the equation $\exp(x) - 2 = 0$, with and without specifying its derivative, we can use the following.

```
In [66]: x_root_guess = 2
In [67]: f = lambda x: np.exp(x) - 2
In [68]: fprime = lambda x: np.exp(x)
In [69]: optimize.newton(f, x_root_guess)
Out[69]: 0.69314718056
In [70]: optimize.newton(f, x_root_guess, fprime=fprime)
Out[70]: 0.69314718056
```

Note that this method gives us less control over which root is being computed if the function has multiple roots. For instance, there is no guarantee that the root the function returns is closest to the initial guess; we cannot know in advance if the root is larger or smaller than the initial guess.

The SciPy `optimize` module provides additional functions for root finding. In particular, the `optimize.brentq` and `optimize.brenth` functions, which are variants of the bisection method, also work on an interval where the function changes sign. The `optimize.brentq` function is generally considered the preferred all-around root-finding function in SciPy. To find a root of the same equation considered previously, use the `optimize.brentq` and `optimize.brenth` functions, as follows.

```
In [71]: optimize.brentq(lambda x: np.exp(x) - 2, -2, 2)
Out[71]: 0.6931471805599453
In [72]: optimize.brenth(lambda x: np.exp(x) - 2, -2, 2)
Out[72]: 0.6931471805599381
```

Note that these two functions take a Python function for the equation as the first argument and the lower and upper values of the sign-changing interval as the second and third arguments.

Systems of Nonlinear Equations

In contrast to a linear system of equations, we cannot generally write a system of nonlinear equations as a matrix-vector multiplication. Instead, we represent a system of multivariate nonlinear equations as a vector-valued function, for example, $f: \mathbb{R}^N \rightarrow \mathbb{R}^N$, that takes an N -dimensional vector and maps it to another N -dimensional vector. Multivariate systems of equations are much more complicated to solve than univariate equations, partly because there are so many more possible behaviors. As a consequence, no method strictly guarantees convergence to a solution, such as the bisection method for a univariate nonlinear equation, and the existing methods are much more computationally demanding than the univariate case, especially as the number of variables increases.

Not all methods discussed for univariate equation solving can be generalized to the multivariate case. For example, the bisection method cannot be directly generalized to a multivariate equation system. On the other hand, Newton's method can be used for multivariate problems. In this case, its iteration formula is $x_{k+1} = x_k - J_f(x_k)^{-1}f(x_k)$, where $J_f(x_k)$ is the Jacobian matrix of the $f(x)$ function, with elements $[J_f(x_k)]_{ij} = \partial f_i(x_k)/\partial x_j$. Instead of inverting the Jacobian matrix, it is sufficient to solve the linear equation system $J_f(x_k)\delta x_k = -f(x_k)$ and update x_k using $x_{k+1} = x_k + \delta x_k$. Like the secant variant of the Newton method for univariate equation systems, there are also variants of the multivariate method that avoid computing the Jacobian by estimating it from previous function evaluations. Broyden's method is a popular example of this secant updating method for multivariate equation systems. In the SciPy `optimize` module, `broyden1` and `broyden2` provide two implementations of Broyden's method using different approximations of the Jacobian and the `optimize.fsolve` function implements a Newton-like method, where optionally, the Jacobian can be specified, if available. The functions all have a similar function signature: The first argument is a Python function that represents the equation to be solved, and it should take a NumPy array as the first argument and return an array of the same shape. The second argument is an initial guess for the solution as a NumPy array. The `optimize.fsolve` function also takes an optional keyword argument `fprime`, which can be used to provide a function that returns the Jacobian of the $f(x)$ function. In addition, all these functions take numerous optional keyword arguments for tuning their behavior (see the docstrings for details).

For example, consider the following system of two multivariate and nonlinear equations

$$\begin{cases} y - x^3 - 2x^2 + 1 = 0 \\ y + x^2 - 1 = 0 \end{cases},$$

which can be represented by the vector-valued function $f([x_1, x_2]) = [x_2 - x_1^3 - 2x_1^2 + 1, x_2 + x_1^2 - 1]$. To solve this equation system using SciPy, we need to define a Python function for $f([x, x_2])$ and call, for example, the `optimize.fsolve` using the function and an initial guess for the solution vector.

```
In [73]: def f(x):
...:     return [x[1] - x[0]**3 - 2 * x[0]**2 + 1, x[1] + x[0]**2 - 1]
In [74]: optimize.fsolve(f, [1, 1])
Out[74]: array([ 0.73205081,  0.46410162])
```

The `optimize.broyden1` and `optimize.broyden2` can be used similarly. To specify a Jacobian for `optimize.fsolve`, we need to define a function that evaluates the Jacobian for a given input vector. This requires that we first derive the Jacobian by hand or, for example, using SymPy.

```
In [75]: x, y = sympy.symbols("x, y")
In [76]: f_mat = sympy.Matrix([y - x**3 - 2*x**2 + 1, y + x**2 - 1])
In [77]: f_mat.jacobian(sympy.Matrix([x, y]))

Out[77]:  $\begin{pmatrix} -3x^2 - 4x & 1 \\ 2x & 1 \end{pmatrix}$ 
```

which we can then easily implement as a Python function that can be passed to the `optimize.fsolve` function.

```
In [78]: def f_jacobian(x):
...:     return [-3*x[0]**2 - 4*x[0], 1], [2*x[0], 1]
In [79]: optimize.fsolve(f, [1, 1], fprime=f_jacobian)
Out[79]: array([ 0.73205081,  0.46410162])
```

As with Newton's method for a univariate nonlinear equation system, the initial guess for the solution is important, and different initial guesses may result in different solutions being found for the equations. There is no guarantee that any solution is found, although the proximity of the initial guess to the true solution is

often correlated with convergence to that particular solution. When possible, it is often a good approach to graph the equations being solved to visually indicate the number of solutions and their locations. For example, the following code demonstrates how three different solutions can be found to the equation systems we are considering here, by using different initial guesses with the `optimize.fsolve` function. The result is shown in Figure 5-8.

```
In [80]: def f(x):
...:     return [x[1] - x[0]**3 - 2 * x[0]**2 + 1,
...:             x[1] + x[0]**2 - 1]
...:
...: x = np.linspace(-3, 2, 5000)
...: y1 = x**3 + 2 * x**2 - 1
...: y2 = -x**2 + 1
...:
...: fig, ax = plt.subplots(figsize=(8, 4))
...:
...: ax.plot(x, y1, 'b', lw=1.5, label=r'$y = x^3 + 2x^2 - 1$')
...: ax.plot(x, y2, 'g', lw=1.5, label=r'$y = -x^2 + 1$')
...:
...: x_guesses = [[-2, 2], [1, -1], [-2, -5]]
...: for x_guess in x_guesses:
...:     sol = optimize.fsolve(f, x_guess)
...:     ax.plot(sol[0], sol[1], 'r*', markersize=15)
...:
...:     ax.plot(x_guess[0], x_guess[1], 'ko')
...:     ax.annotate("", xy=(sol[0], sol[1]), xytext=(x_guess[0], x_guess[1]),
...:                 arrowprops=dict(arrowstyle="->", linewidth=2.5))
...:
...: ax.legend(loc=0)
...: ax.set_xlabel(r'$x$', fontsize=18)
```




Figure 5-8. Graph of a system of two nonlinear equations. The solutions are indicated with red stars, and the initial guess with a black dot and an arrow to the solution each initial guess eventually converged to

By systematically solving the equation systems with different initial guesses, we can build a visualization of how different initial guesses converge to different solutions. This is done in the following code example, and the result is shown in Figure 5-9. This example demonstrates that even for this relatively simple example, the regions of initial guesses that converge to different solutions are highly nontrivial, and there are also missing dots corresponding to initial guesses for which the algorithm fails to converge to any solution. Nonlinear equation solving is a complex task, and visualizations of different types can often be a valuable tool when discovering a particular problem's characteristics.

```
In [81]: fig, ax = plt.subplots(figsize=(8, 4))
...:
...: ax.plot(x, y1, 'k', lw=1.5)
...: ax.plot(x, y2, 'k', lw=1.5)
...:
...: sol1 = optimize.fsolve(f, [-2, 2])
...: sol2 = optimize.fsolve(f, [ 1, -1])
...: sol3 = optimize.fsolve(f, [-2, -5])
...: sols = [sol1, sol2, sol3]
...: colors = ['r', 'b', 'g']
...: for idx, s in enumerate(sols):
...:     ax.plot(s[0], s[1], colors[idx]+'*', markersize=15)
...:
...: for m in np.linspace(-4, 3, 80):
...:     for n in np.linspace(-15, 15, 40):
...:         x_guess = [m, n]
...:         sol = optimize.fsolve(f, x_guess)
...:         idx = (abs(sols - sol)**2).sum(axis=1).argmin()
...:         ax.plot(x_guess[0], x_guess[1], colors[idx]+'.')
...:
...: ax.set_xlabel(r'$x$', fontsize=18)
```




Figure 5-9. Visualization of the convergence to different solutions for different initial guesses. Each dot represents an initial guess, and its color encodes which solution the solver eventually converges to. The solutions are marked with correspondingly color-coded stars

Summary

This chapter explored methods for solving algebraic equations using the SymPy and SciPy libraries. Equation solving is one of the most elementary mathematical tools for computational sciences, and it is an important component in many algorithms and methods and has direct applications in many problem-solving situations. In some cases, analytical algebraic solutions exist, especially for equations that are polynomials or contain certain combinations of elementary functions. Such equations can often be handled symbolically with SymPy. Numerical methods are usually the only feasible approach for equations with no algebraic solution and larger systems of equations. Linear equation systems can always be systematically solved. For this reason, there are many important applications for linear equation systems, be it for originally linear systems or as approximations to originally nonlinear systems. Nonlinear equation solving requires a different set of methods, and it is generally much more complex and computationally demanding compared to linear equation systems. Solving linear equation systems is an important step in the iterative methods employed in many of the methods that exist to solve nonlinear equation systems. For numerical equation solving, we can use the linear algebra and optimization modules in SciPy, which provide efficient and well-tested methods for numerical root finding and equation solving of both linear and nonlinear systems.

Further Reading

Equation solving is a basic numerical technique whose methods are covered in most introductory numerical analysis texts. A good example of a book that covers these topics is *Scientific Computing* by M. Heath (McGraw-Hill, 2001), and *Numerical Recipes: The Art of Scientific Computing* by W. H. Press (Cambridge University Press, 2007) gives a practical introduction with implementation details.