

CHAPTER 17



Signal Processing

This chapter explores signal processing, a subject with applications in diverse branches of science and engineering. A signal in this context can be a quantity that varies in time (temporal signal) or as a function of space coordinates (spatial signal). An audio signal is a typical example of a temporal signal, while an image is a typical example of a spatial signal in two dimensions. In practice, signals are often continuous functions. But it is common in computational applications to work with discretized signals, where the original continuous signal is sampled at discrete points with uniform distances. The sampling theorem gives rigorous and quantitative conditions for when a discrete sequence of samples can accurately represent a continuous signal.

Signal processing is essential in scientific computing because of its broad applicability and because there are very efficient computational methods for fundamental and highly practical problems. In particular, the fast Fourier transform (FFT) is an important algorithm for many signal-processing problems. Moreover, it is one of the most important numerical algorithms in all computing. This chapter examines how FFTs can be used in spectral analysis. But beyond this basic application, there is also broad usage of FFT both directly and indirectly as a component in other algorithms. Other signal-processing methods, such as convolution and correlation analysis and linear filters, have widespread applications, particularly in engineering fields such as control theory.

The chapter discusses spectral analysis and basic applications of linear filters using the SciPy library.

Importing Modules

This chapter mainly works with the `fftpack` and `signal` modules from the SciPy library. As usual with modules from the SciPy library, we import the modules using the following pattern.

```
In [1]: from scipy import fftpack
In [2]: from scipy import signal
```

We also use the `io.wavfile` module from SciPy to read and write WAV audio files in one of the examples. We import this module in the following way.

```
In [3]: import scipy.io.wavfile
In [4]: from scipy import io
```

For basic numerics and graphics, we also require the NumPy, Pandas, and Matplotlib libraries.

```
In [5]: import numpy as np
In [6]: import pandas as pd
In [7]: import matplotlib.pyplot as plt
In [8]: import matplotlib as mpl
```

Spectral Analysis

Let's begin this exploration of signal processing by considering spectral analysis. Spectral analysis is a fundamental application of Fourier transforms, a mathematical integral transform that allows us to take a signal from the time domain—which is described as a function of time—to the frequency domain, which is described as a function of frequency. The frequency-domain representation of a signal is useful for many purposes, for example, extracting features such as dominant frequency components of a signal, applying filters to signals, and solving differential equations (see Chapter 9), to mention a few.

Fourier Transforms

The following is the mathematical expression for the Fourier transform $F(v)$ of a continuous signal $f(t)$:¹

$$F(v) = \int_{-\infty}^{\infty} f(t) e^{-2\pi i v t} dt,$$

and the inverse Fourier transform is given by the following:

$$f(t) = \int_{-\infty}^{\infty} F(v) e^{2\pi i v t} dv.$$

Here $F(v)$ is the complex-valued amplitude spectrum of the signal $f(t)$, and v is the frequency. From $F(v)$, we can compute other types of spectrums, such as the power spectrum $|F(v)|^2$. In this formulation, $f(t)$ is a continuous signal with infinite duration. In practical applications, we are often more interested in approximating $f(t)$ function using a finite number of samples for a finite time duration. For example, we might sample the $f(t)$ at N uniformly spaced points in the time interval $t \in [0, T]$, resulting in a sequence of samples that we denote (x_0, x_1, \dots, x_N) . The continuous Fourier transform shown in the preceding text can be adapted to the discrete case: the discrete Fourier transform (DFT) of a sequence of uniformly spaced samples is as follows:

$$X_k = \sum_{n=0}^{N-1} x_n e^{-2\pi i n k / N}.$$

Similarly, we have the inverse DFT:

$$x_n = \frac{1}{N} \sum_{k=0}^{N-1} X_k e^{2\pi i n k / N},$$

where X_k is the discrete Fourier transform of the samples x_n , and k is a frequency bin number that can be related to an actual frequency. The DFT for a sequence of samples can be computed very efficiently using the fast Fourier transform (FFT) algorithm. The SciPy `fftpack` module² provides implementations of the FFT algorithm. The `fftpack` module contains FFT functions for various cases. Here, we focus on demonstrating the usage of the `fft` and `ifft` functions and several of the helper functions in the `fftpack` module. However, the general usage is similar for all FFT functions in Table 17-1.

¹ There are several alternative definitions of the Fourier transform, which vary in the coefficient in the exponent and the normalization of the transform integral.

² FFT is also implemented in the `fft` module in NumPy. It provides mostly the same functions as `scipy.fftpack`, which we use here. As a rule, when SciPy and NumPy provide the same functionality, it is generally preferable to use SciPy if available and fall back to the NumPy implementation when SciPy is unavailable.

Table 17-1. Summary of Selected Functions from the *fftpack* Module in SciPy

Function	Description
<code>fft</code> , <code>ifft</code>	General FFT and inverse FFT of a real- or complex-valued signal. The resulting frequency spectrum is complex valued.
<code>rfft</code> , <code>irfft</code>	The FFT and inverse FFT of a real-valued signal.
<code>dct</code> , <code>idct</code>	The discrete cosine transform (DCT) and its inverse.
<code>dst</code> , <code>idst</code>	The discrete sine transform (DST) and its inverse.
<code>fft2</code> , <code>ifft2</code> , <code>fftn</code> , <code>ifftn</code>	The two-dimensional and the N-dimensional FFT for complex-valued signals and their inverses.
<code>fftshift</code> , <code>ifftshift</code> , <code>rfftshift</code> , <code>irfftshift</code>	Shift the frequency bins in the result vector produced by <code>fft</code> and <code>rfft</code> , respectively, so that the spectrum is arranged such that the zero-frequency component is in the middle of the array.
<code>fftfreq</code>	Calculate the frequencies corresponding to the FFT bins in the result returned by <code>fft</code> .

For detailed usage of each function, including their arguments and return values, see their docstrings, which are available using, for example, `help(fftpack.fft)`.

The DFT takes discrete samples as input and outputs a discrete frequency spectrum. To be able to use DFT for processes that are originally continuous, we first must reduce the signals to discrete values using sampling. According to the sampling theorem, a continuous signal with bandwidth B (i.e., the signal does not contain frequencies higher than B) can be reconstructed entirely from discrete samples with sampling frequency $f_s \geq 2B$. This is a significant result in signal processing because it tells us under what circumstances we can work with discrete instead of continuous signals. It allows us to determine a suitable sampling rate when measuring a continuous process since it is often possible to know or approximately guess the bandwidth of a process, for example, from physical arguments. While the sampling rate determines the maximum frequency we can describe with a discrete Fourier transform, the spacing of samples in frequency space is determined by the total sampling time T or equivalently from the number of sample points once the sampling frequency is determined, $T = N/f_s$.

As an introductory example, consider a simulated signal with pure sinusoidal components at 1 Hz and 22 Hz on top of a normal-distributed noise floor. We begin by defining a `signal_samples` function that generates noisy samples of this signal.

```
In [9]: def signal_samples(t):
...:     return (2 * np.sin(2 * np.pi * t) + 3 * np.sin(22 * 2 * np.pi * t) +
...:             2 * np.random.randn(*np.shape(t)))
```

We can get a vector of samples by calling this function with an array with sample times as an argument. Say that we are interested in computing the frequency spectrum of this signal up to frequencies of 30 Hz. We then need to choose the sampling frequency $f_s = 60$ Hz, and if we want to obtain a frequency spectrum with a resolution of $\Delta f = 0.01$ Hz, we need to collect at least $N = f_s / \Delta f = 6000$ samples, corresponding to a sampling period of $T = N/f_s = 100$ seconds.

```
In [10]: B = 30.0
In [11]: f_s = 2 * B
In [12]: delta_f = 0.01
In [13]: N = int(f_s / delta_f); N
```

```
Out[13]: 6000
In [14]: T = N / f_s; T
Out[14]: 100.0
```

Next, we sample the signal function at N uniformly spaced points in time by creating a t array containing the sample times and then using it to evaluate the `signal_samples` function.

```
In [15]: t = np.linspace(0, T, N)
In [16]: f_t = signal_samples(t)
```

The resulting signal is plotted in Figure 17-1. The signal is rather noisy when viewed over the entire sampling time and for a shorter period, and the added random noise mostly masks the pure sinusoidal signals when viewed in the time domain.

```
In [17]: fig, axes = plt.subplots(1, 2, figsize=(8, 3), sharey=True)
...: axes[0].plot(t, f_t)
...: axes[0].set_xlabel("time (s)")
...: axes[0].set_ylabel("signal")
...: axes[1].plot(t, f_t)
...: axes[1].set_xlim(0, 5)
...: axes[1].set_xlabel("time (s)")
```




Figure 17-1. Simulated signal with random noise. Full signal to the left and zoom into early times on the right

To reveal the sinusoidal components in the signal, we can use the FFT to compute the signal's spectrum (or, in other words, its frequency-domain representation). We obtain the discrete Fourier transform of the signal by applying the `fft` function to the array of discrete samples, f_t .

```
In [18]: F = fftpack.fft(f_t)
```

The result is an `F` array, which contains the spectrum's frequency components at frequencies determined by the sampling rate and the number of samples. When computing these frequencies, it is convenient to use the `fftfreq` helper function, which takes the number of samples and the time duration between successive samples as parameters and returns an array of the same size as `F` that contains the frequencies corresponding to each frequency bin.

```
In [19]: f = fftpack.fftfreq(N, 1.0/f_s)
```

The frequency bins for the amplitude values returned by the `fft` function contain both positive and negative frequencies up to the frequency corresponding to half the sampling rate, $f_s/2$. For real-valued signals, the spectrum is symmetric at positive and negative frequencies, and we are, for this reason, often only interested in the positive-frequency components. Using the frequency array `f`, we can conveniently create a mask that can be used to extract the part of the spectrum that corresponds to the frequencies we are interested in. Here, we create a mask for selecting the positive-frequency components.

```
In [20]: mask = np.where(f >= 0)
```

The spectrum for the positive-frequency components is shown in Figure 17-2. The top panel contains the entire positive-frequency spectrum and is plotted on a log scale to increase the contrast between the signal and the noise. We can see sharp peaks near 1 Hz and 22 Hz, corresponding to the sinusoidal components in the signal. These peaks stand out from the noise floor in the spectrum. Despite the noise concealing the sinusoidal components in the time-domain signal, we can detect their presence in the frequency-domain representation. The two lower panels in Figure 17-2 show magnifications of the two peaks at 1 Hz and 22 Hz, respectively.

```
In [21]: fig, axes = plt.subplots(3, 1, figsize=(8, 6))
...: axes[0].plot(f[mask], np.log(abs(F[mask])))
...: axes[0].plot(B, 0, 'r*', markersize=10)
...: axes[0].set_ylabel("$\log(|F|)$", fontsize=14)
...: axes[1].plot(f[mask], abs(F[mask])/N)
...: axes[1].set_xlim(0, 2)
...: axes[1].set_ylabel("$|F|/N$", fontsize=14)
...: axes[2].plot(f[mask], abs(F[mask])/N)
...: axes[2].set_xlim(21, 23)
...: axes[2].set_xlabel("frequency (Hz)", fontsize=14)
...: axes[2].set_ylabel("$|F|/N$", fontsize=14)
```




Figure 17-2. Spectrum of the simulated signal with frequency components at 1 Hz and 22 Hz

Frequency-Domain Filter

Just like we can compute the frequency-domain representation from the time-domain signal using the FFT function `fft`, we can compute the time-domain signal from the frequency-domain representation using the *inverse* FFT function `ifft`. For example, applying the `ifft` function to the `F` array reconstructs the `f_t` array. Modifying the spectrum before applying the inverse transform allows us to realize frequency-domain filters. For example, selecting only frequencies below 2 Hz in the spectrum amounts to applying a 2 Hz low-pass filter, which suppresses high-frequency components in the signal (higher than 2 Hz in this case).

```
In [22]: F_filtered = F * (abs(f) < 2)
In [23]: f_t_filtered = fftpack.ifft(F_filtered)
```

Computing the inverse FFT for the filtered signal results in a time-domain signal where the high-frequency oscillations are absent, as shown in Figure 17-3. This simple example summarizes the essence of many frequency-domain filters. Some of the many types of filters commonly used in signal-processing analysis are covered later in this chapter.

```
In [24]: fig, ax = plt.subplots(figsize=(8, 3))
...: ax.plot(t, f_t, label='original')
...: ax.plot(t, f_t_filtered.real, color="red", lw=3, label='filtered')
...: ax.set_xlim(0, 10)
...: ax.set_xlabel("time (s)")
...: ax.set_ylabel("signal")
...: ax.legend()
```




Figure 17-3. The original time-domain signal and the reconstructed signal after applying a low-pass filter to the frequency-domain representation of the signal

Windowing

In the previous section, we directly applied the FFT to the signal. This can give acceptable results, but it is often possible to further improve the frequency spectrum's quality and contrast by applying a *window function* to the signal before applying the FFT. A window function is a function that, when multiplied with the signal, modulates its magnitude so that it approaches zero at the beginning and the end of the sampling duration. There are many possible window functions, and the SciPy signal module provides implementations of many common window functions, including the Blackman function, the Hann function, the Hamming function, Gaussian window functions (with variable standard deviation), and the Kaiser window function.³ These functions are all plotted in Figure 17-4. This graph shows that while all these window functions are slightly different, the overall shape is very similar.

```
In [25]: fig, ax = plt.subplots(1, 1, figsize=(8, 3))
...: N = 100
...: ax.plot(signal.blackman(N), label="Blackman")
...: ax.plot(signal.hann(N), label="Hann")
...: ax.plot(signal.hamming(N), label="Hamming")
...: ax.plot(signal.gaussian(N, N/5), label="Gaussian (std=N/5)")
...: ax.plot(signal.kaiser(N, 7), label="Kaiser (beta=7)")
...: ax.set_xlabel("n")
...: ax.legend(loc=0)
```

³Several other window functions are also available. See the docstring for the `scipy.signal` module for a complete list.




Figure 17-4. Examples of commonly used window functions

The alternative window functions all have slightly different properties and objectives, but for the most part, they can be used interchangeably. The main purpose of window functions is to reduce spectral leakage between nearby frequency bins, which occurs in discrete Fourier transform computation when the signal contains components with periods that are not exactly divisible with the sampling period. Signal components with such frequencies can, therefore, not fit a full number of cycles in the sampling period. Since the discrete Fourier transform assumes that the signal is periodic, the resulting discontinuity at the period boundary can give rise to spectral leakage. Multiplying the signal with a window function reduces this problem. Alternatively, we could increase the sample points (increase the sampling period) to obtain a higher frequency resolution, but this might not always be practical.

To see how we can use a window function before applying the FFT to a time-series signal, let's consider the outdoor temperature measurements that we looked at in Chapter 12. First, we use the Pandas library to load the dataset and resample it to evenly spaced hourly samples, using the `fillna` method to aggregate the elements.

```
In [26]: df = pd.read_csv('temperature_outdoor_2014.tsv', delimiter="\t",
...:                     names=["time", "temperature"])
In [27]: df.time = (pd.to_datetime(df.time.values, unit="s").
...:               tz_localize('UTC').tz_convert('Europe/Stockholm'))
In [28]: df = df.set_index("time")
In [29]: df = df.resample("H").ffill()
In [30]: df = df[(df.index >= "2014-04-01")*(df.index < "2014-06-01")].dropna()
```

Once the Pandas data frame has been created and processed, we extract the underlying NumPy arrays to be able to process the time-series data using the `fftpack` module.

```
In [31]: time = df.index.astype('int64')/1.0e9
In [32]: temperature = df.temperature.values
```

Now, we wish to apply a window function to the data in the array `temperature` before we compute the FFT. Here, we use the Blackman window function, which is available as the `blackman` function in the `signal` module in SciPy. As an argument to the window function, we need to pass the length of the sample array, and it returns an array of that same length.

```
In [33]: window = signal.blackman(len(temperature))
```


To apply the window function, we multiply it with the array containing the time-domain signal and use the result in the subsequent FFT computation. However, before we proceed with the FFT for the windowed temperature signal, we first plot the original temperature time series and the windowed version. The result is shown in Figure 17-5. The result of multiplying the time series with the window function is a signal approaching zero near the sampling period boundaries. It can be viewed as a periodic function with smooth transitions between period boundaries, and as such, the FFT of the windowed signal has better-behaved properties.

```
In [34]: temperature_windowed = temperature * window
In [35]: fig, ax = plt.subplots(figsize=(8, 3))
...: ax.plot(df.index, temperature, label="original")
...: ax.plot(df.index, temperature_windowed, label="windowed")
...: ax.set_ylabel("temperature", fontsize=14)
...: ax.legend(loc=0)
```




Figure 17-5. Windowed and original temperature time-series signal

After preparing the windowed signal, the rest of the spectral analysis proceeds as before: we can use the `fft` function to compute the spectrum and the `fftfreq` function to calculate the frequencies corresponding to each frequency bin.

```
In [36]: data_fft_windowed = fftpack.fft(temperature_windowed)
In [37]: f = fftpack.fftfreq(len(temperature), time[1]-time[0])
```

Here, we also select the positive frequencies by creating a mask array from the array `f` and plot the resulting positive-frequency spectrum as shown in Figure 17-6. It shows peaks at the frequency corresponding to 1 day (1/86,400 Hz) and its higher harmonics (2/86,400 Hz, 3/86,400 Hz, etc.).

```
In [38]: mask = f > 0
In [39]: fig, ax = plt.subplots(figsize=(8, 3))
...: ax.set_xlim(0.000005, 0.00004)
...: ax.axvline(1./86400, color='r', lw=0.5)
...: ax.axvline(2./86400, color='r', lw=0.5)
...: ax.axvline(3./86400, color='r', lw=0.5)
...: ax.plot(f[mask], np.log(abs(data_fft_windowed[mask])), lw=2)
...: ax.set_ylabel("$\\log|F|$", fontsize=14)
...: ax.set_xlabel("frequency (Hz)", fontsize=14)
```




Figure 17-6. Spectrum of the windowed temperature time series. The dominant peak occurs at the frequency corresponding to a 1-day period and its higher harmonics

To get the most accurate spectrum from a given set of samples, applying a window function to the time-series signal is generally advisable before applying an FFT. Most of the window functions available in SciPy can be used interchangeably, and the choice of window function is usually not critical. A popular choice is the Blackman window function, which is designed to minimize spectral leakage. For more information about the properties of different window functions, see Chapter 9 of *The Scientist and Engineer's Guide to Digital Signal Processing* by S. Smith.

Spectrogram

As a final example in this section on spectral analysis, let's analyze the spectrum of an audio signal sampled from a guitar.⁴ First, we load sampled data from the `guitar.wav` file using the `io.wavfile.read` function from the SciPy library.

```
In [40]: sample_rate, data = io.wavfile.read("guitar.wav")
```

The `io.wavfile.read` function returns a tuple containing the sampling rate, `sample_rate`, and a NumPy array containing the audio intensity. For this particular file, we get the sampling rate of 44.1 kHz, and the audio signal was recorded in stereo, which is represented by a data array with two channels. Each channel contains 1,181,625 samples.

```
In [41]: sample_rate
Out[41]: 44100
In [42]: data.shape
Out[42]: (1181625, 2)
```

Here, we are only concerned with analyzing a single audio channel, so we form the average of the two channels to obtain a mono-channel signal.

```
In [43]: data = data.mean(axis=1)
```

⁴The data used in this example was obtained from <https://www.freesound.org/people/guitarguy1985/sounds/52047>.

We can calculate the total duration of the audio recording by dividing the number of samples by the sampling rate. The result suggests that the recording is about 26.8 seconds.

```
In [44]: data.shape[0] / sample_rate
Out[44]: 26.79421768707483
```

It is often the case that we like to compute the spectrum of a signal in segments instead of the entire signal at once, for example, if the nature of the signal varies in time on a long timescale but contains nearly periodic components on a short timescale. This is particularly true for music, which can be considered nearly periodic on short timescales from the point of view of human perception (subsecond timescales) but varies on longer timescales. In the guitar sample, we would like to apply the FFT on a sliding window in the time-domain signal. The result is a time-dependent spectrum, often visualized as an equalizer graph on music equipment and applications. Another approach is to visualize the time-dependent spectrum using a two-dimensional heatmap graph, which in this context is known as a spectrogram. The following computes the spectrogram of the guitar sample.

Before proceeding with the spectrogram visualization, we calculate the spectrum for a small part of the sample. Let's begin by determining the number of samples from the full array. If we want to analyze 0.5 seconds at the time, we can use the sampling rate to compute the number of samples to use.

```
In [45]: N = int(sample_rate/2.0) # half a second -> 22050 samples
```

Next, given the number of samples and the sampling rate, we can compute the frequencies f for the frequency bins for the result of the forthcoming FFT calculation and the sampling times t for each sample in the time-domain signal. We also create a frequency mask for selecting positive frequencies smaller than 1000 Hz, which is used later to select a subset of the computed spectrum.

```
In [46]: f = fftpack.fftfreq(N, 1.0/sample_rate)
In [47]: t = np.linspace(0, 0.5, N)
In [48]: mask = (f > 0) * (f < 1000)
```

Next, we extract the first N samples from the full sample array `data` and apply the `fft` function.

```
In [49]: subdata = data[:N]
In [50]: F = fftpack.fft(subdata)
```

The time- and frequency-domain signals are shown in Figure 17-7. The time-domain signal in the left panel is zero in the beginning before the first guitar string is plucked. The frequency-domain spectrum shows several dominant frequencies corresponding to the different tones produced by the guitar.

```
In [51]: fig, axes = plt.subplots(1, 2, figsize=(12, 3))
...: axes[0].plot(t, subdata)
...: axes[0].set_ylabel("signal", fontsize=14)
...: axes[0].set_xlabel("time (s)", fontsize=14)
...: axes[1].plot(f[mask], abs(F[mask]))
...: axes[1].set_ylabel("$|F|$", fontsize=14)
...: axes[1].set_xlabel("Frequency (Hz)", fontsize=14)
```




Figure 17-7. Signal and spectrum for samples half a second duration of a guitar sound

The next step is to repeat the analysis for successive segments from the full sample array. The time evolution of the spectrum can be visualized as a spectrogram, with frequency on the x axis and time on the y axis. To be able to plot the spectrogram with the `imshow` function from Matplotlib, we create a two-dimensional NumPy array `spectrogram_data` for storing the spectra for the successive sample segments. The shape of the `spectrogram_data` array is (n_max, f_values) , where n_max is the number of segments of length N in the sample array data, and f_values are the number of frequency bins with frequencies that match the condition used to compute `mask` (positive frequencies less than 1000 Hz).

```
In [52]: n_max = int(data.shape[0] / N)
In [53]: f_values = np.sum(mask)
In [54]: spectrogram_data = np.zeros((n_max, f_values))
```

To improve the resulting spectrogram's contrast, we apply a Blackman window function to each subset of the sample data before we compute the FFT. Here, we choose the Blackman window function for its spectral leakage-reducing properties, but many other window functions give similar results. The length of the window array must be the same as the length of the subdata array, so we pass its length argument to the Blackman function.

```
In [55]: window = signal.blackman(len(subdata))
```

Finally, we can compute the spectrum for each segment in the sample by looping over the array slices of size N , apply the window function, compute the FFT, and store the subset of the result for the frequencies we are interested in in the `spectrogram_data` array.

```
In [56]: for n in range(0, n_max):
...:     subdata = data[(N * n):(N * (n + 1))]
...:     F = fftpack.fft(subdata * window)
...:     spectrogram_data[n, :] = np.log(abs(F[mask]))
```

When the `spectrogram_data` array is computed, we can visualize the spectrogram using the `imshow` function from Matplotlib. The result is shown in Figure 17-8.

```
In [57]: fig, ax = plt.subplots(1, 1, figsize=(8, 6))
...: p = ax.imshow(spectrogram_data, origin='lower',
...:               extent=(0, 1000, 0, data.shape[0] / sample_rate),
...:               aspect='auto',
...:               cmap=matplotlib.cm.RdBu_r)
...: cb = fig.colorbar(p, ax=ax)
```

```
...: cb.set_label("\log|F|$", fontsize=14)
...: ax.set_ylabel("time (s)", fontsize=14)
...: ax.set_xlabel("Frequency (Hz)", fontsize=14)
```

The spectrogram in Figure 17-8 contains a lot of information about the sampled signal and how it evolves. The narrow vertical stripes correspond to tones produced by the guitar, and those signals slowly decay over time. The broad horizontal bands correspond roughly to periods when strings are being plucked on the guitar, giving a very broad frequency response for a short time. Note, however, that the color axis represents a logarithmic scale, so small variations in the color represent considerable variation in the actual intensity.




Figure 17-8. Spectrogram of an audio sampling of a guitar sound

Signal Filters

One of the main objectives of signal processing is to manipulate and transform temporal or spatial signals to change their characteristics. Typical applications are noise reduction, sound effects in audio signals, and effects such as blurring, sharpening, contrast enhancement, and color balance adjustments in image data. Many common transformations can be implemented as filters that act on the frequency/domain representation of the signal, for example, by suppressing specific frequency components. The previous section presented an example of a low-pass filter, which we implemented by taking the Fourier transform of the signal, removing the high-frequency components, and finally taking the inverse Fourier transform

to obtain a new time-domain signal. With this approach, we can implement arbitrary frequency filters, but we cannot necessarily apply them in real-time on a streaming signal since they require buffering sufficient samples to be able to perform the discrete Fourier transform. In many applications, it is desirable to apply filters and continuously transform a signal, for example, when processing signals in transmission or live audio signals.

Convolution Filters

Certain frequency filters can be implemented directly in the time domain using a convolution of the signal with a function that characterizes the filter. An important property of Fourier transformations is that the (inverse) Fourier transform of the product of two functions (e.g., the spectrum of a signal and the filter shape function) is a convolution of the two functions (inverse) Fourier transforms. Therefore, if we want to apply a filter, H_k , to the spectrum X_k of a signal x_n , we can instead compute the convolution of x_n with h_m , the inverse Fourier transform of the H_k filter function. In general, we can write a filter on convolution form as follows:

$$y_n = \sum_{k=-\infty}^{\infty} x_k h_{n-k},$$

where x_k is the input, y_n is the output, and h_{n-k} is the convolution kernel that characterizes the filter. Note that in this general form, the signal y_n at time step n depends on both earlier and later values of the input x_k . To illustrate this point, let's return to the first example in this chapter, where we applied a low-pass filter to a simulated signal with components at 1 Hz and 22 Hz. That example Fourier transformed the signal and multiplied its spectrum with a step function that suppressed all high-frequency components. Finally, we inverse Fourier transformed the signal back into the time domain. The result was a smoothed version of the original noisy signal (see Figure 17-3). An alternative approach using convolution is to inverse Fourier transform the frequency response function for the filter H and use the result h as a kernel with which we convolve the original time-domain signal f_t .

```
In [58]: t = np.linspace(0, T, N)
In [59]: f_t = signal_samples(t)
In [60]: H = abs(f) < 2
In [61]: h = fftpack.fftfreq(N, T).real
In [62]: f_t_filtered_conv = signal.convolve(f_t, h, mode='same')
```

To carry out the convolution, the `convolve` function from the `signal` module in SciPy was used. It takes as arguments two NumPy arrays containing the signals for which to compute the convolution. Using the optional keyword argument `mode`, we can set the size of the output array to be the same as the first input (`mode='same'`), the full convolution output after having zero-padded the arrays to account for transients (`mode='full'`), or to contain only elements that do not rely on zero-padding (`mode='valid'`). Here, we use `mode='same'`, to easily compare and plot the result with the original signal, `f_t`. The result of applying this convolution filter, `f_t_filtered_conv`, is shown in Figure 17-9, together with the corresponding result computed using `fft` and `ifft` with a modified spectrum (`f_t_filtered`). As expected, the two methods give identical results.

```
In [63]: fig = plt.figure(figsize=(8, 6))
...: ax = plt.subplot2grid((2,2), (0,0))
...: ax.plot(f, H)
...: ax.set_xlabel("frequency (Hz)")
...: ax.set_ylabel("Frequency filter")
...: ax.set_ylim(0, 1.5)
```

```

...: ax = plt.subplot2grid((2,2), (0,1))
...: ax.plot(t - t[-1]/2.0, h.real)
...: ax.set_xlabel("time (s)")
...: ax.set_ylabel("convolution kernel")
...: ax = plt.subplot2grid((2,2), (1,0), colspan=2)
...: ax.plot(t, f_t, label='original', alpha=0.25)
...: ax.plot(t, f_t_filtered.real, 'r', lw=2,
...:         label='filtered in frequency domain')
...: ax.plot(t, f_t_filtered_conv.real, 'b--', lw=2,
...:         label='filtered with convolution')
...: ax.set_xlim(0, 10)
...: ax.set_xlabel("time (s)")
...: ax.set_ylabel("signal")
...: ax.legend(loc=2)

```

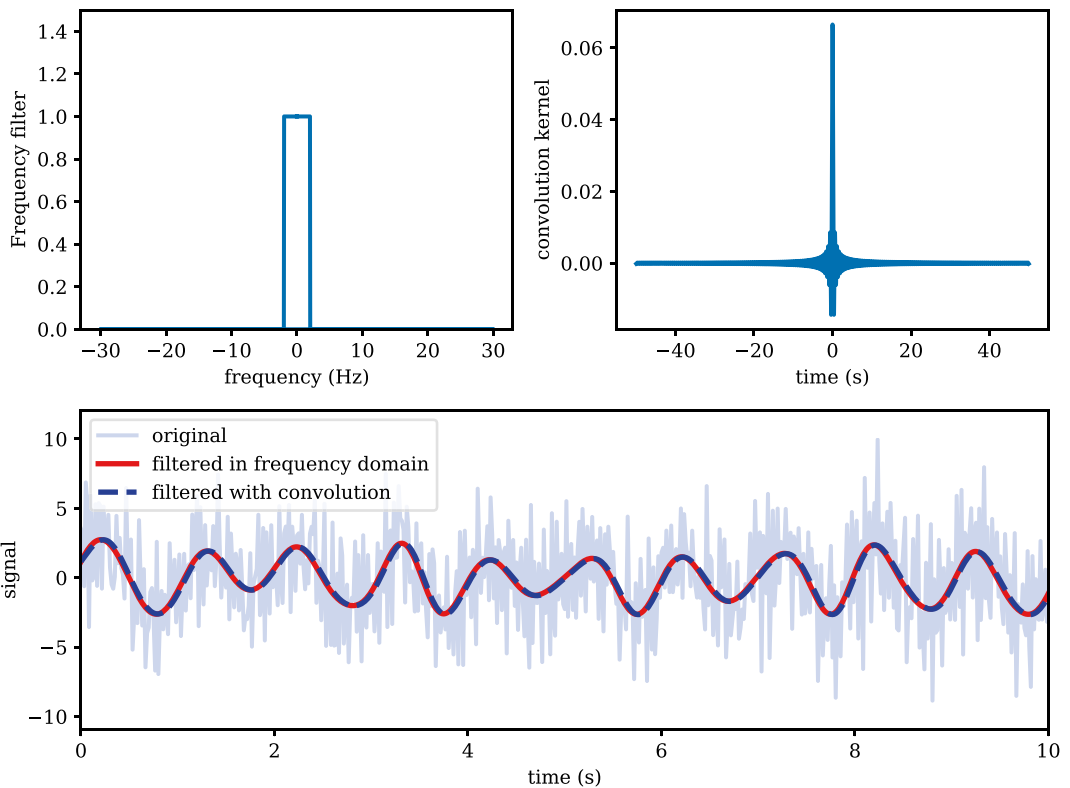


Figure 17-9. Top left: frequency filter. Top right: convolution kernel corresponding to the frequency filter (its inverse discrete Fourier transform). Bottom: simple lowpass filter applied via convolution

FIR and IIR Filters

In the convolution filter example, there was no computational advantage for using convolution to implement the filter rather than a sequence of a call to `fft`, spectrum modifications, followed by a call to `ifft`. The convolution here is, in general, more computationally expensive than the extra FFT transformation, and the SciPy signal module even provides a `fftconvolve` function, which implements the convolution using FFT and its inverse. Furthermore, the convolution kernel of the filter has many undesirable properties, such as being noncasual, where the output signal depends on future input values (see the upper-right panel in Figure 17-9). However, important special cases of convolution-like filters can be efficiently implemented with dedicated digital signal processors (DSPs) and general-purpose processors.

An important family of such filters is the *finite impulse response* (FIR) filters, which take the form $y_n = \sum_{k=0}^M b_k x_{n-k}$.

This time-domain filter is casual because the output y_n only depends on input values at earlier time steps.

Another similar type of filter is the *infinite impulse response* (IIR) filter, which can be written in the form

$a_0 y_n = \sum_{k=0}^M b_k x_{n-k} - \sum_{k=1}^N a_k y_{n-k}$. This is not strictly a convolution since it additionally includes past values of the output when computing a new output value (a feedback term). But it is nonetheless in a similar form. Both FIR and IIR filters can be used to evaluate new output values given the recent history of the signal and the output and can, therefore, be evaluated sequentially in the time domain if we know the finite sequences of values of b_k and a_k .

Computing the values of b_k and a_k given a set of requirements on filter properties is known as filter design. The SciPy signal module provides many functions for this purpose. For example, using the `firwin` function, we can compute the b_k coefficients for an FIR filter given frequencies of the band boundaries, where, for example, the filter transitions from a pass to a stop filter (for a low-pass filter). The `firwin` function takes the number of values in the a_k sequence as the first argument (also known as *taps* in this context). The second argument, `cutoff`, defines the low-pass transition frequency in units of the Nyquist frequency (half the sampling rate). The scale of the Nyquist frequency can optionally be set using the `nyq` argument, which defaults to 1. Finally, we can specify the type of window function to use with the `window` argument.

```
In [64]: n = 101
In [65]: f_s = 1 / 3600
In [66]: nyq = f_s/2
In [67]: b = signal.firwin(n, cutoff=nyq/12, nyq=nyq, window="hamming")
```

The result is the sequence of coefficients b_k that defines an FIR filter and can be used to implement the filter with a time-domain convolution. Given the coefficients b_k , we can evaluate the amplitude and phase response of the filter using the `freqz` function from the signal module. It returns arrays containing frequencies and the corresponding complex-valued frequency response, which are suitable for plotting purposes, as shown in Figure 17-10.

```
In [68]: f, h = signal.freqz(b)
In [69]: fig, ax = plt.subplots(1, 1, figsize=(12, 3))
...: h_ampl = 20 * np.log10(abs(h))
...: h_phase = np.unwrap(np.angle(h))
...: ax.plot(f/max(f), h_ampl, 'b')
...: ax.set_ylim(-150, 5)
```



```

...: ax.set_ylabel('frequency response (dB)', color="b")
...: ax.set_xlabel(r'normalized frequency')
...: ax = ax.twinx()
...: ax.plot(f/max(f), h_phase, 'r')
...: ax.set_ylabel('phase response', color="r")
...: ax.axvline(1.0/12, color="black")

```




Figure 17-10. The amplitude and phase response of a low-pass FIR filter

The low-pass filter shown in Figure 17-10 is designed to pass through signals with frequencies less than $f_s/24$ (indicated with a vertical line) and suppress higher-frequency signal components. The finite transition region between pass and stop bands and the nonperfect suppression above the cutoff frequency is a price we have to pay to be able to represent the filter in FIR form. The accuracy of the FIR filter can be improved by increasing the number of coefficients b_k , at the expense of higher computational complexity.

The effect of an FIR filter, given the coefficients b_k , and an IIR filter, given the coefficients b_k and a_k , can be evaluated using the `lfiltfilt` function from the `signal` module. As the first argument, this function expects the array with coefficients b_k , and as the second argument, the array with the coefficients a_k in an IIR filter or the scalar 1 in an FIR filter. The third argument to the function is the input signal array, and the return value is the filter output. For example, we can use the following to apply the FIR filter we created in the preceding text to the array with hourly temperature measurements `temperature`.

```
In [70]: temperature_filt = signal.lfilter(b, 1, temperature)
```

The effect of applying the low-pass FIR filter to the signal is to smoothen the function by eliminating the high-frequency oscillations, as shown in Figure 17-11. Another approach to achieve a similar result is to apply a moving average filter, in which the output is a weighted average or median of a few nearby input values. The `medfilt` function from the `signal` module applies a median filter of a given input signal, using the number of past nearby values specified with the second argument to the function.

```
In [71]: temperature_median_filt = signal.medfilt(temperature, 25)
```




Figure 17-11. Output of an FIR filter and a median filter

The result of applying the FIR low-pass filter and the median filter to the hourly temperature measurement dataset is shown in Figure 17-11. Note that the output of the FIR filter is shifted from the original signal by a time delay corresponding to the number of taps in the FIR filter. The median filter implemented using `medfilt` does not suffer from this issue because it is computed from both past and future values, making it a noncasual filter that cannot be evaluated on the fly on streaming input data.

```
In [72]: fig, ax = plt.subplots(figsize=(8, 3))
...: ax.plot(df.index, temperature, label="original", alpha=0.5)
...: ax.plot(df.index, temperature_filt, color="red", lw=2, label="FIR")
...: ax.plot(df.index, temperature_median_filt, color="green", lw=2,
...:         label="median filter")
...: ax.set_ylabel("temperature", fontsize=14)
...: ax.legend(loc=0)
```

To design an IIR filter, we can use the `iirdesign` function from the `signal` module or use one of the many predefined IIR filter types, including the Butterworth filter (`signal.butter`), Chebyshev filters of types I and II (`signal.cheby1` and `signal.cheby2`), and elliptic filter (`signal.ellip`). For example, we can use the following to create a Butterworth high-pass filter that allows frequencies above the critical frequency 7/365 Hz to pass while lower frequencies are suppressed.

```
In [73]: b, a = signal.butter(2, 7/365.0, btype='high')
```

The first argument to this function is the order of the Butterworth filter, and the second argument is the critical frequency of the filter (where it goes from the bandstop to the bandpass function). For this example, the optional argument `btype` can be used to specify if the filter is a low-pass filter (`low`) or high-pass filter (`high`). More options are described in the function's docstring: see, for example, `help(signal.butter)`. The outputs `a` and `b` are the a_k and b_k coefficients that define the IIR filter, respectively. Here, we have computed a Butterworth filter of second order, so `a` and `b` each have three elements.

```
In [74]: b
Out[74]: array([ 0.95829139, -1.91658277,  0.95829139])
In [75]: a
Out[75]: array([ 1.          , -1.91484241,  0.91832314])
```

Like before, we can apply the filter to an input signal (here, we again use the hourly temperature dataset as an example).

```
In [76]: temperature_iir = signal.lfilter(b, a, temperature)
```

Alternatively, we can apply the filter using the `filtfilt` function, which applies the filter forward and backward in time, resulting in a noncasual filter.

```
In [77]: temperature_filtfilt = signal.filtfilt(b, a, temperature)
```

The results of both types of filters are shown in Figure 17-12. Eliminating the low-frequency components detrends the time series and only retains the high-frequency oscillations and fluctuations. The filtered signal can, therefore, be viewed as measuring the volatility of the original signal. This example highlights that the daily variations are greater during the spring months of March, April, and May compared to the winter months of January and February.

```
In [78]: fig, ax = plt.subplots(figsize=(8, 3))
...: ax.plot(df.index, temperature, label="original", alpha=0.5)
...: ax.plot(df.index, temperature_iir, color="red", label="IIR filter")
...: ax.plot(df.index, temperature_filtfilt, color="green",
...:         label="filtfilt filtered")
...: ax.set_ylabel("temperature", fontsize=14)
...: ax.legend(loc=0)
```




Figure 17-12. Output from an IIR high-pass filter and the corresponding `filtfilt` filter (applied both forward and backward)

These techniques can be directly applied to audio and image data. For example, to apply a filter to the audio signal of the guitar samples, we can use the `lfilter` functions. The coefficients b_k for the FIR filter can sometimes be constructed manually. For example, to apply a naive echo sound effect, we can create an FIR filter that repeats past signals with some time delay: $y_n = x_n + x_{n-N}$, where N is a time delay in units of time steps. The corresponding coefficients b_k are easily constructed and can be applied to the audio signal data.

```
In [79]: b = np.zeros(10000)
...: b[0] = b[-1] = 1
...: b /= b.sum()
In [80]: data_filt = signal.lfilter(b, 1, data)
```

To listen to the modified audio signal, we can write it to a WAV file using the `write` function from the `io.wavfile` module in SciPy.

```
In [81]: io.wavfile.write("guitar-echo.wav", sample_rate,
...:                    np.vstack([data_filt, data_filt]).T.astype(np.int16))
```

Similarly, we can implement many types of image processing filters using the tools from the `signal` module. SciPy also provides a module `ndimage`, which contains many common image manipulation functions and filters specially adapted for applying two-dimensional image data. The Scikit-Image library⁵ provides a more advanced framework for working with image processing in Python.

Summary

Signal processing is a vast field with applications in most fields of science and engineering. As such, we have only covered a few fundamental applications of signal processing in this chapter, focusing on introducing methods for approaching this type of problem with computational methods using Python and the libraries and tools available within the Python ecosystem for scientific computing. In particular, the chapter explored spectral analysis of time-dependent signals using fast Fourier transform and the design and application of linear filters to signals using the `signal` module in the SciPy library.

Further Reading

For a comprehensive review of the signal processing theory, see *The Scientist and Engineer's Guide to Digital Signal Processing* by Steven Smith (1999), which can be viewed online at www.dspguide.com/pdfbook.htm. For a Python-oriented discussion of signal processing, see *Python for Signal Processing* by J. Unpingco (Springer, 2014), from which content is available as IPython notebooks at <http://nbviewer.org/github/unpingco/Python-for-Signal-Processing>.

⁵ See the project's web page at <http://scikit-image.org> for more information.